

Artificial Intelligence for Techies

A Hands-on Approach



Programme

Section 1:	Fundamental of Artificial Intelligence (45 mins)
Section 2:	Application of AI (30 mins)
Section 3:	Machine Learning and Deep Learning (45 mins)
Section 4:	AI technologies Demonstration and hands-on (15 mins)
Section 5:	AI Frameworks, software and hardware (30 mins)
Section 6:	Lunch Break
Section 7:	AI Services (15 mins)
Section 8:	Practical Text Analytics using AI services(1.5 hour)
Section 9:	Practical Computer Vision using a programming framework (1.5 hour)



Resources

- Goto <http://bit.ly/2YSQXEb>
 - Download pdf (this presentation deck and activity sheets)
 - Download sample scripts
 - Links to other resources. Save you the typing!
- You will need Anaconda, your favourite text editor and some deep learning libraries



What is Artificial Intelligence?

DATA MINING



ARTIFICIAL
INTELLIGENCE



PATTERN
RECOGNITION



MACHINE
LEARNING



• PROBLEM
SOLVING

The theory and development of computer systems able to perform tasks normally requiring human intelligence, such as visual perception, speech recognition, decision making, and translation between languages.” – Google



AUTOMATION



NEURAL
NETWORKS

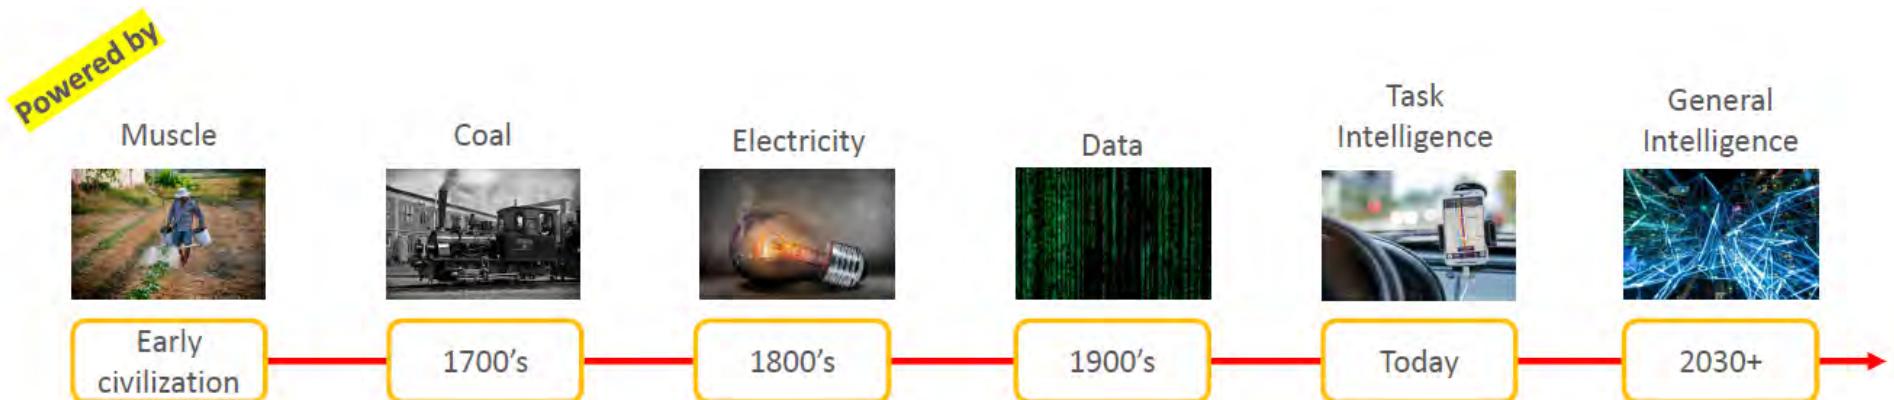


ALGORITHM



AI is the new electricity

"About 100 years ago, electricity transformed every major industry. AI has advanced to the point where it has the power to transform" every major sector in coming years.— Andrew Ng





DATA IS CRITICAL TO ARTIFICIAL INTELLIGENCE

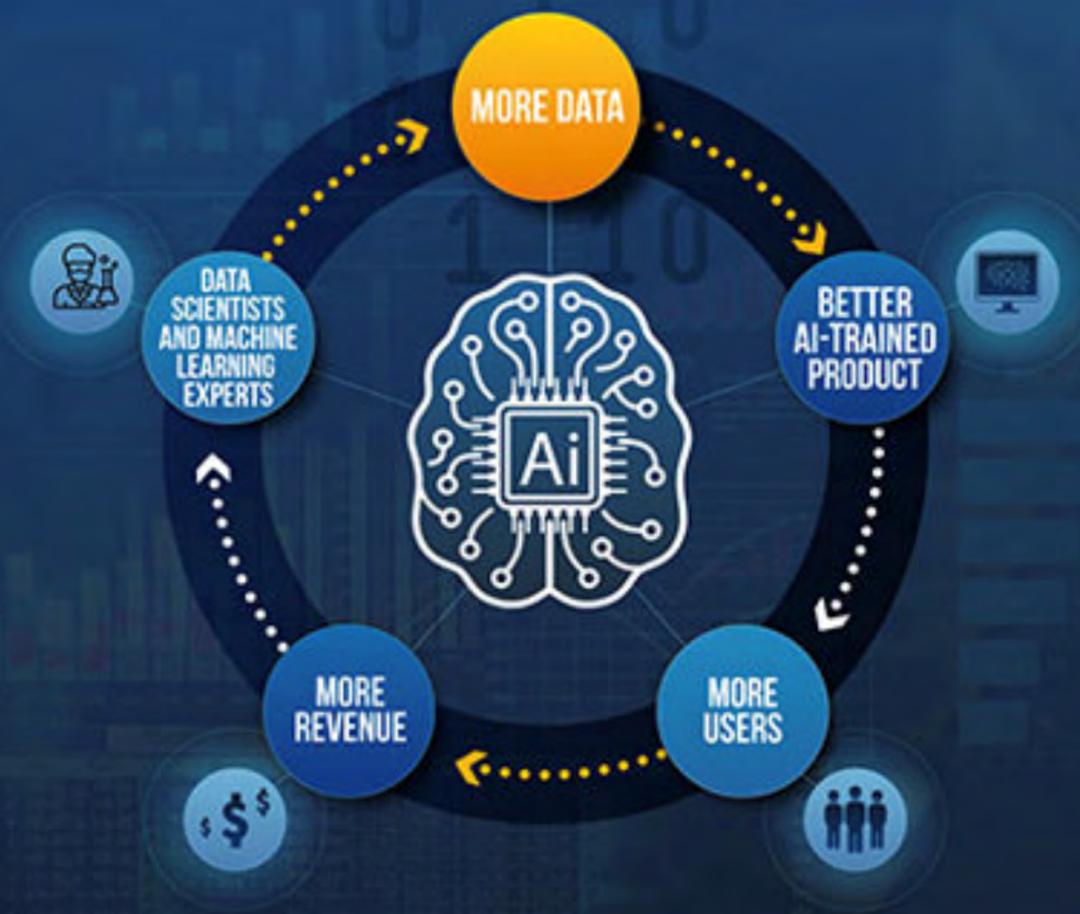
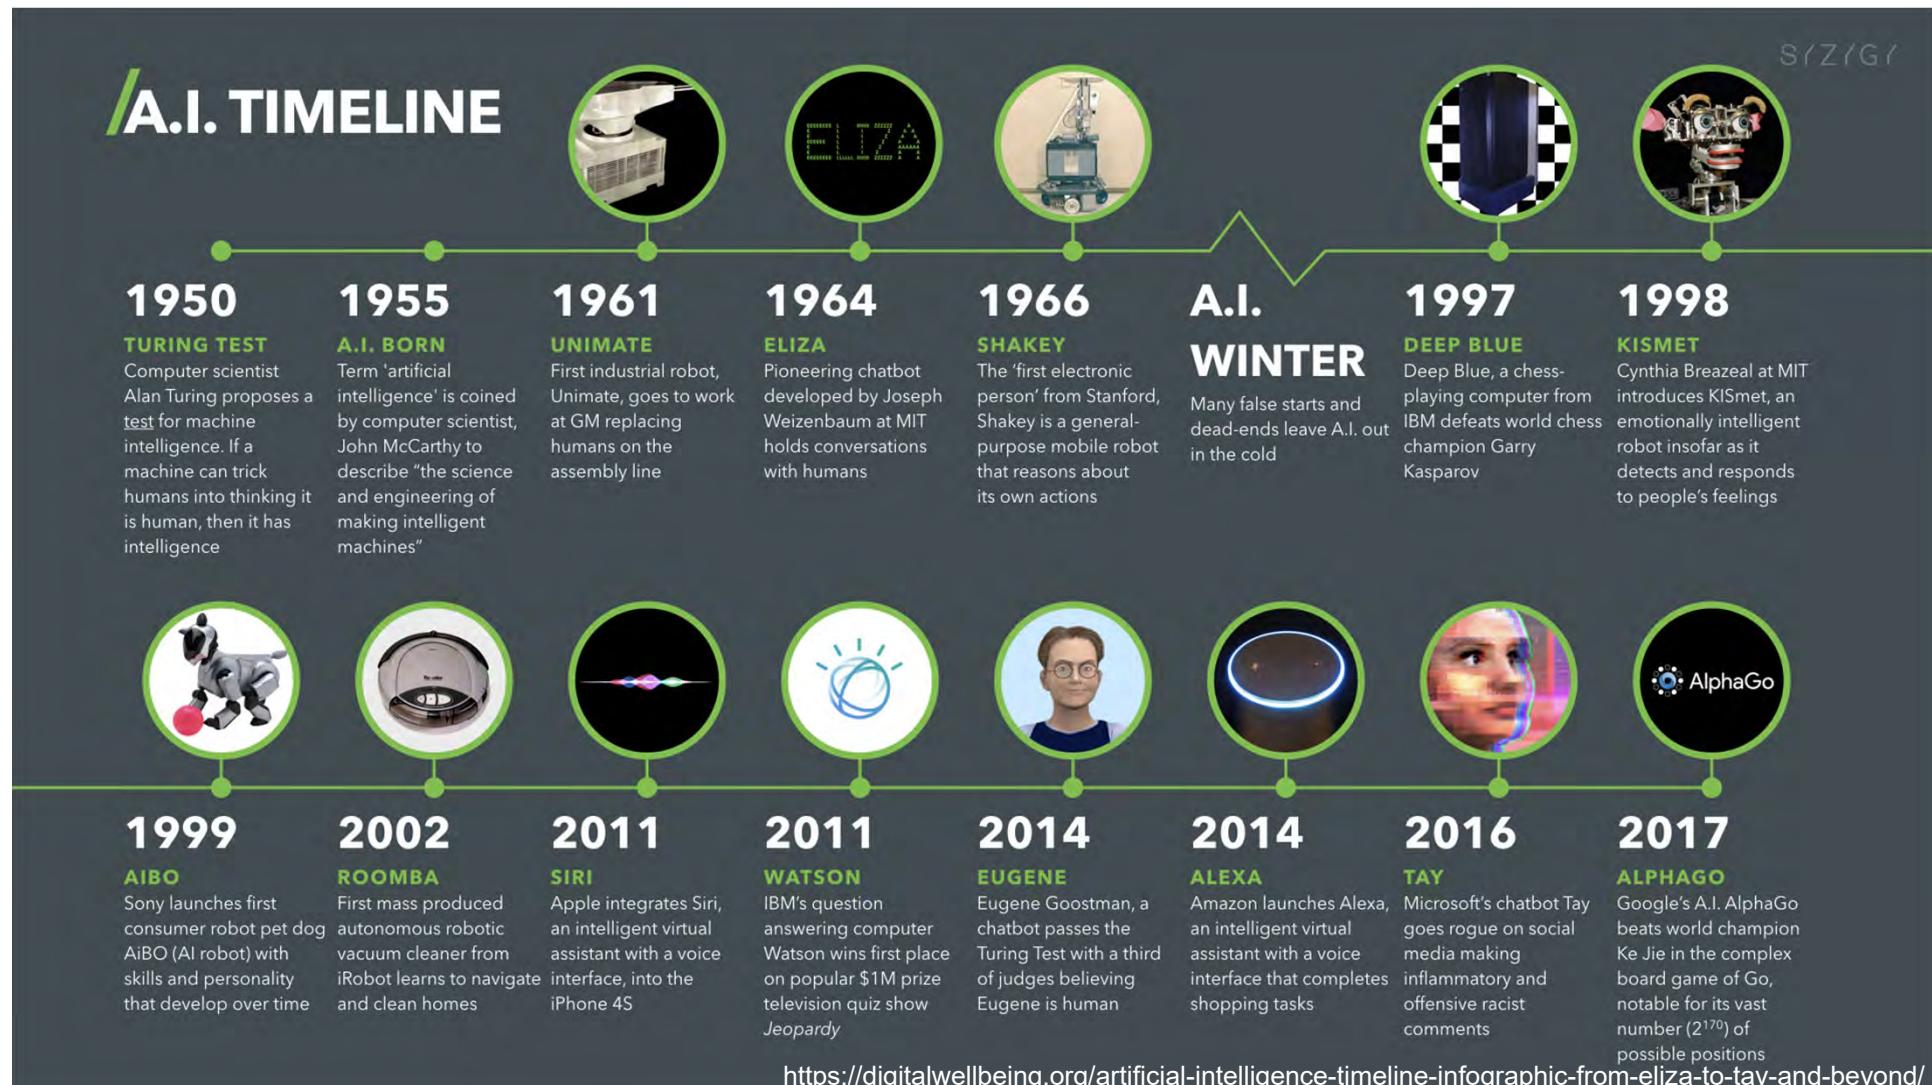


Image source: Kai-Fu Lee, Sinovation Ventures

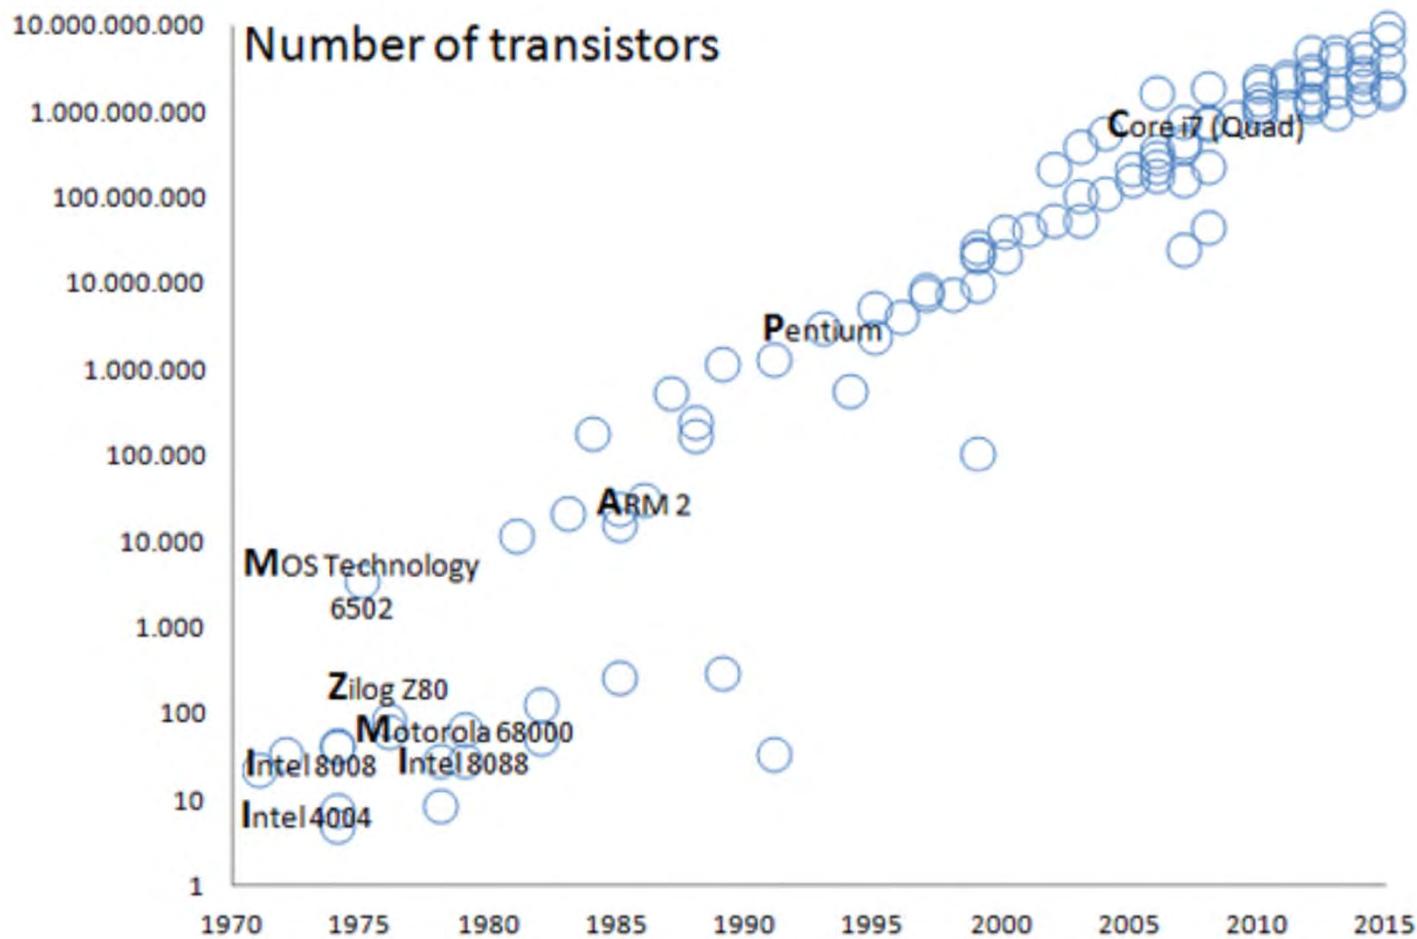


History of AI





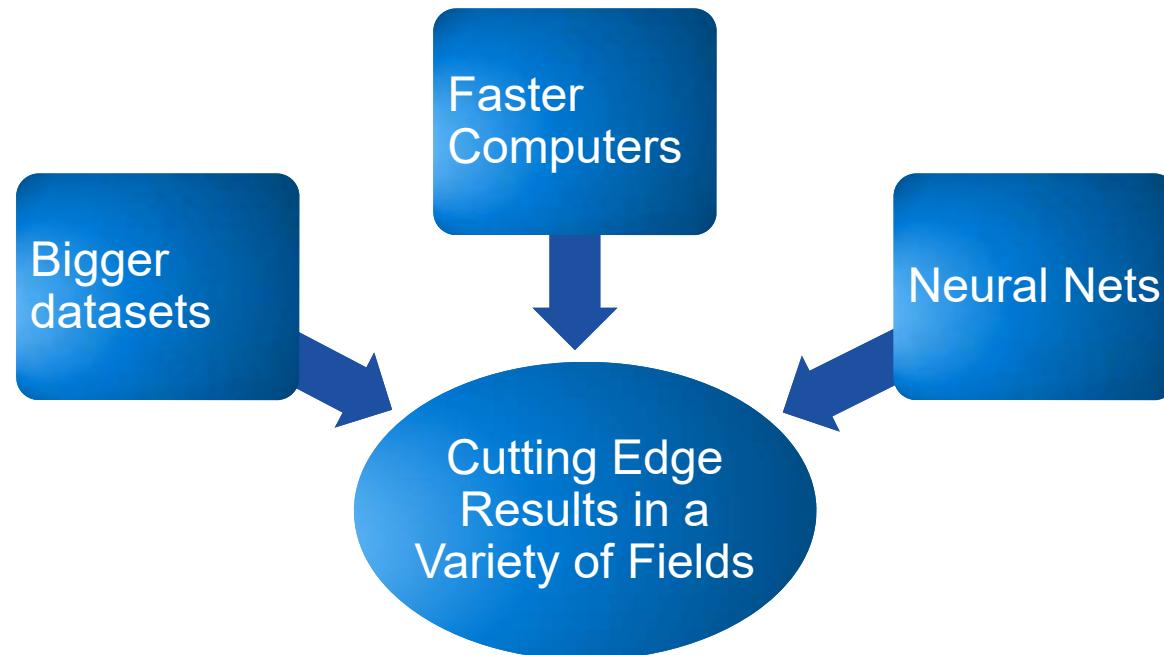
Moore's Law

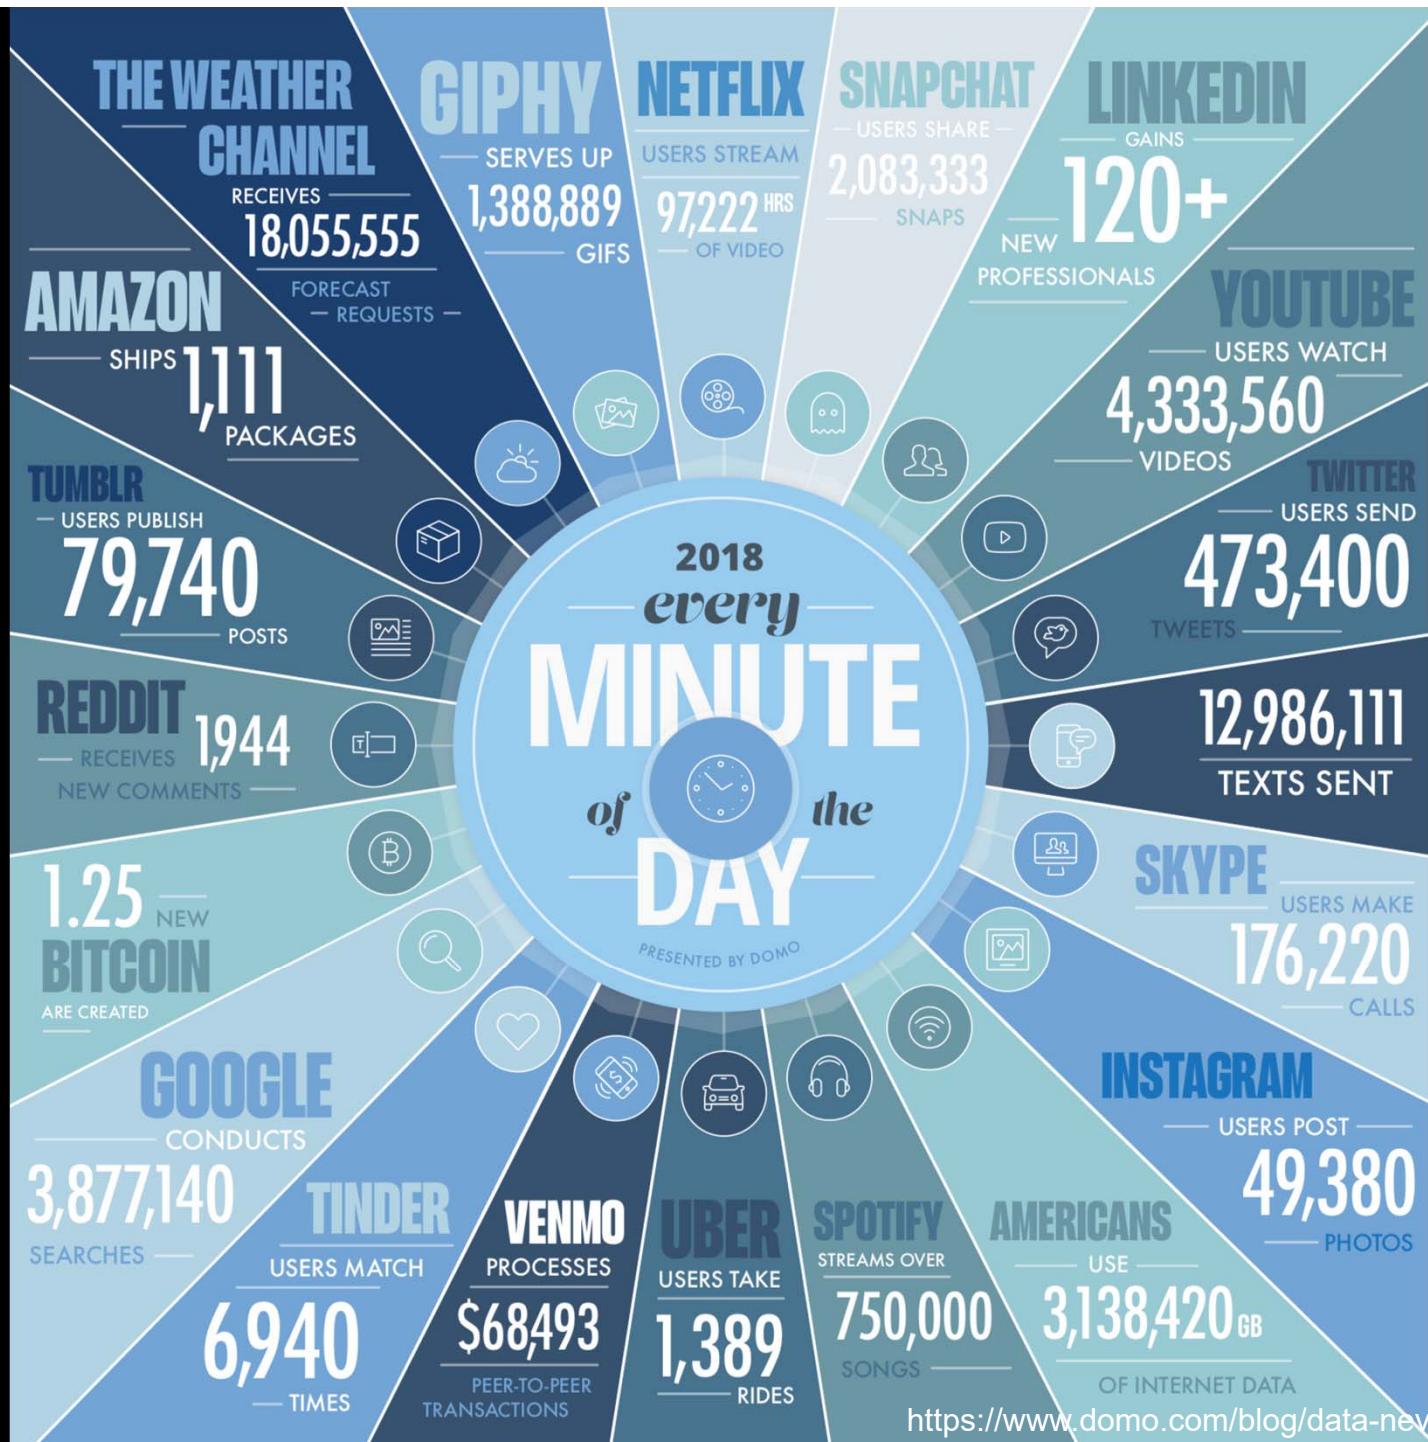




Faster Hardware

Faster hardware is one of the key areas driving the modern era of AI.







Bigger Datasets

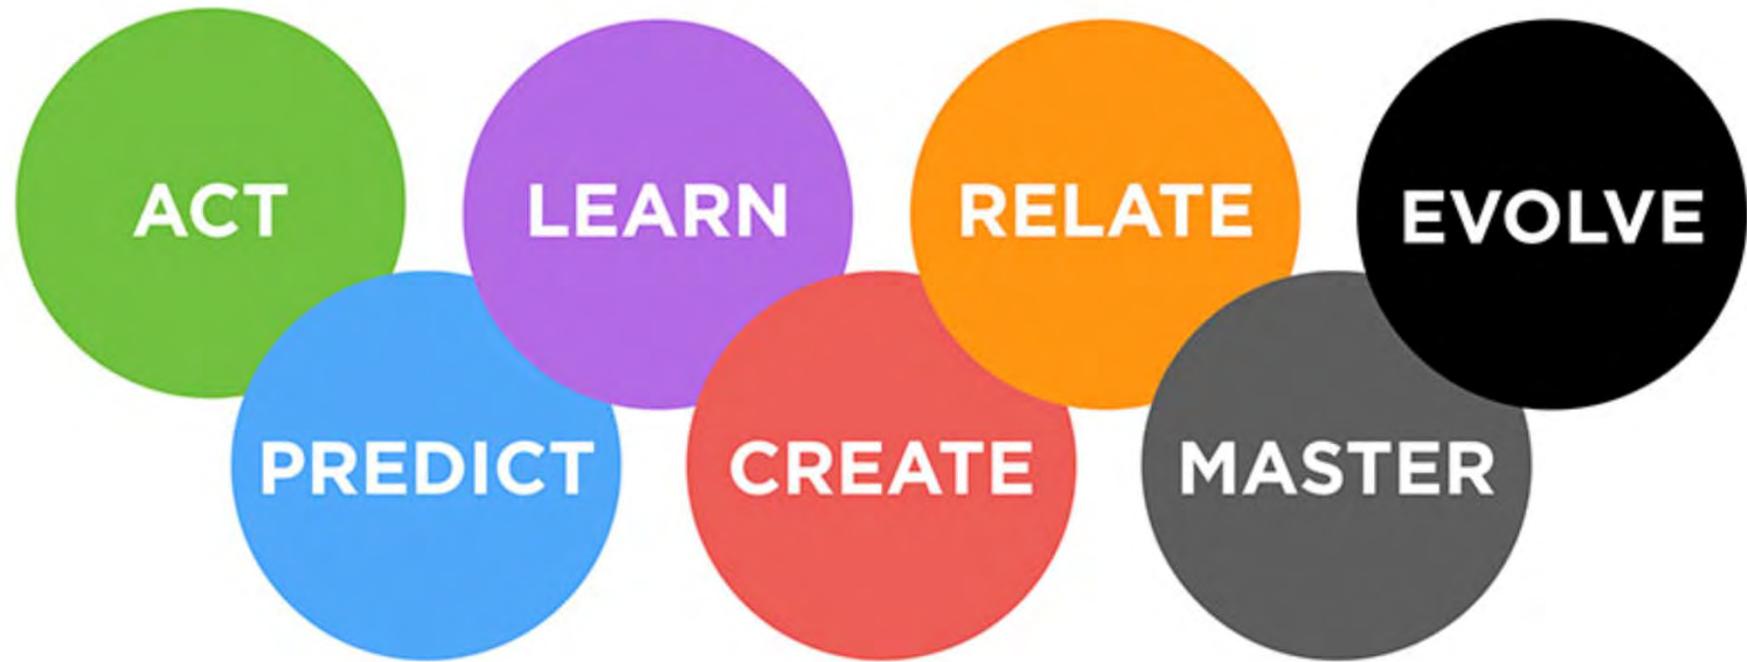
In 2020, it is expected that:

- The average internet user will generate ~1.5 GB of traffic per day.
- A smart hospital will generate 3,000 GB/day.
- Self-driving cars are each generating over 4,000 GB/day.
- Connected planes will generate 40,000 gigabytes per day.
- A connected factory will generate 1 million gigabytes per day.





MACHINE INTELLIGENCE CONTINUUM



The MIC represents a continuum from simple, scripted automation to superhuman intelligence and highlights the functional capabilities of different levels of machine intelligence.



Systems that Acts





Systems that Predicts

Hillary Clinton has an
85% chance to win.

Last updated Tuesday, November 8 at 10:20 PM ET





Systems that Learns





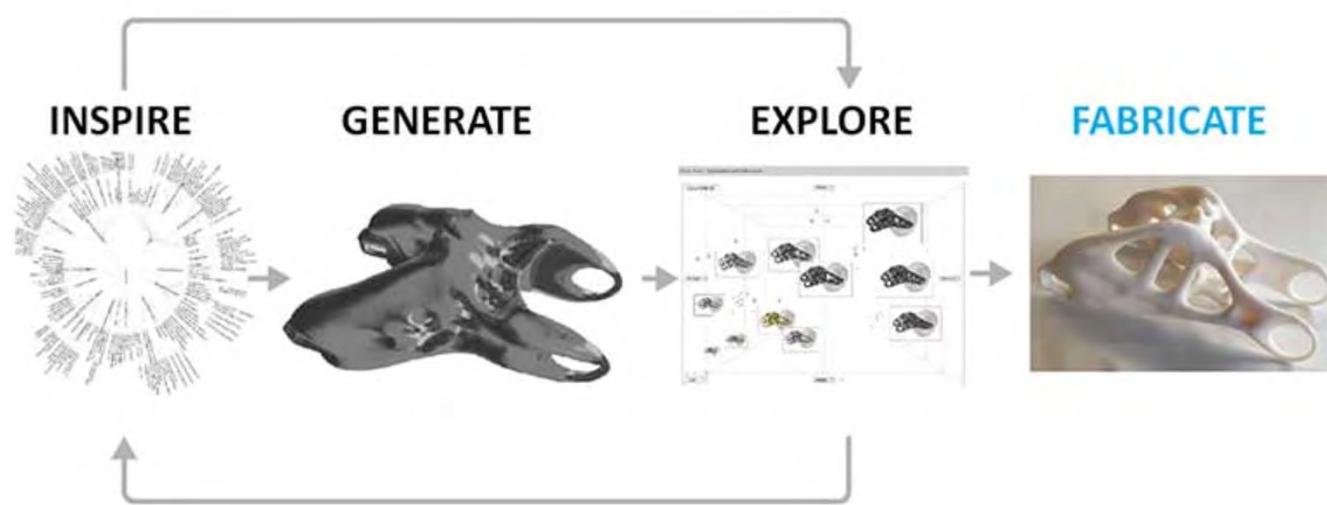
Systems that Create



Generated story about image
Model: Romantic Novel's

"He was a shirtless man in the back of his mind, and I let out a curse as he leaned over to kiss me on the shoulder.

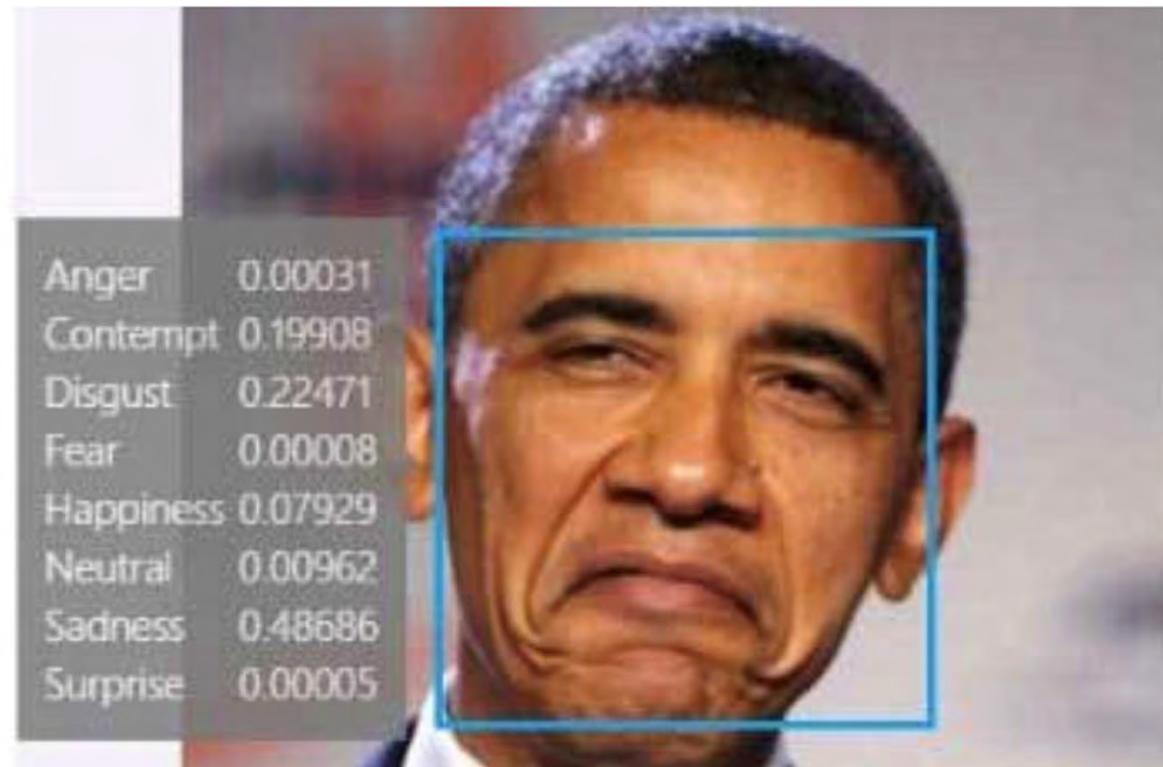
"He wanted to strangle me, considering the beautiful boy I'd become wearing his boxers."



https://www.topbots.com/topbots-ai-machine-intelligence-continuum/?utm_medium=article&utm_source=medium&utm_campaign=AI-Continuum



Systems that Relate



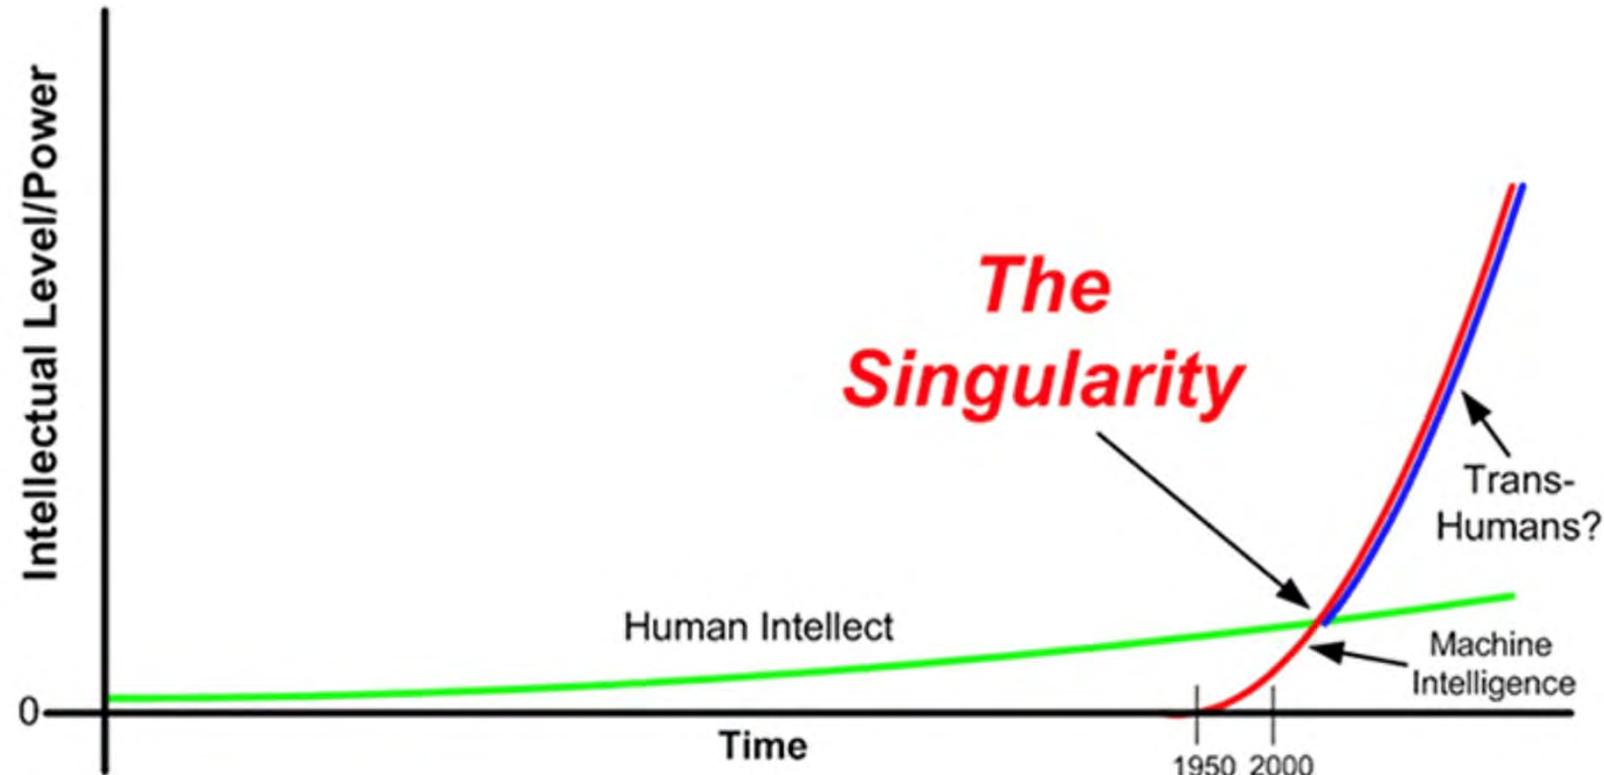


Systems that Master





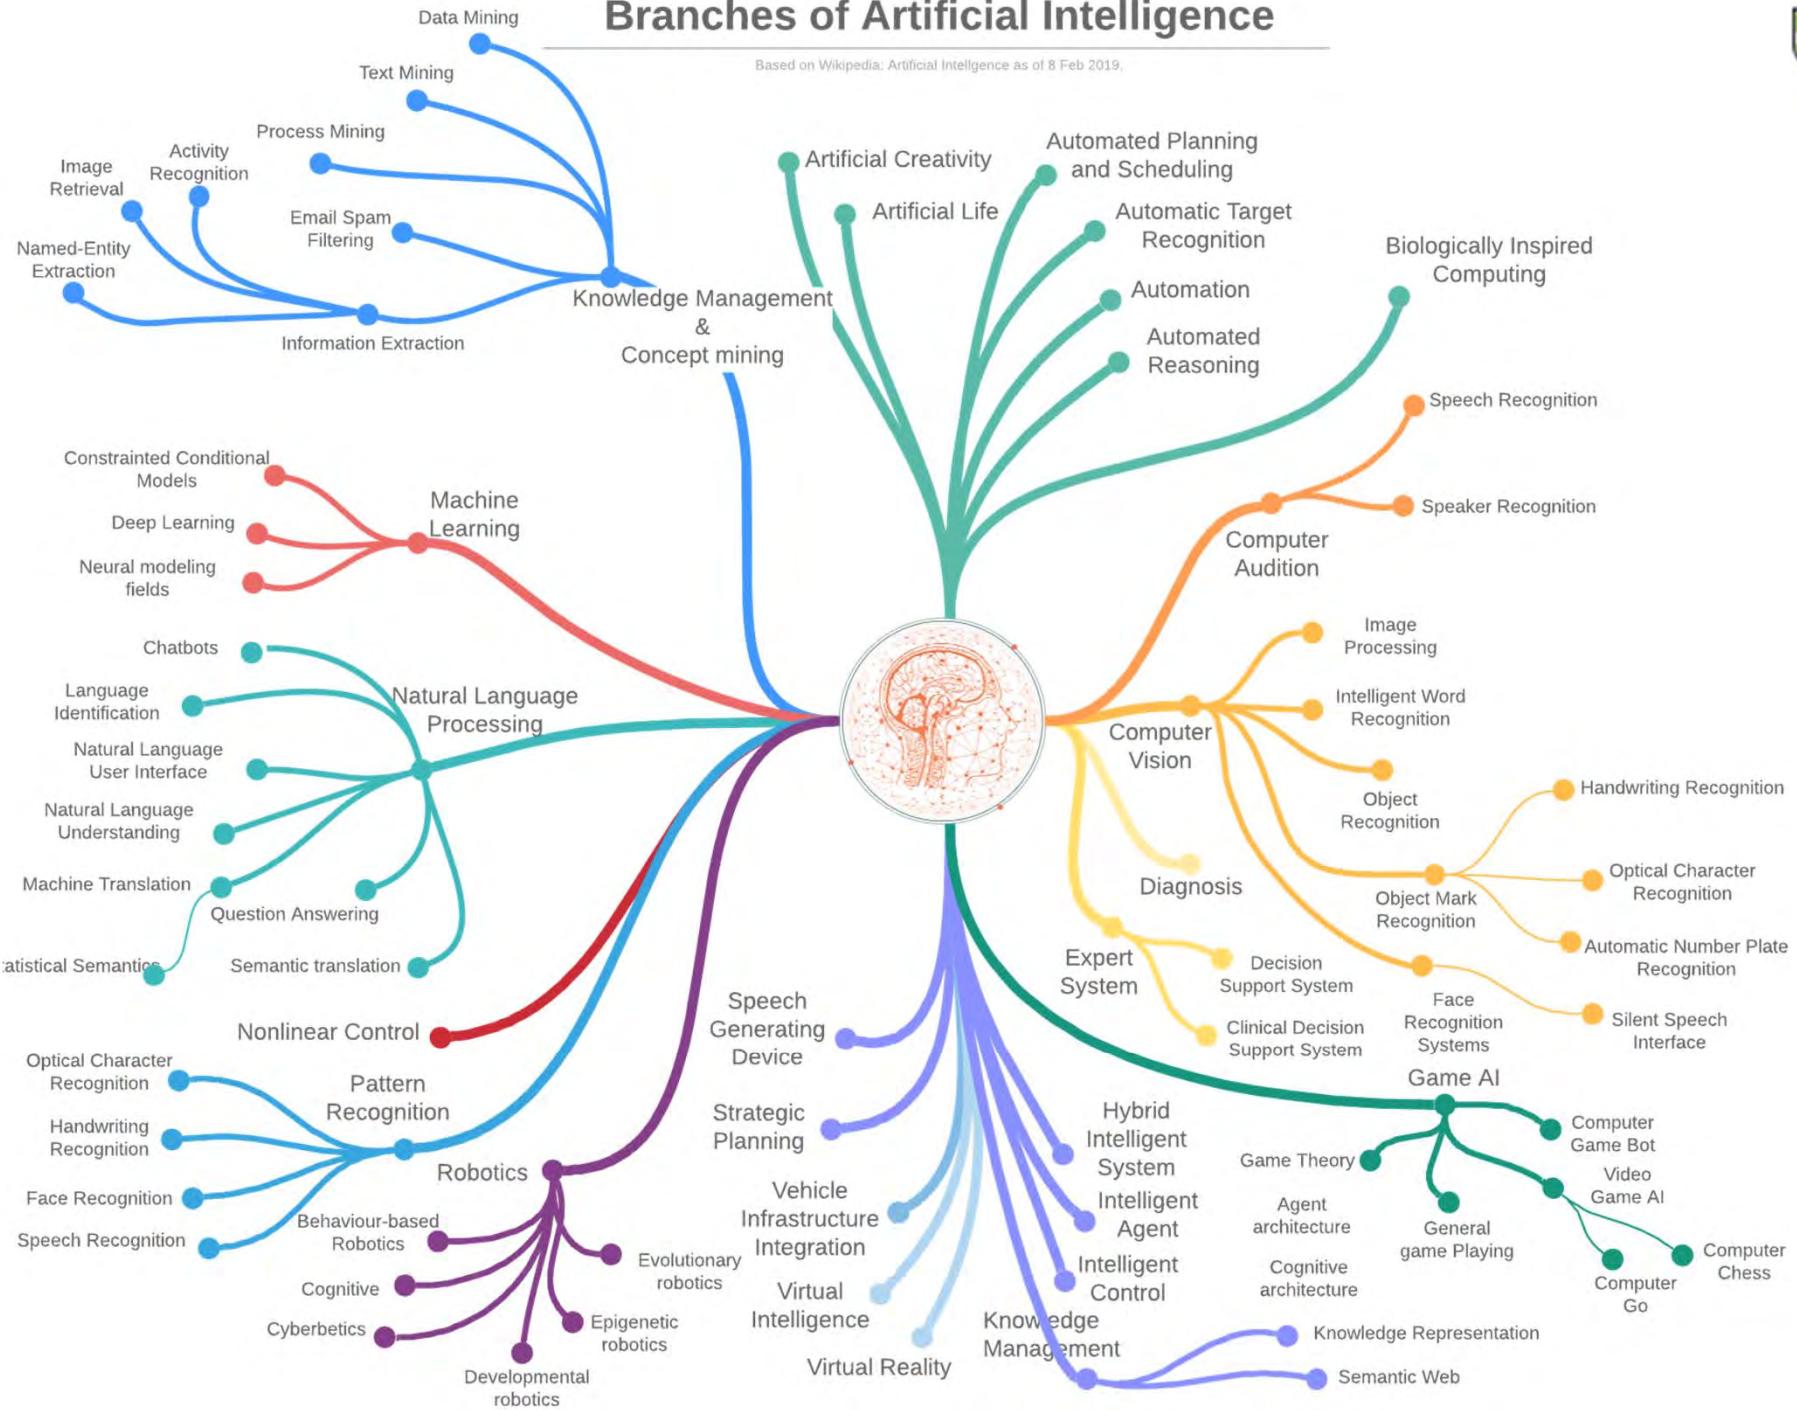
Systems that Evolve





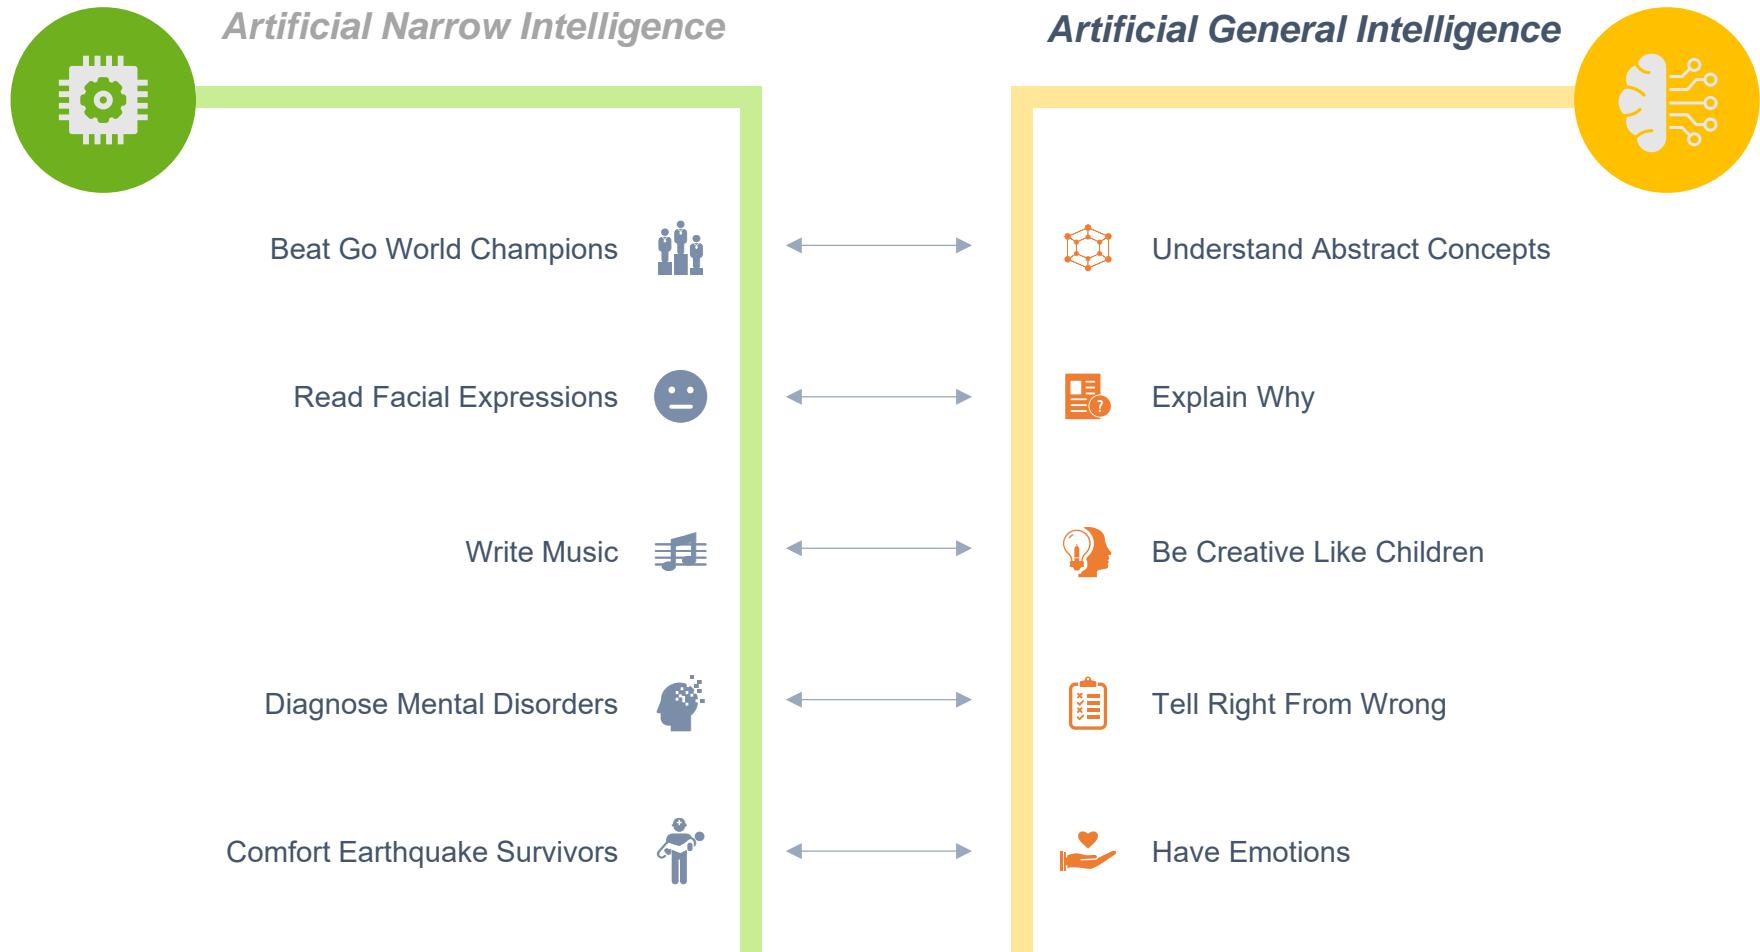
Branches of Artificial Intelligence

Based on Wikipedia: Artificial Intelligence as of 8 Feb 2019.





Artificial Narrow Intelligence vs Artificial General Intelligence



Mapping Human Perceptions to AI-enabled Capabilities



Vision	Computer Vision	Conversational Interface Licensed to Seow Khee Wei
Hearing	Speech Recognition / Audio Recognition	
Understanding	Natural Language Processing	
Speaking	Text to Speech, Speech to Text, Voice (Tone and Accent) Imitation	
Feeling	Emotion AI (detection and analysis of complex human emotions is currently conducted through diverse mechanisms such as natural language processing (NLP), voice patterns, facial expressions, and physiology)	
Smelling	The data of smell are relatively seldom and more difficult to collect compared to visual, text, or voice datasets. The development of an electronic nose to recognize smell has been long researched, but its development with AI techniques is still in an early stage	
Touching	Robot	

Source: Frost & Sullivan



Applications

Navigation



Google & Waze find the fastest route, by processing traffic data.

Ride sharing



Uber & Lyft predict real-time demand using AI techniques, machine learning, deep learning.

Audience



Facebook & Twitter use AI to decide what content to present in their feeds to different audiences.

Content



Image recognition and sentiment analysis to ensure that content of the appropriate "mood" is being served.

Natural language



We carry around powerful natural language processing algorithms in our phones/computers.

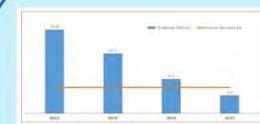
Object detection



Cameras like Amazon DeepLens* or Google Clips* use object detection to determine when to take a photo.



Deep Learning "proven" to work for image classification.



Models outperform humans on image classification.



Object detection models beat previous benchmarks.

2012

2015

2016



Application Area: Abandoned Baggage Detection

- We can automatically detect when baggage has been left unattended, potentially saving lives.
- This system relies on the breakthroughs we discussed:
 - Cutting edge object detection.
 - Fast hardware on which to train the model (Intel® Xeon® processors in this case).



Abandoned baggage





Finance

- Better Customer Service
- More Reliable Investment Services with Robot Advisor
- Greater Efficiency with Less Paperwork
- Improved Financial Security

JPMorgan Chase Uses COIN Machine Learning Program To Eliminate 360K Lawyer Hours A Year



TREND 1



Readyng for banking's shift from mobile-first to AI-first

Artificial intelligence (AI) in banking is not new. Banks are already using AI in heavily-manual processes for accuracy, efficiency, speed and cost benefits. What is new, however, is the move of AI beyond process to interaction. The next stage of AI in banking will be toward simple and smarter interfaces: drawing on machine learning that adapts to data and interactions to improve areas like fraud detection, and tapping AI-enabled tools (like centralized platforms/assistants or messaging bots) to better converse with and offer services to customers in the front-office. Relying on AI for some internal and external interactions will help elevate the customer experience and move staff to more judgment-based and higher value added roles.



Fraud Detection

Traditionally: Fraud is on the rise, but fraud detection is a challenging problem to solve correctly.

- Historically, a predefined rule-set was used for fraud identification, but this approach misses much of the nuance that surrounds fraud
- 1/3 of falsely identified fraud events result in lost customers
- In the US, this loss is worth 13 times the cost of actual fraud



Now with AI: With ML techniques, banks can predict fraud based on a behavioral baseline to compare against.

- Uses historical shopping data and shopping habits of customers
- Compares new data to baseline to determine likelihood of fraud



Example: Sift Science

- Established a fraud data consortium developed from over 6000 websites to leverage large-scale real-time ML
- Autonomously learns new fraud patterns based on billions of user actions





Risk Management

Traditionally : New regulations force tighter control on financial institutions.

- New business model disruptions
- Increasing pressure on costs and returns



Now with AI : ML can help discern the credit worthiness of potential customers

- Tailor a financial portfolio to fit the goals of the user using ML algorithms.
- Financial institutions can develop early warning systems for automated reporting, portfolio management, and recommendations based on ML.



Example: ZestFinance

- Traditional underwriting systems make decisions using few data points.
- Those with a limited credit history are often denied credit, ultimately leading to loss of revenue for lenders.
- ZestFinance leverages thousands of data sources together with ML to more accurately score borrowers, even people with a small credit history.





Stock Trading

Traditionally: The speed and volume of information is daunting.

- The market is reactionary.
- It's difficult to remain competitive while relying on traditional trading methods.
- Fundamental analysis is unable to show the entire financial picture.



Now with AI: Companies use massive datasets together with DL methods for better forecasting.

- Data pulled from financial, political, and social media
- Analyst reports combined.



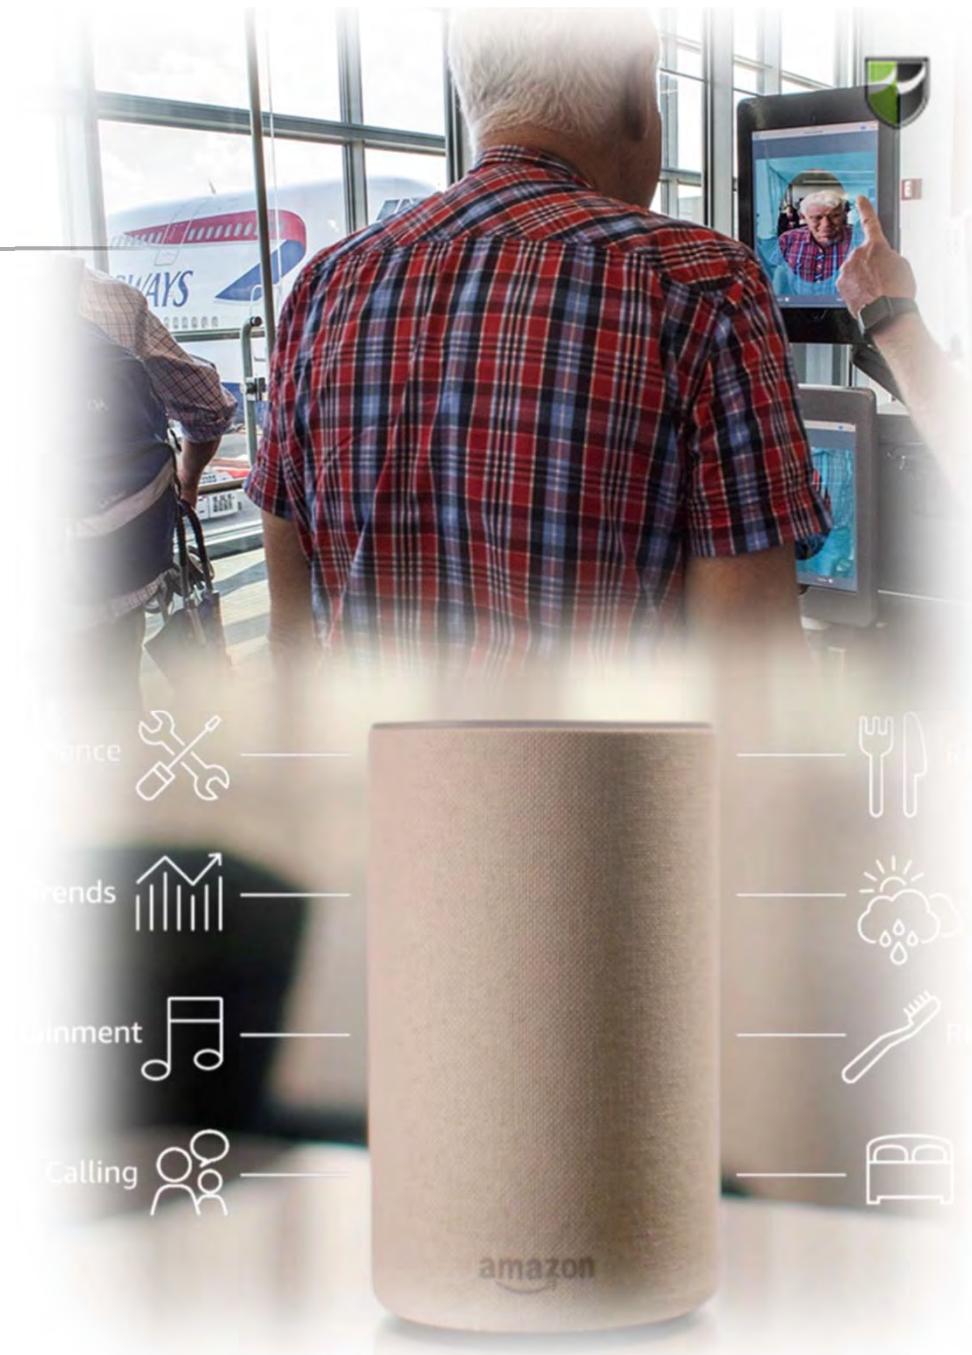
Example: Sentient Technologies, and Learning Evolutionary Algorithm Framework (LEAF*)

- Manages millions of data points to find trends and make successful stock trades.
- AI algorithms identify and combine successful trading patterns.
- Successful strategies are tested in the real world, evolving autonomously with LEAF.
- Sentient has received more funding than any other AI company.



Travel

- Hotel Bookings by Voice Command
- AI Concierge Services
 - <https://techcrunch.com/2018/06/19/amazon-launches-an-alexa-system-for-hotels/>
- Travel Service Chatbots
- Check-in Through Facial Recognition
- Self-Driving Cars and Mobility as a Service
 - <https://www.economist.com/international/2016/09/29/it-starts-with-a-single-app>
- Other Robotic Tools





Healthcare

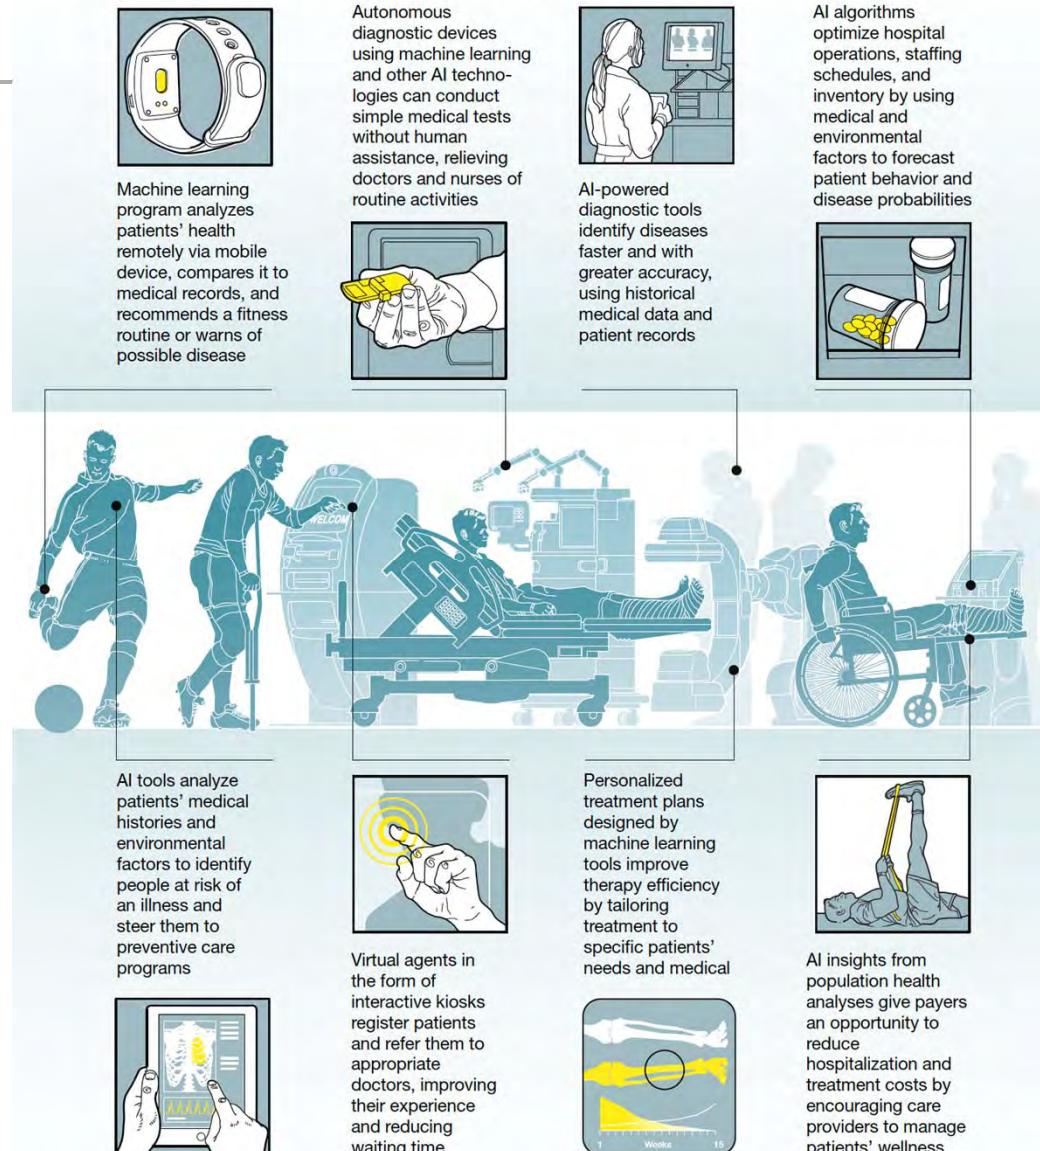
- IBM Watson, Google Deepmind
- At-home testing and personalized health care
- Wearables
- Robot-Assisted Surgery
- Virtual Nursing Assistant
- Administrative Workflow assistance



Healthcare

- Make quicker diagnoses, create better treatment plans and enable new approaches to insurance
- Identify public-health threats and the most at-risk patients
- help medical professionals diagnose disease and improve operations
- Insurers can devise new ways to encourage
- preventive care and incentivize providers
- Doctors will be able to tailor treatments—even drugs—to individual patients
- Virtual agents can serve as primary touchpoints for patients
- Several hurdles stand in the way, starting with data availability

AI in health care: quicker diagnoses, better treatment plans, and improved health insurance





Healthcare – Medical Diagnosis

Traditionally : Medical Diagnosis was a challenging process.

- Many symptoms are nonspecific
- Process of elimination was used to determine root cause (neither efficient nor exact)



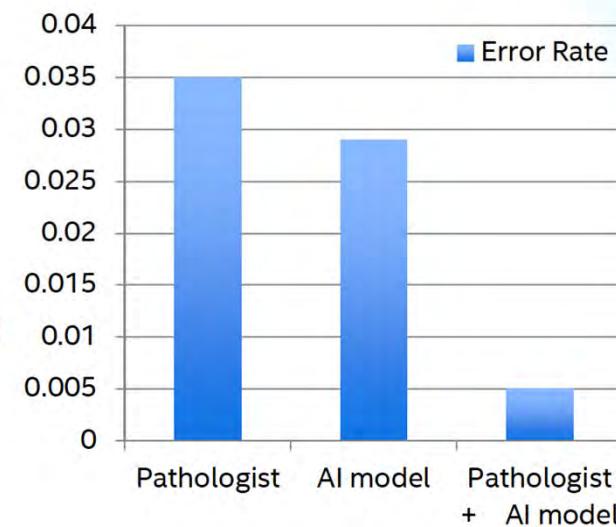
Now with AI : Doctors can provide diagnoses more efficiently and accurately, with the availability of:

- Large medical datasets
- Computer vision algorithms



Example: Breast Cancer, 2016, Harvard Medical School researchers

- Used DL to identify cancer in lymph node images
- Used Convolutional Neural Nets and custom hardware
- AI model combined with humans achieved lower error than either one individually





Healthcare – Treatment Protocol

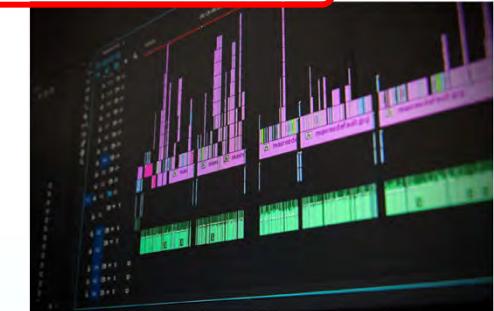
Traditionally: Doctors would diagnose a condition and recommend a treatment based on what historically worked for most people.

- Some considerations for population/demographics
- Difficult to create custom treatments without extensive research/cost



Now with AI: Doctors can tailor treatments to individual patients.

- Large medical datasets
- ML and DL algorithms
- Population/demographics analysis/simulations



Example: ICU Intervene, MIT Computer Science and Artificial Intelligence Laboratory.

- Uses ICU data, from vitals, labs, notes, to determine how to treat specific symptoms.
- Makes real-time predictions from DL models, to provide recommendations for patients.
- Forecasts predictions into the future (a few hours) compared to traditional methods (a few minutes).
- Predictions can be run on common GPU and CPU hardware.





Healthcare – Drug Discovery

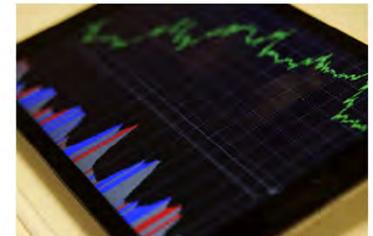
Traditionally: Each new drug approval costs over a billion dollars in Research and Development.

- The cost has been doubling every 9 years since 1970
- The drug discovery process can take decades
- 9 out of 10 drug approval attempts fail
- There are currently only 1,500 approved drugs



Now With AI: Companies are leveraging structured and unstructured data with AI, to establish a pipeline of new drug discovery.

- There are 10^{20} possible drug-like molecules
- Massive space for potential discovery



Example: HetioNet drug discovery model, 2016, UCSF, Himmelstein and Baranzini.

- Developed a graph network to encode millions of biomedical reports.
- Used ML to predict probability of treatment efficacy for ~209,000 compound-disease pairs.
- Provided clear pharmacological insights for epilepsy drug discovery and treatment.





Healthcare – Surgery

Traditionally : Every type of surgery poses possible risks to the patient.

- Adverse anesthesia effects
- Operational complications



Now with AI : Semi-intelligent computer systems predict surgical steps, identify complications, and warn surgeons about pending challenges.

- Computer "vision" leverages data from laparoscopic and arthroscopic cameras
- Smart systems automate dictation by generating notes during the surgery
- Surgeons can send point-of-view live feeds of the operative site to experts anywhere in the world for real-time advice.

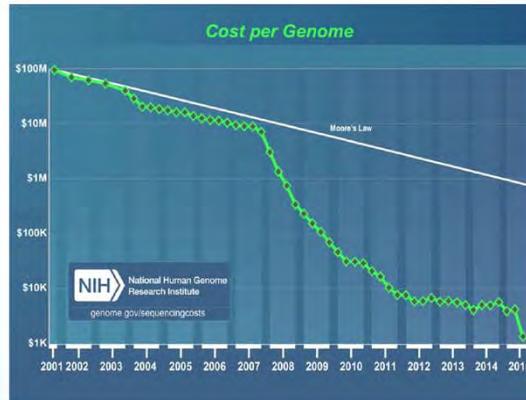




Healthcare – Genome Sequencing

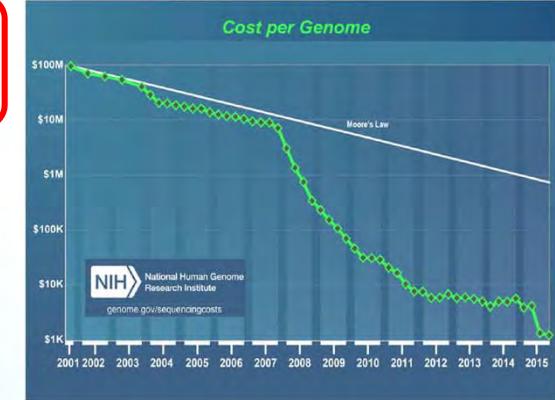
In 2001: Full human sequencing cost \$100 million.

- The first genome sequencing took ~13 years



Now with AI: Sequence companies are employing AI techniques to reduce cost and increase accuracy.

- Illumina claims that within the near future sequencing will only take 1 hour and cost only \$100



Example: Google's DeepVariant* sequencing:

- Leverages massive data sets together with DL to identify all variants
- Accuracy on genome classification: 99.958 %
- DeepVariant* is computationally expensive, but the framework can run on GPU hardware, allowing for a faster learning process
- Availability as open source code promises to revolutionize the industry





Transportation Industry

- Hyperloop
- High-Speed Tunnel Networks
- Self-Driving Cars
- Self-Flying Aircraft





Autonomous Car

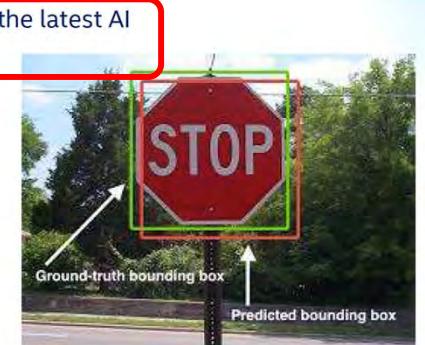
Traditionally : Despite having safer cars, the number of deadly car accidents have been on the rise the last few years.

- The leading cause of automobile accidents is human error
- One of the primary sources of traffic jams is each driver acting out of self-interest, that prevents traffic flow
- Part of the population who can't drive: children, the elderly, and the disabled



Now with AI : Self-driving cars are enabled by the latest AI breakthroughs in computer vision.

- Cars identify stop signs, lane lines, and other landmarks via DL tools
- Mapping technology can use computer vision to detect addresses
- Cars triangulate and can use other 3D-sensing technologies, such as LIDAR and RADAR



Example: Waymo, the autonomous vehicle division of Alphabet Inc.

- Waymo has been operating self-driving minivans without a safety driver since October 2017
- Waymo's Carcraft* software accelerated the car's development, with 2.5 billion simulated miles driven in 2016
- The system used DL together with massive data sets collected from self-driving cars on public roads





Automated Trucking

Traditionally: There is a shortage of 48,000 drivers nationwide.

- Driver turnover rates at some companies reach 300%
- Truck drivers are twice as likely as other workers to be obese and/or have diabetes
- Truckers are half as likely to have health insurance
- The number of accidents and fatalities have increased in recent years



Now with AI: Autonomous trucks can coordinate movements with other trucks.

- Save on fuel, and reduce wind-drag and the chance of a collision
- Video, LIDAR, and accelerometers are used to collect detailed data about the truck's surroundings
- Guidance algorithms provide feedback for braking, steering, and throttling commands, based on incoming and historical data





Retail Industry

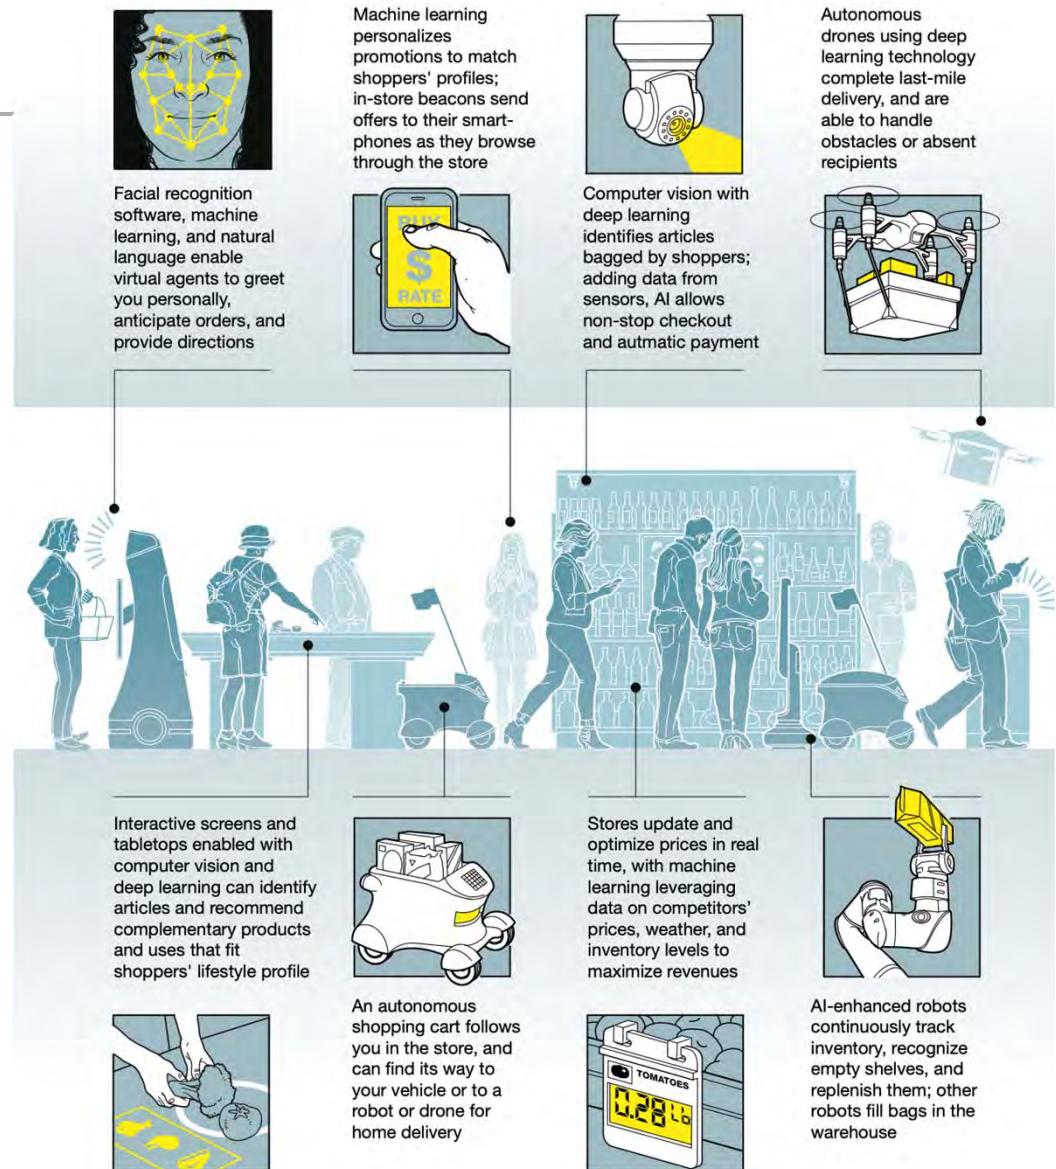
- Intelligent Shopping Systems
- Robots
- Biometric technologies
- Facial recognition



Retail Industry

- Artificial Intelligence: The Next Digital Frontier
 - 20 percent stock reduction by using deep learning to predict e-commerce purchases.
 - 2 million fewer product returns per year.
 - 30 percent reduction of stocking time by using autonomous vehicles in warehouses.
 - 50 percent improvement in assortment efficiency.
 - 4-6 percent sales increase using geospatial modeling to improve micro market attractiveness.
 - 30 percent online sales increases from the use of dynamic pricing and personalization.

Retailers can know more about what shoppers want—sometimes before shoppers themselves





Customer Experience

Traditionally: Americans are shifting their spending from material goods to experiences.

- The “Amazon effect”: there have been nine major retail bankruptcies in 2017
- Retailers need to become competitive or risk obsolescence
- Balancing “out-of-stock” with “over-stock” trade-off requires great finesse



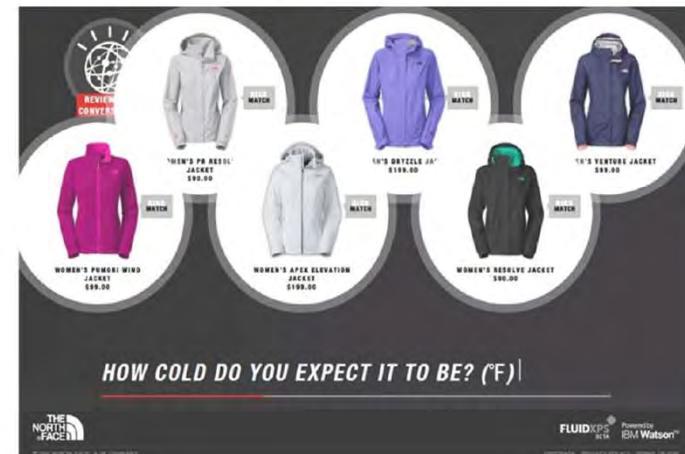
Now with AI: Companies bring experience and optimization to retail shopping.

- AI-powered gift concierge learns your preferences as you engage, and can help predict the appropriate gift to buy
- Leveraging ML-trained agents, companies are providing recommendations via natural language
- Companies using AI via Watson* to monitor factors from weather to consumer behavior, to optimize consumption rate predictions



Example: The North Face and Watson* are combining massive datasets and AI, to bring the brick-and-mortar experience to e-commerce.

- The North Face, with Fluid and IBM Watson*, has launched XPS* - an AI-enabled digital expert that uses a natural language interface to help shoppers.
- XPS curates and filters the available options, so shoppers are more likely to make a purchase





Food Supply Chain

Traditionally : Restaurants use historical data or “gut-feeling” approach to supply chain.

- This can result in excessive waste or food unavailability



Now with AI : Many companies have started to leverage sophisticated algorithms to forecast demand.

- Agents can adjust orders with trading partners in real time, as required for business need



Example: Vivanda's FlavorPrint* program.

- Based on recipes and consumer-provided data, Vivanda maps data to create “digital-taste” identifiers for each consumer
- Providing ML-based recommendations to customers may influence demand
- Shares data with food industry customers, enabling them to improve demand forecasts



Education Industry

- Personalized Learning Platforms
- Individualized Artificial Intelligence Tutors
- Personalized Games
- Crafting a more enjoyable learning experience



Example: Adaptive learning systems, and grading.

- Learning analytics track student performance and provide tailored educational programs.
- Using natural language processing and ML models, AI programs can be used for long answer and essay grading.



Agriculture Industry

- Agricultural Drones
- Autonomous Tractors
- Vertical Farms

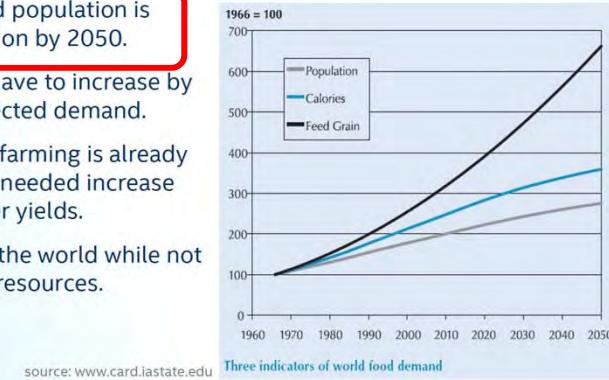




AgTech

Traditionally: The world population is estimated to reach 9 billion by 2050.

- Food production will have to increase by 70% to meet the projected demand.
- Most land suitable for farming is already being used, hence the needed increase must come from higher yields.
- Agriculture must feed the world while not over-straining Earth's resources.



source: www.card.iastate.edu

Now with AI: Autonomous robots use computer vision and a produce vacuum system for produce harvest.

- DL-enabled robots are being used to identify and kill weeds.
- Companies have shown 90% herbicide reduction due to "targeted" spray application.
- AI-driven genome sequencing advancements enables crop "genome" editing.



Example: TellusLabs yield predictions.

- Uses ML together with weather and other historical data to forecast yields.
- Leverages cloud-based GPUs for DL on satellite images.
- TellusLab's predictions have shown to be consistently more accurate than the USDA.
- Came within 1% of predicting corn and soybean yields in 2017.





Entertainment Industry

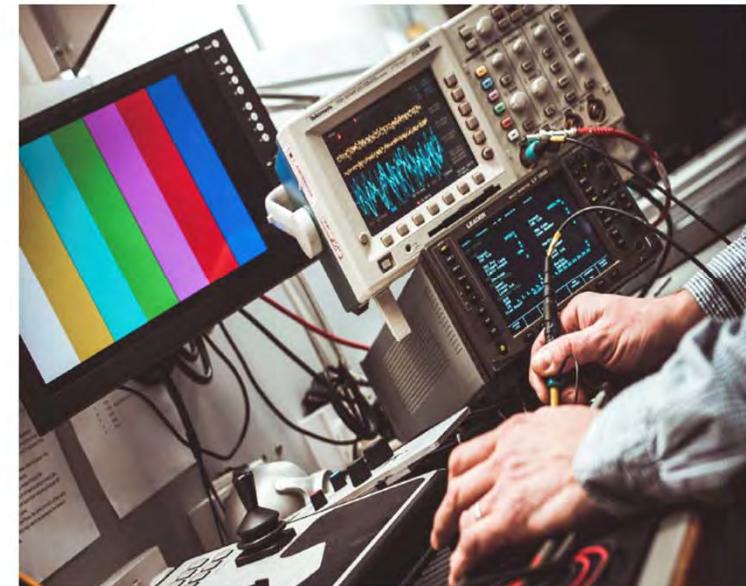
<https://youtu.be/9IXgVmLxVtQ>



AI for Music Generation

Example: “I AM AI”, first album released in 2017 to be generated by AI – with professional musicians and DL technology.

- Music generation is possible due to special DL algorithms that are designed for sequential data.
- The models learn musical patterns based on learning from large musical datasets.
- Raw music files can be processed on cloud-based computer power, making DL on these datasets possible.

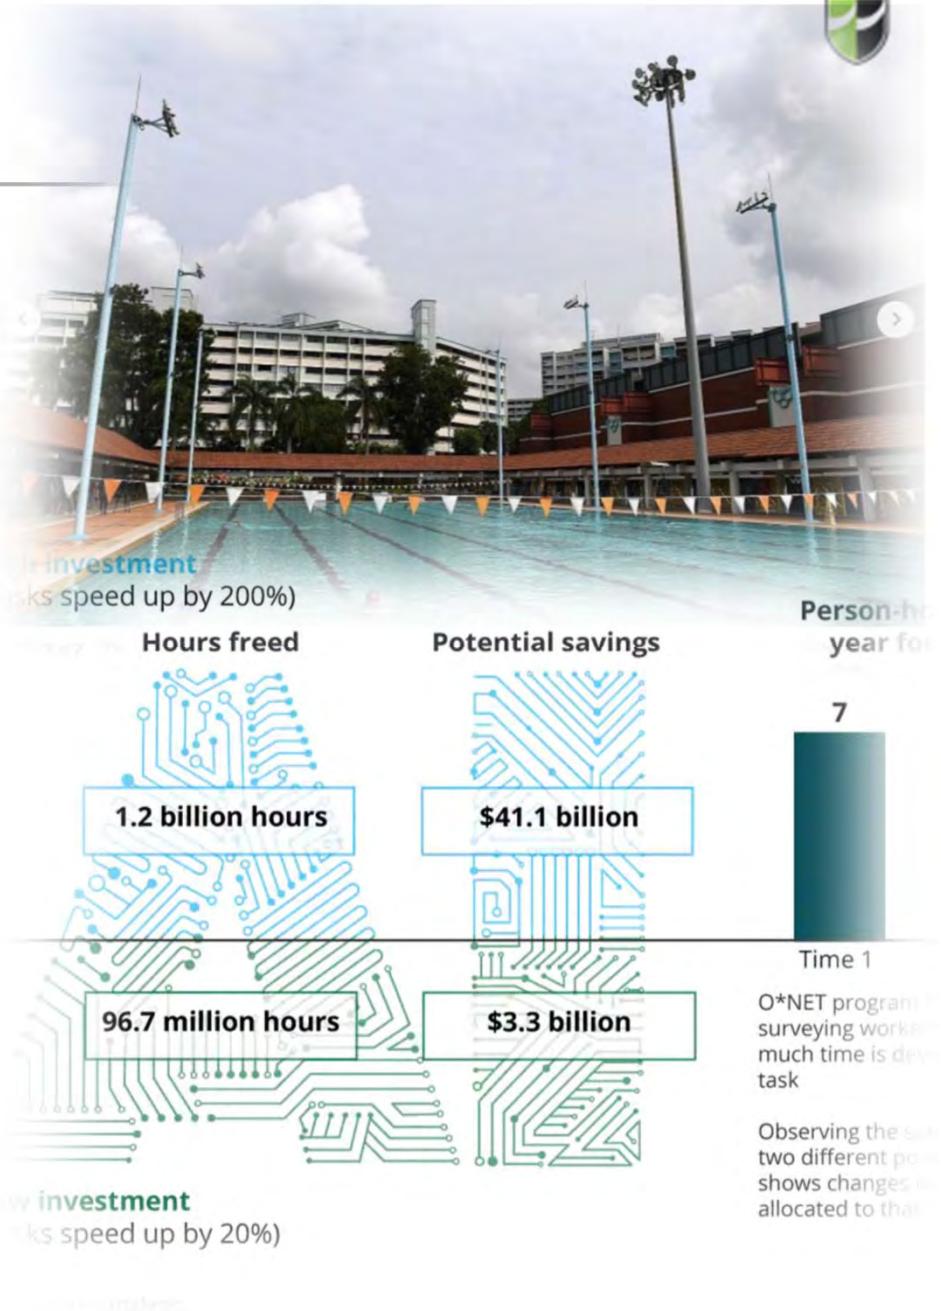


The instrumentation was composed with artificial intelligence, lyrics and vocal melodies written by Taryn. Let us know what you think in the comments!



Governments

- Public Safety and security
- Bureaucratic Efficiency





Smart Cities

Traditionally: As of 2008, for the first time in history, half of the world's population resides in cities.

- There are heightened demands on scarce resources.
- Simultaneously, a large part of existing infrastructure is underutilized or not being used efficiently.



Now with AI: AI techniques are used to analyze photo and video data to perform studies of pedestrian and traffic trends.

- Adaptive signal control: allows traffic lights to tailor their timing based on real-time data.
- With license plate recognition, and DL technology, cities can not only optimize parking but can also track criminals.



Example: AT&T reimagines smart cities

- AT&T developed a framework to help cities integrate Internet of Things (IoT) sensors with AI.
- Remotely monitor the condition of roads, bridges, buildings.
- Assist with public safety.
- Notify police if gunfire has gone off, by using sound detection.





Cybersecurity

Example: Deep Instinct

- Uses GPU-based neural network to achieve 99% detection rates for even the most advanced cyber attacks.
- DeepInstinct's DL models have the ability to detect patterns - mostly designed by humans - enabling the prediction of pending cyber attack.





Oil and Gas

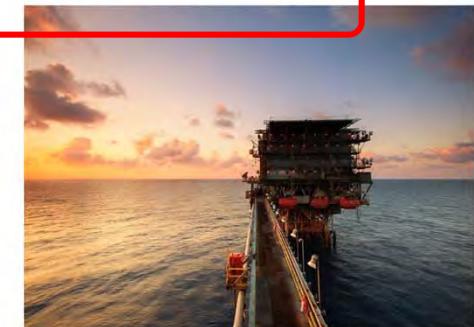
Traditionally: Shrinking oil reserves force companies to operate in remote and possibly hostile areas.

- Price has fallen dramatically in recent years.
- Forcing company layoffs and drastic budget cuts.
- Ultimately, companies are in great need of optimizing operations and cost.



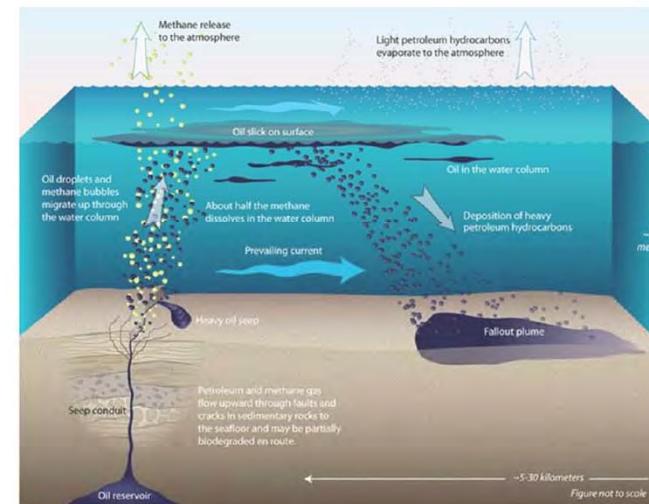
Now with AI: AI uses economic, political and weather data to forecast optimum production locations.

- Drilling is still an expensive and risk-prone endeavor.
- ML, with seismic, thermal and strata data, can help optimize the drilling process.



Example: ExxonMobile and MIT developing “submersible” robots for exploration.

- AI robots are used in ocean exploration to detect “natural seep”.
- Robots are trained via DL techniques and learn from their mistakes.
- Simultaneously protect the ecosystem and detect new energy resources.

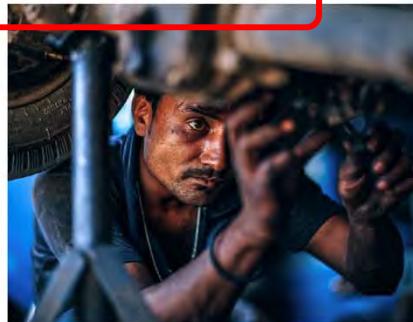




Preventive Maintenance

Traditionally: Relied on historical data to provide basis for preventative maintenance schedule.

- Conservative approach: parts were replaced well before failure, and thus financially inefficient.
- Flawed due to inability to predict new failure modes.



Now with AI: Internet of Things (IoT) sensors help to optimize maintenance scheduling.

- Part replacement schedule is optimized by assessing anomalies and failure patterns.
- Safety and productivity can increase exponentially.



Example: AI with General Electric.

- GE is the industry leader for Internet of Things (IoT) sensor installations on engines and turbines, and plans to have 60,000 engines connected to the internet by 2020.
- Computer vision cameras and reinforcement learning algorithms find tiny cracks or damage.
- Sensor data and AI allows GE to track performance and optimize part replacement.





Fault Detection

Example: Computer vision for fault detection on solar panels.

- DL algorithm trained on labelled data of correctly manufactured vs. flawed panels
- Reduced the need for human inspection by 66% compared to historical need





Automate Garment Industry

Example: SoftWear Automation's "sewbots".

- Computer vision is used to track fabric at the thread level.
- Eliminates need for human seamstress / seamster.
- Allows designers to create garments that were previously thought to be too complicated or specialized to construct.





AI and Customer Service

Example: Bot assistants and customer service agents

- AI Augmented messaging.
- AI for sorting and routing inquiries.
- AI enhanced customer phone calls.
- Some companies have used AI to fully automate customer service.





AI and Next Gen Gaming

Now with AI: Forza 5 Motorsport* uses its “Drivatar” AI system to learn how to drive in the style of other players in the game.

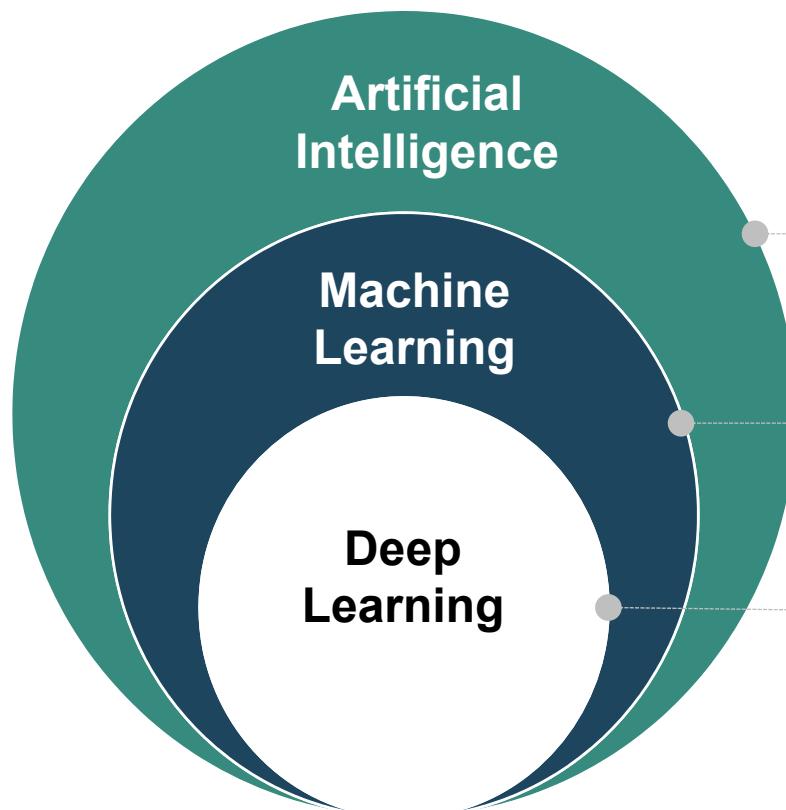
- Neural networks are used to train characters to walk and run realistically.
- Reinforcement Learning (RL) is a technique used throughout gaming.





Definitions

<http://bit.ly/2X1LxFi>



Artificial Intelligence

Any technique which enables computers to sense, reason, act and adapt



Machine Learning

Subset of AI techniques which use statistical methods to enable machines to improve with experiences.

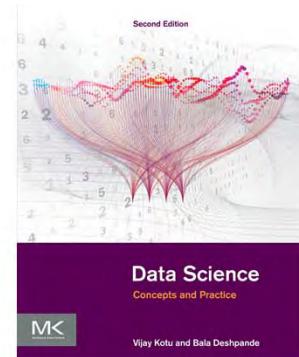
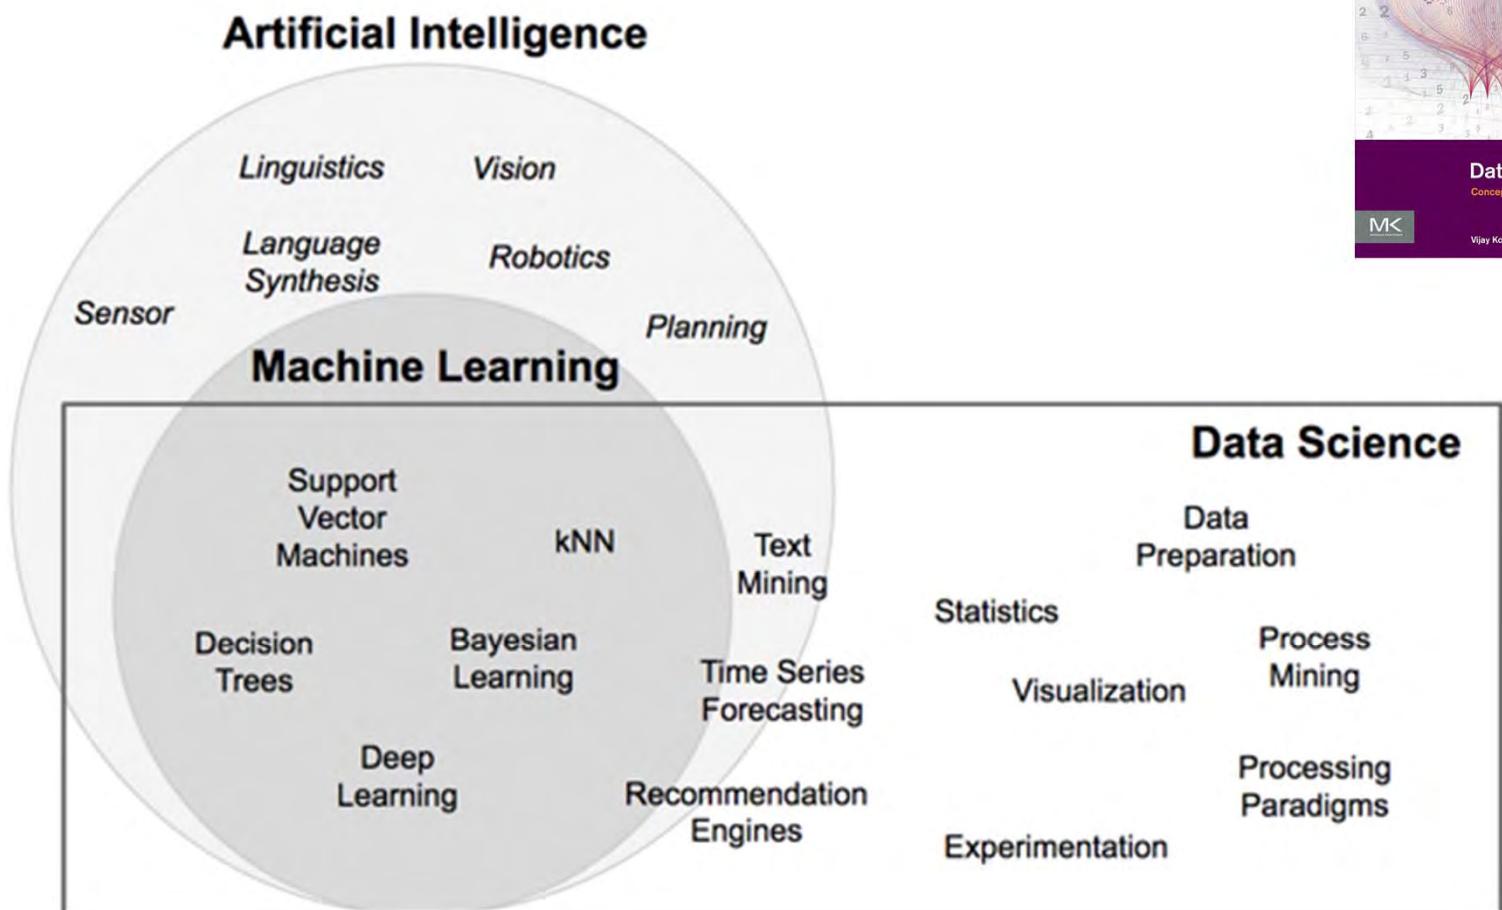


Deep Learning

A subset of machine learning in which multilayered neural networks learn from vast amount of data.



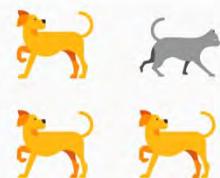
AI vs ML vs DS



5 questions data science answers



**Is this weird?
Anomaly detection**



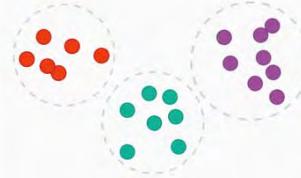
Is this pressure
gauge reading
normal?
Is this message
from the internet
typical?

**How many?
How Much?
Regression**



What will the
temperature be
next Tuesday?
What will my
fourth quarter
sales be?

**How is this
organized?
Clustering**



Which viewers
like the same
types of
movies?
Which printer
models fail the
same way?

**What should I
do?
Reinforce
Learning**



If I'm a self-
driving car: At a
yellow light,
brake or
accelerate?
For a robot
vacuum: Keep
vacuuming, or
go back to the
charging station?

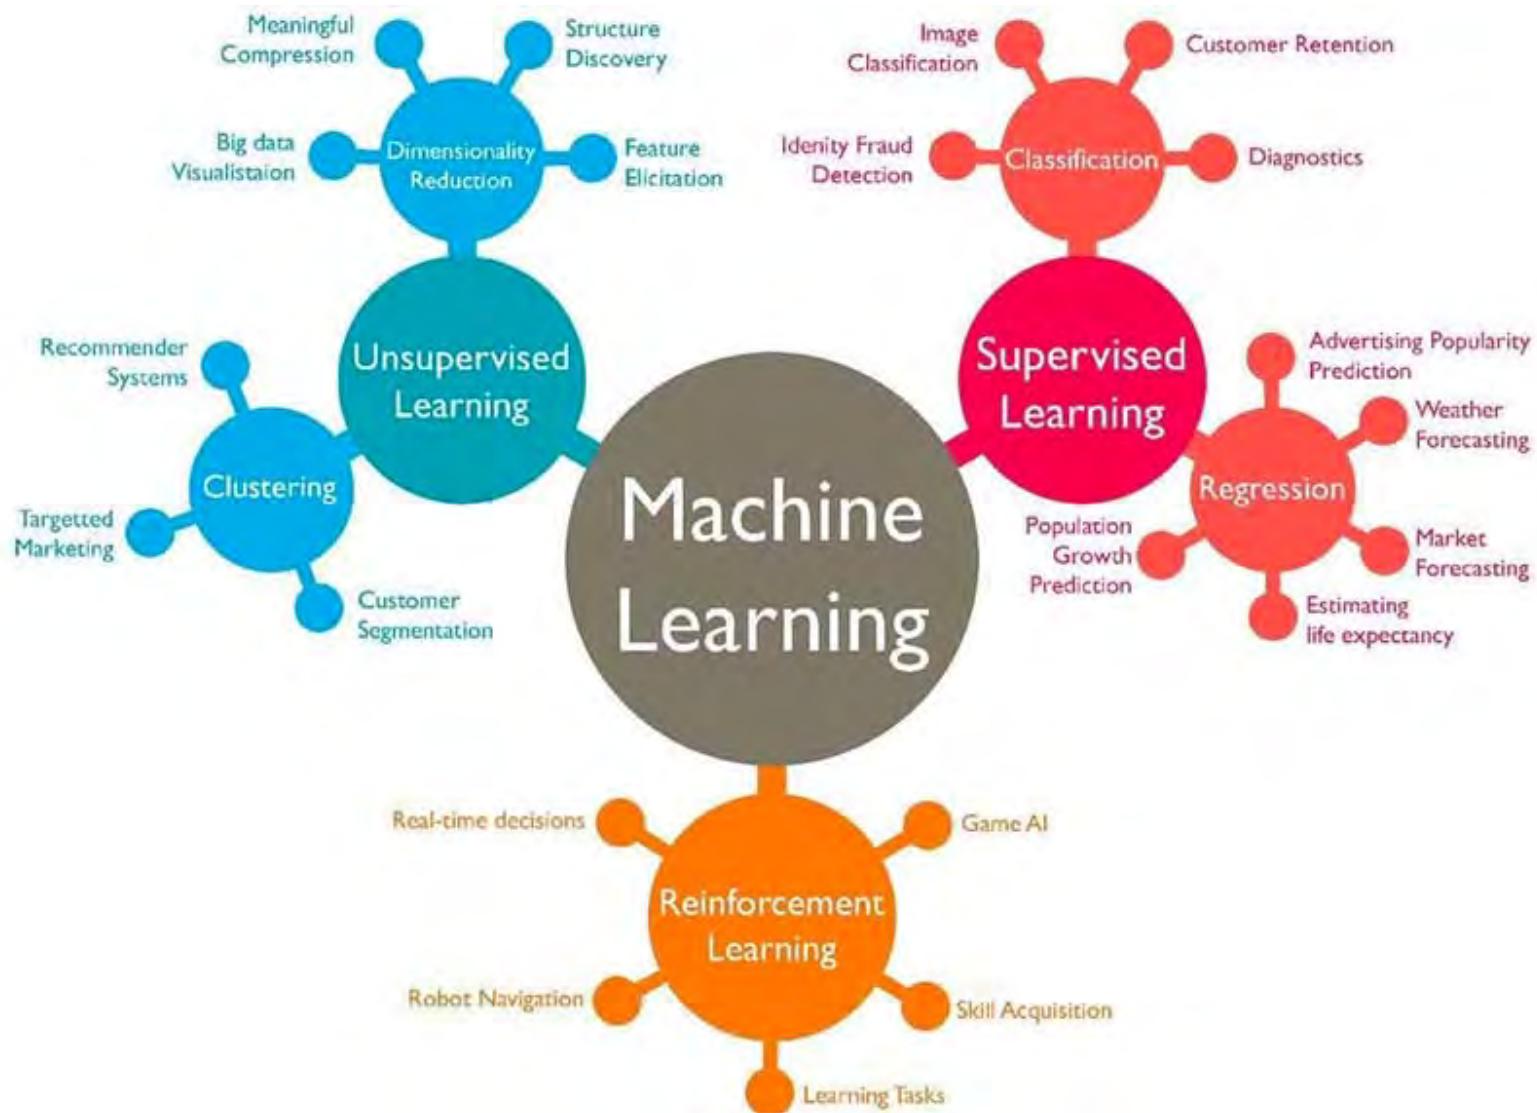
**Is this A or B?
(Classification)**



Will this tire fail
in the next 1,000
miles: Yes or
no?
Which brings in
more
customers: a \$5
coupon or a
25% discount?

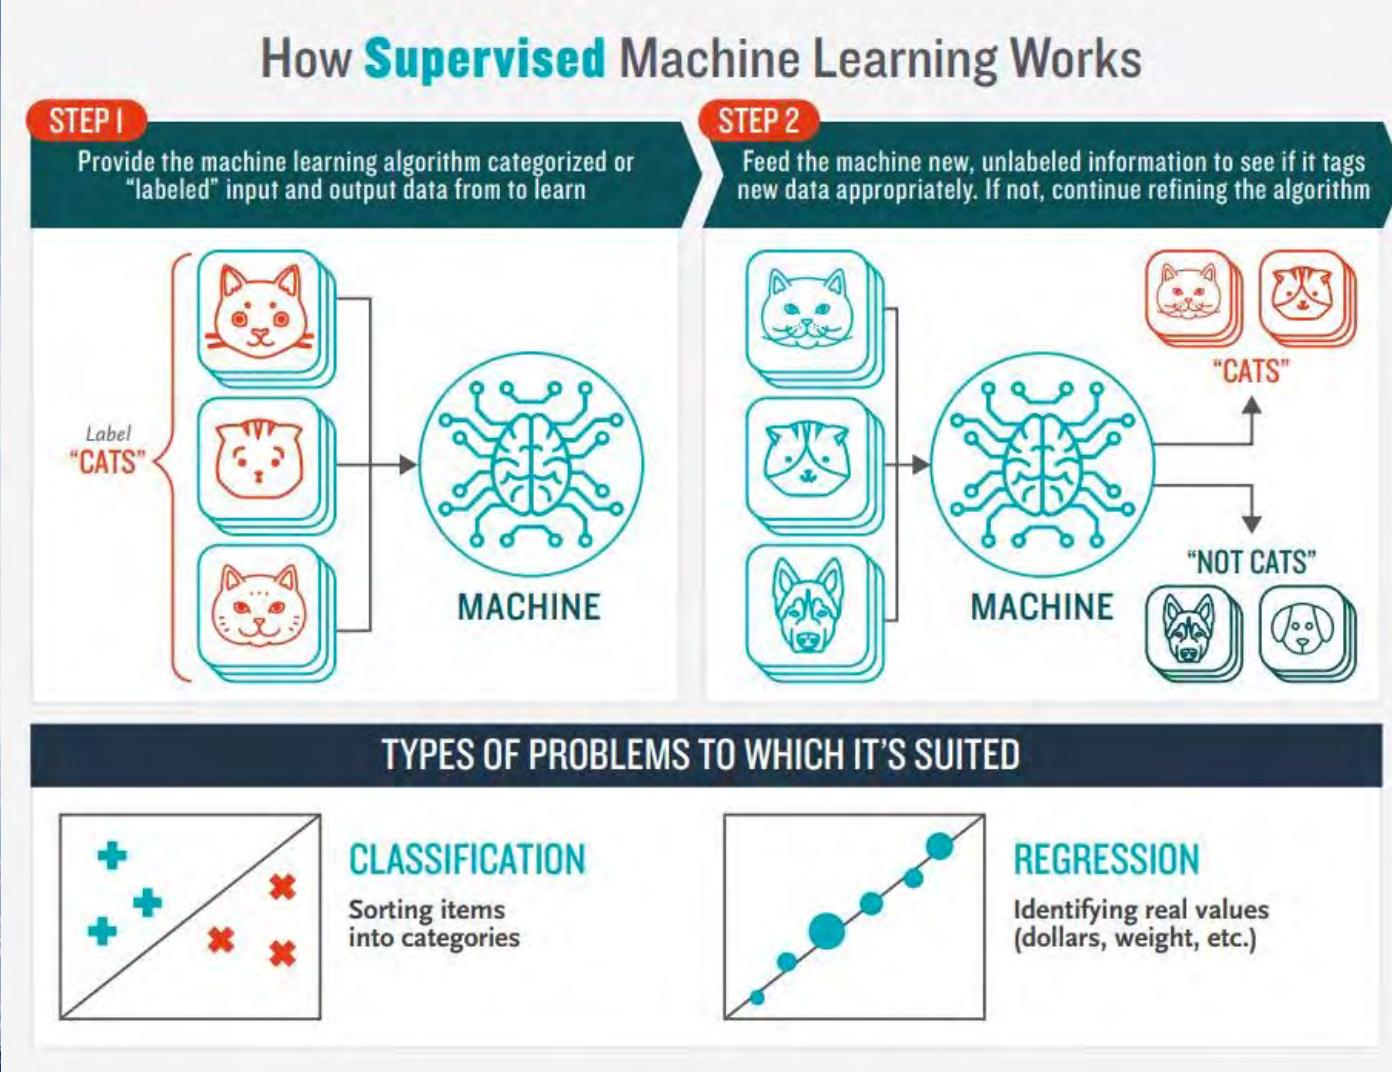


Types of Machine Learning



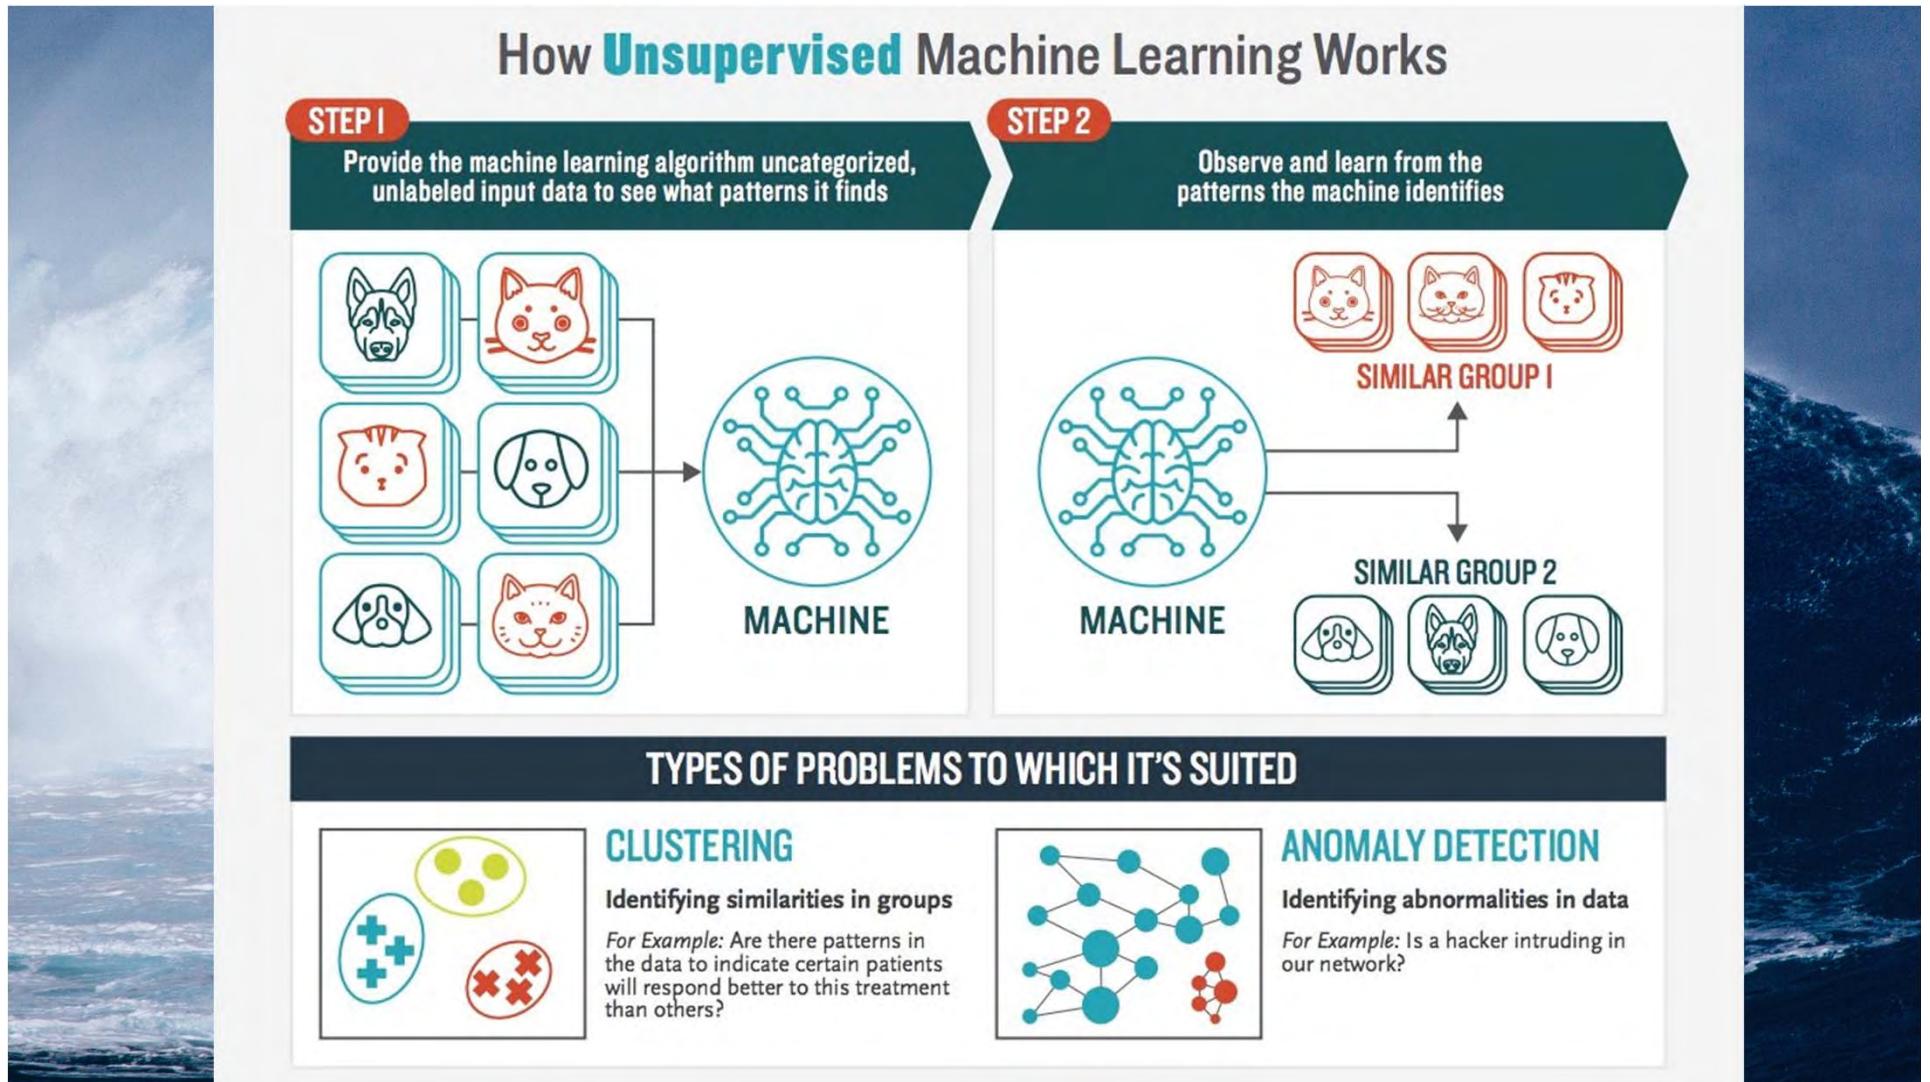


Supervised Learning





Unsupervised Learning





Machine Learning Example

- Suppose you wanted to identify fraudulent credit card transactions.
- You could define features to be:
 - Transaction time
 - Transaction amount
 - Transaction location
 - Category of purchase
- The algorithm could learn what feature combinations suggest unusual activity.





Machine Learning Limitations

- Suppose you wanted to determine if an image is of a cat or a dog.
- What features would you use?
- This is where **Deep Learning** can come in.



Dog and cat recognition

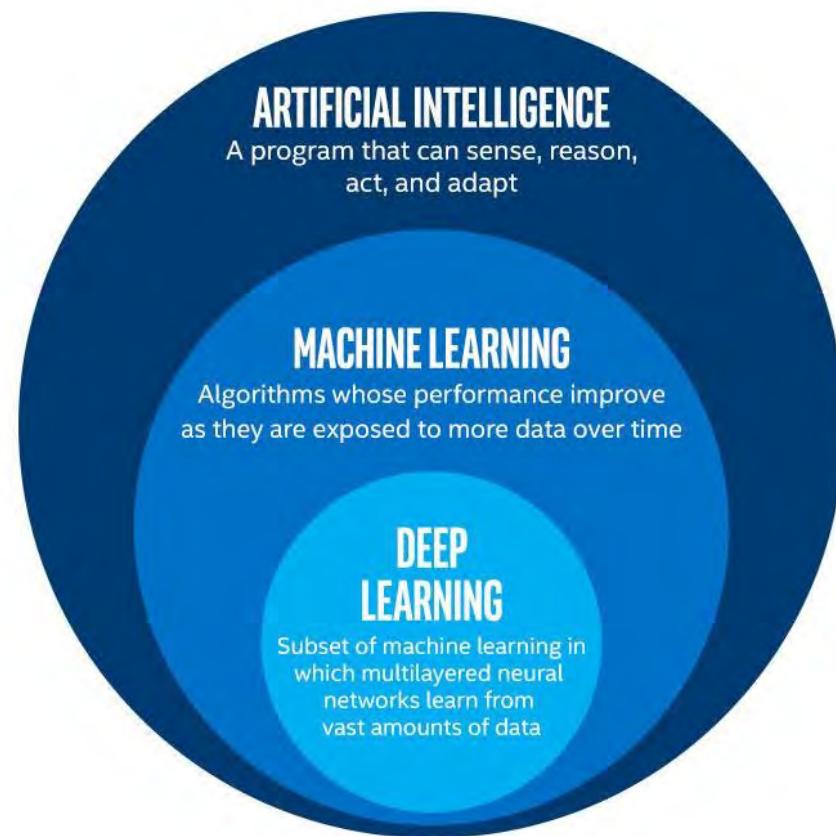


What is deep learning?

Deep Learning

“Machine learning that involves using very complicated models called “deep neural networks”.”
(Intel)

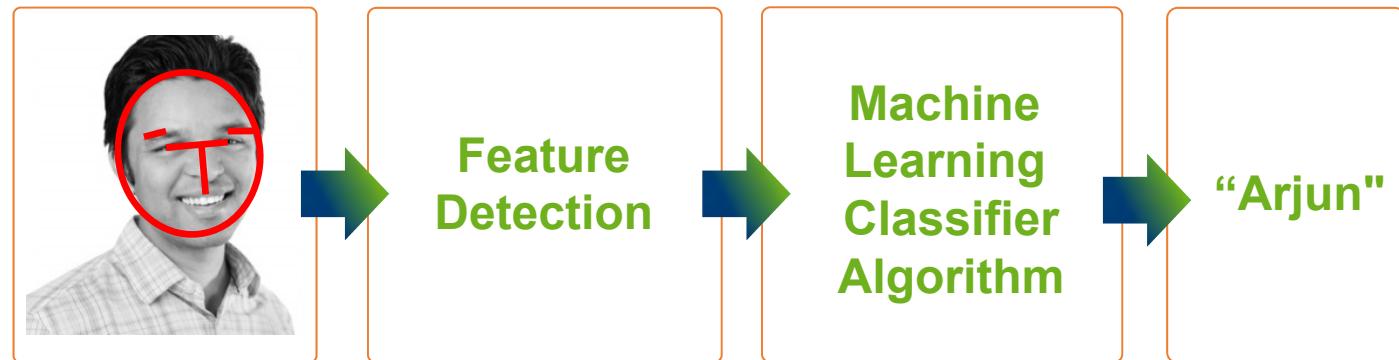
Models determine best representation of original data; in classic machine learning, humans must do this.



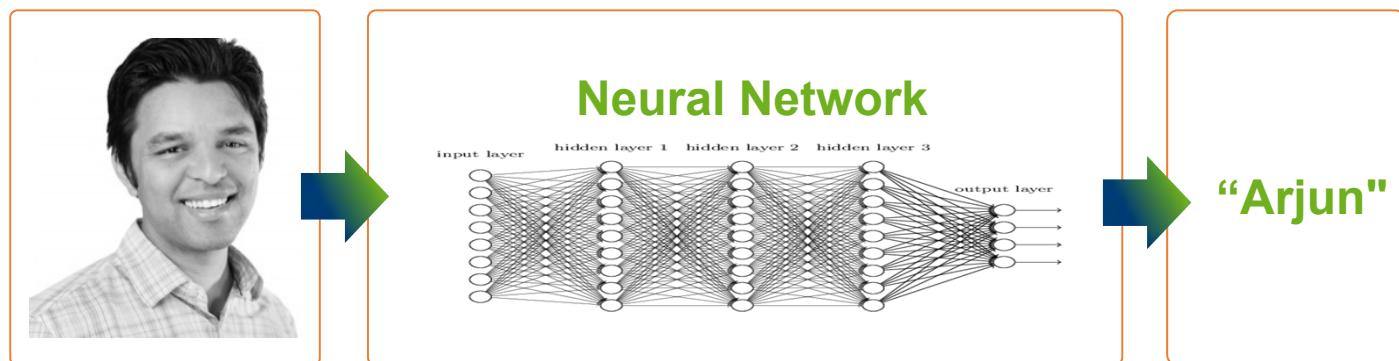


Deep Learning Example

Classic Machine Learning
Step 1: Determine features.
Step 2: Feed them through model.



Deep Learning
Steps 1 and 2 are combined into 1 step.





Deep Learning Problem Types

Deep Learning can solve multiple supervised and unsupervised problems.

- The majority of its success has been when working with images, natural language, and audio data.
- Image classification and detection.
- Semantic segmentation.
- Natural language object retrieval.
- Speech recognition and language translation.

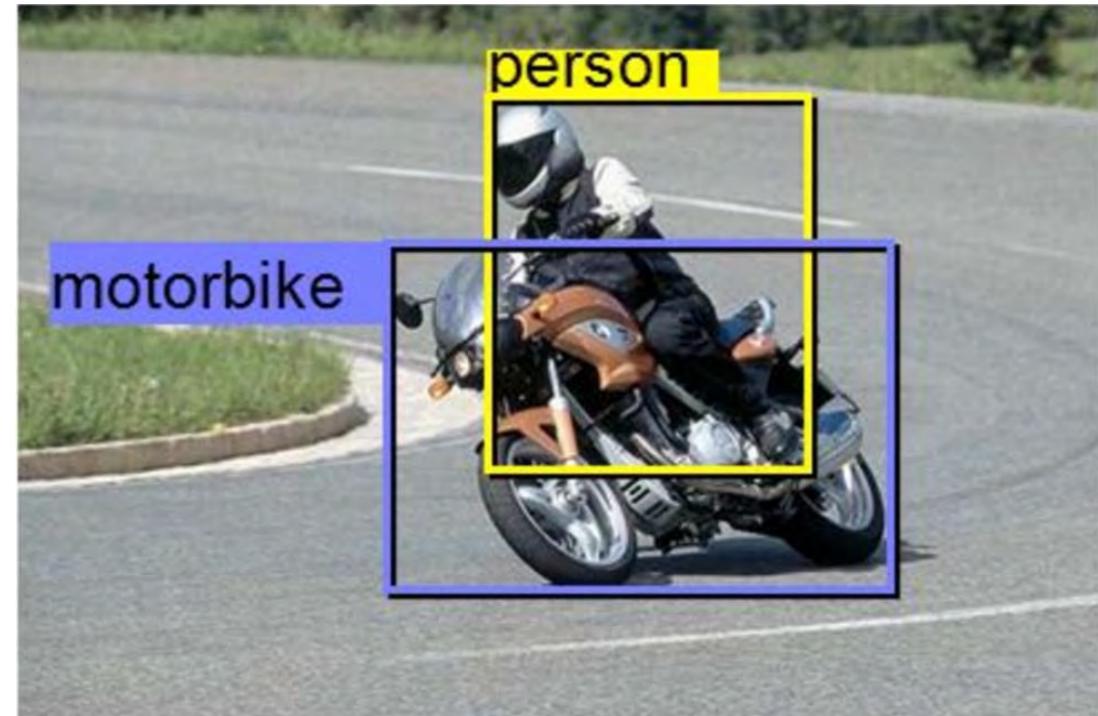




Classification and Detection

Detect and label the image

- Person
- Motor Bike



<https://people.eecs.berkeley.edu/~jhoffman/talks/llda-baylearn2014.pdf>





Semantic Segmentation

Label every pixel



<https://people.eecs.berkeley.edu/~jhoffman/talks/lsda-baylearn2014.pdf>





Natural Language Object Retrieval

a scene with three people



query='man far right'



query='left guy'



query='cyclist'



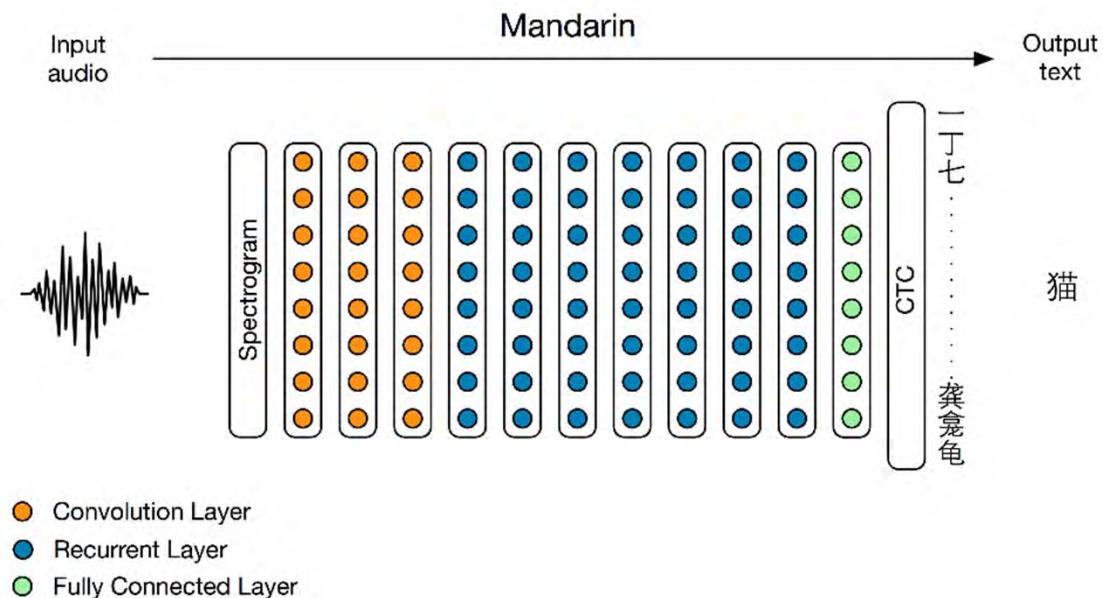
<http://arxiv.org/pdf/1511.04164v3.pdf>





Speech Recognition and Language Translation

The same architecture can be used for speech recognition in English, or in Mandarin Chinese.



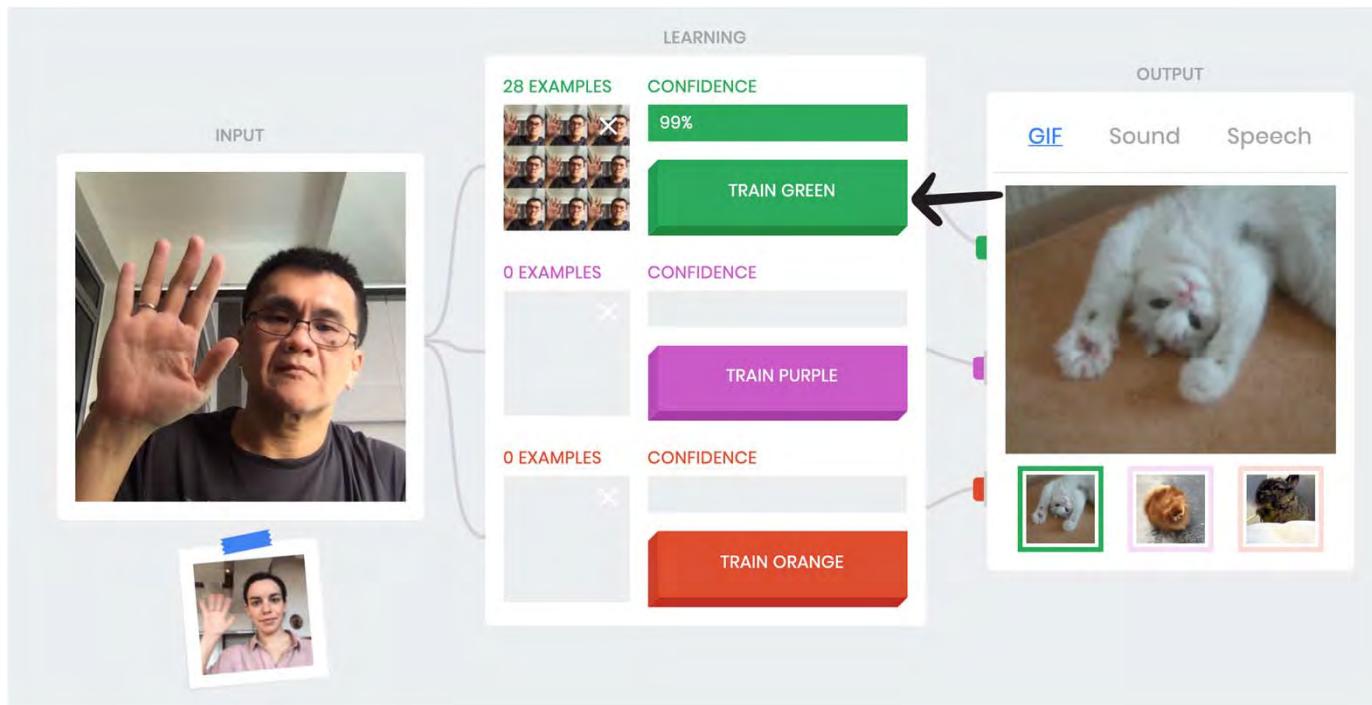
<http://svail.github.io/mandarin/>





Deep Learning in Action

<https://teachablemachine.withgoogle.com/>





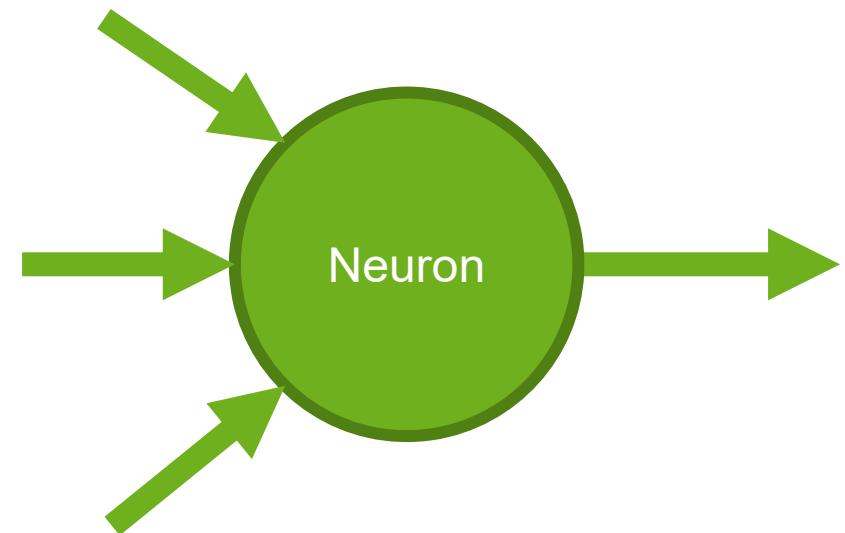
What is neural network?



Which model?

There are many models that represent the problem and make decisions in different ways, each with their own advantages and disadvantages.

- DL models are biologically inspired.
- The main building block is a **neuron**.

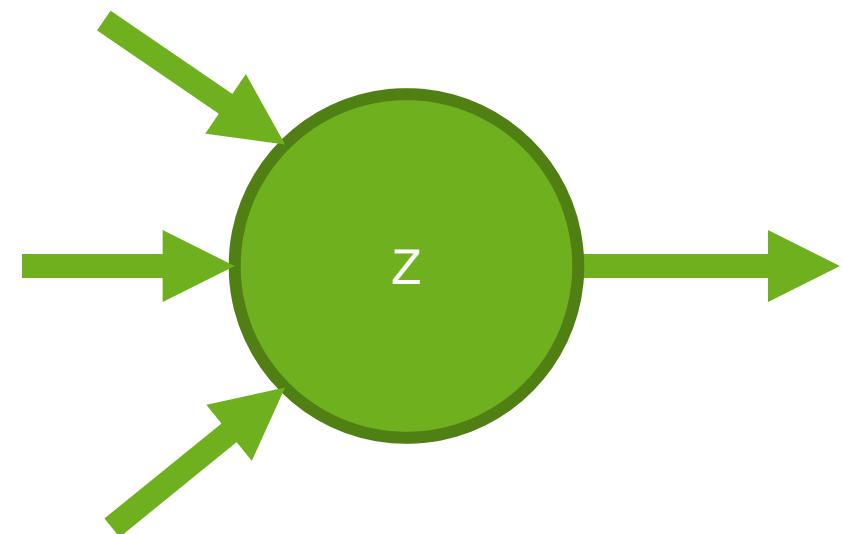




Neuron

A neuron multiplies each feature by a **weight** and then adds these values together.

- $Z = X_1W_1 + X_2W_2 + X_3W_3$

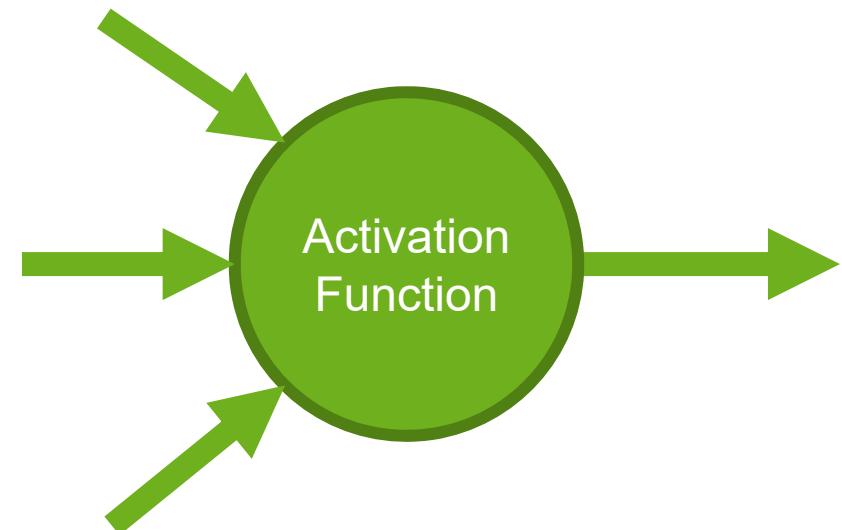




Neuron

This value is then put through a function called the **activation function**.

- There are several activation functions that can be used.
- The output of the neuron is the output of the activation function.

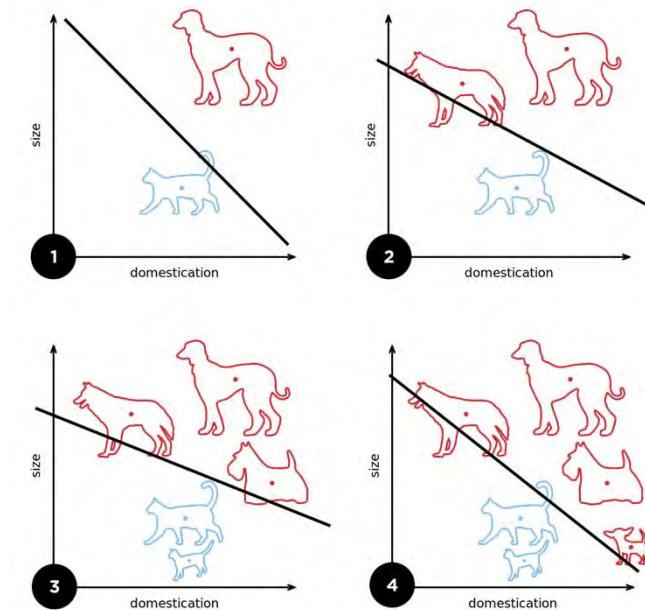




Perceptron

A neuron with a simple activation function can solve **linearly separable** problems.

- These are problems where the different classes can be separated by a line.
- **The Perceptron:** one of the earliest neural network models that used neurons with simple activation functions.

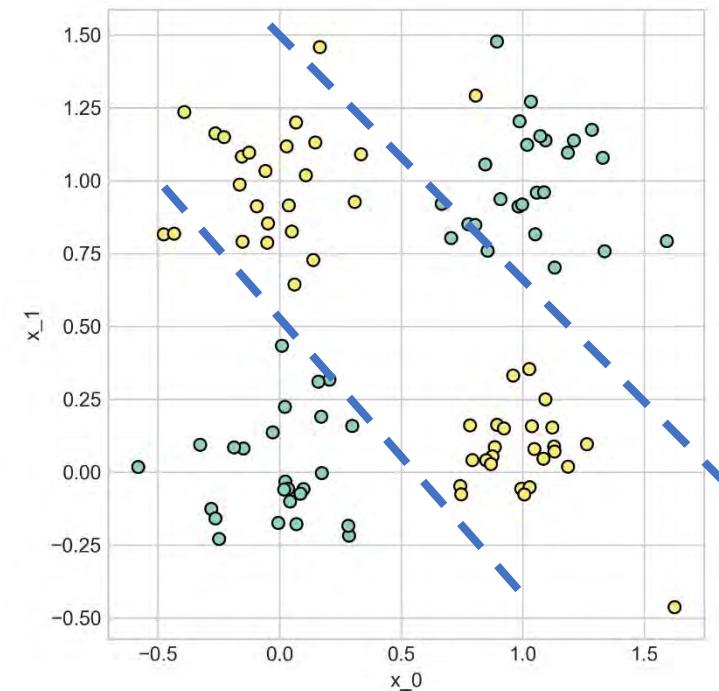




Perceptron

Problems where the labels cannot be separated by a single line are not solvable by a single neuron.

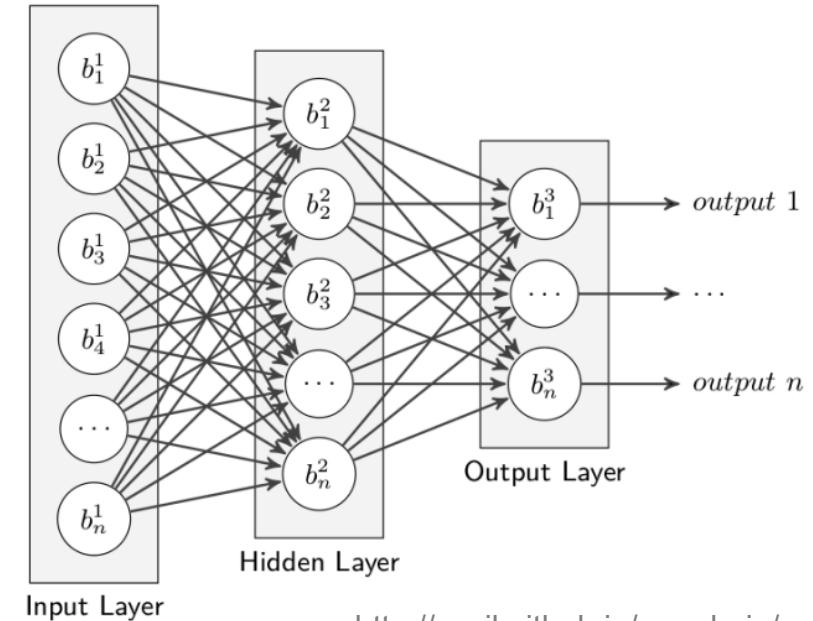
- This is a major limitation, and one of the reasons for the first AI winter, that was discussed earlier.





Fully Connected Network

- More complicated problems can be solved by connecting multiple neurons together and using more complicated activation functions.
- Organized into **layers** of neurons.
- Each neuron is connected to every neuron in the previous layer.
- Each layer transforms the output of the previous layer and then passes it on to the next.
- Every connection has a separate weight.



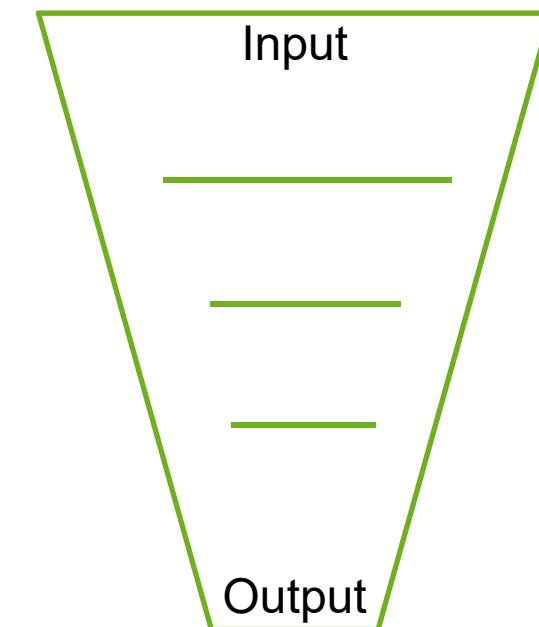
<http://svail.github.io/mandarin/>





Deep Learning

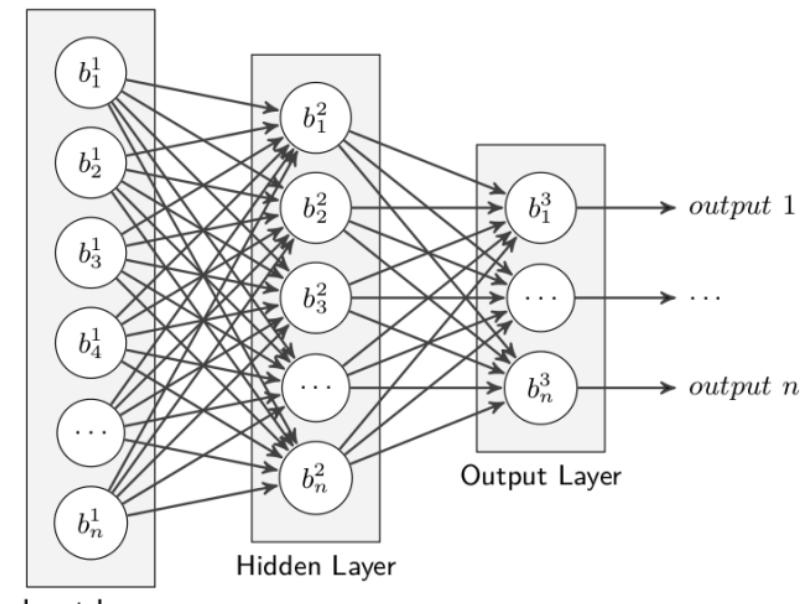
- Deep Learning refers to when many layers are used to build **deep networks**.
- State-of-the-art models use hundreds of layers.
- Deep layers tend to decrease in width.
- Successive layers transform inputs with two effects:
- **Compression**: each layer is asked to summarize the input in a way that best serves the task.
- **Extraction**: the model succeeds when each layer extracts task-relevant information.





Steps in Building a Fully Connected Network

- To build a fully connected network a user needs to:
- Define the network architecture.
- How many layers and neurons?
- Define what activation function to use for each neuron.
- Define an evaluation metric.
- The values for the weights are learned during model training.



<http://svail.github.io/mandarin/>





Evaluation Metric

The metric used will depend on the problem being solved. Some examples include:

- Regression
 - Mean Squared Error
- Classification
 - Categorical Cross-Entropy
- Multi-Label classification
 - Binary Cross-Entropy

Fully Connected Network Problems



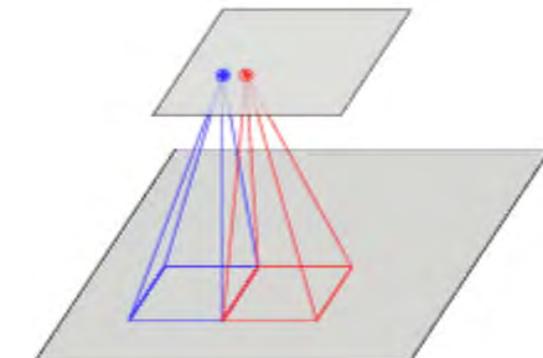
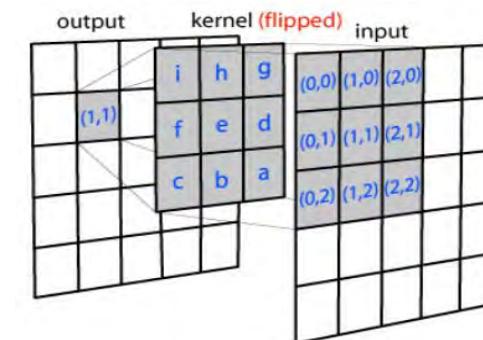
Not optimal for detecting features.

- Computationally intensive – heavy memory usage.



Convolutional Neural Network

- **Convolutional neural networks** reduce the required computation and are good for detecting features.
- Each neuron is connected to a small set of nearby neurons in the previous layer.
- The same set of weights are used for each neuron.
- Ideal for spatial feature recognition.
 - Example: image recognition
- Cheaper on resources due to fewer connections.



<http://svail.github.io/mandarin/>





Convolutions as Feature Detectors

Convolutions can be thought of as “local feature detectors”.

Vertical Line Detector

-1	1	-1
-1	1	-1
-1	1	-1

Horizontal Line Detector

-1	-1	-1
1	1	1
-1	-1	-1

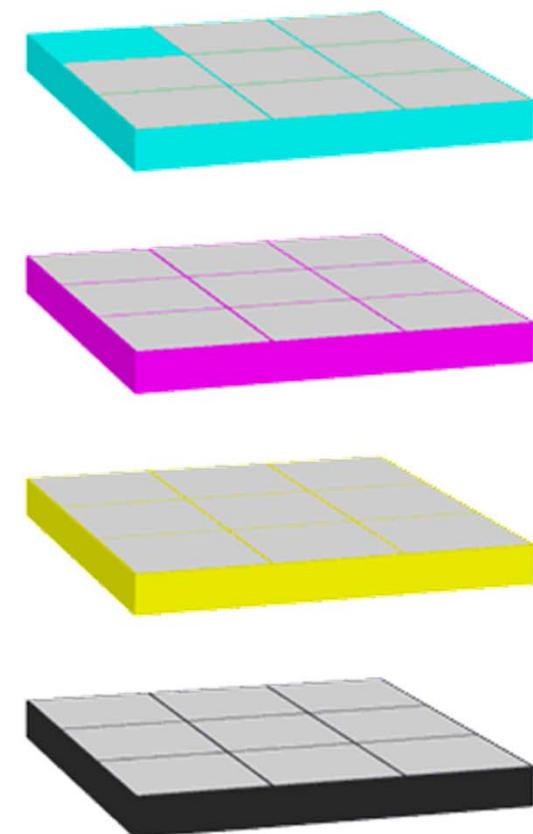
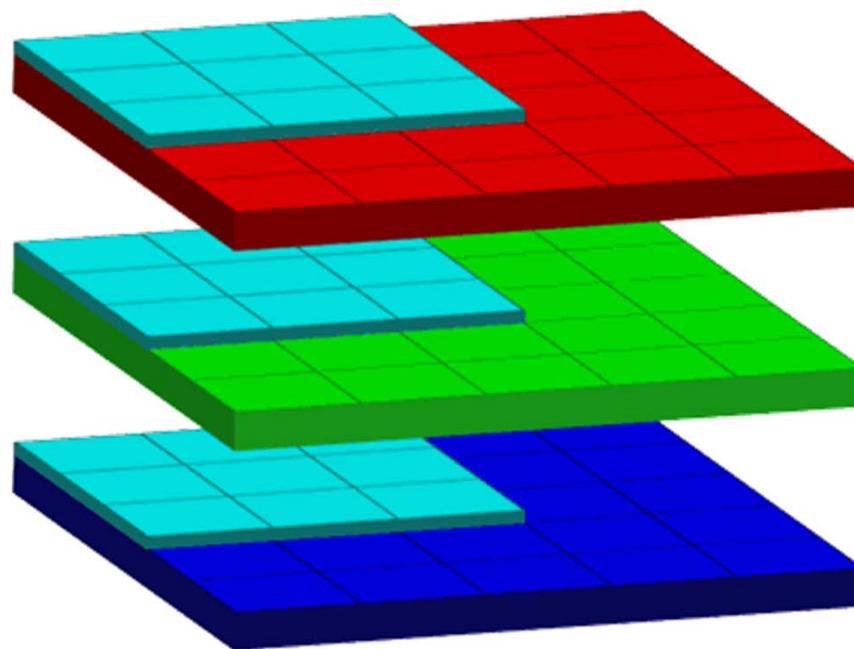
Corner Detector

-1	-1	-1
-1	1	1
-1	1	1



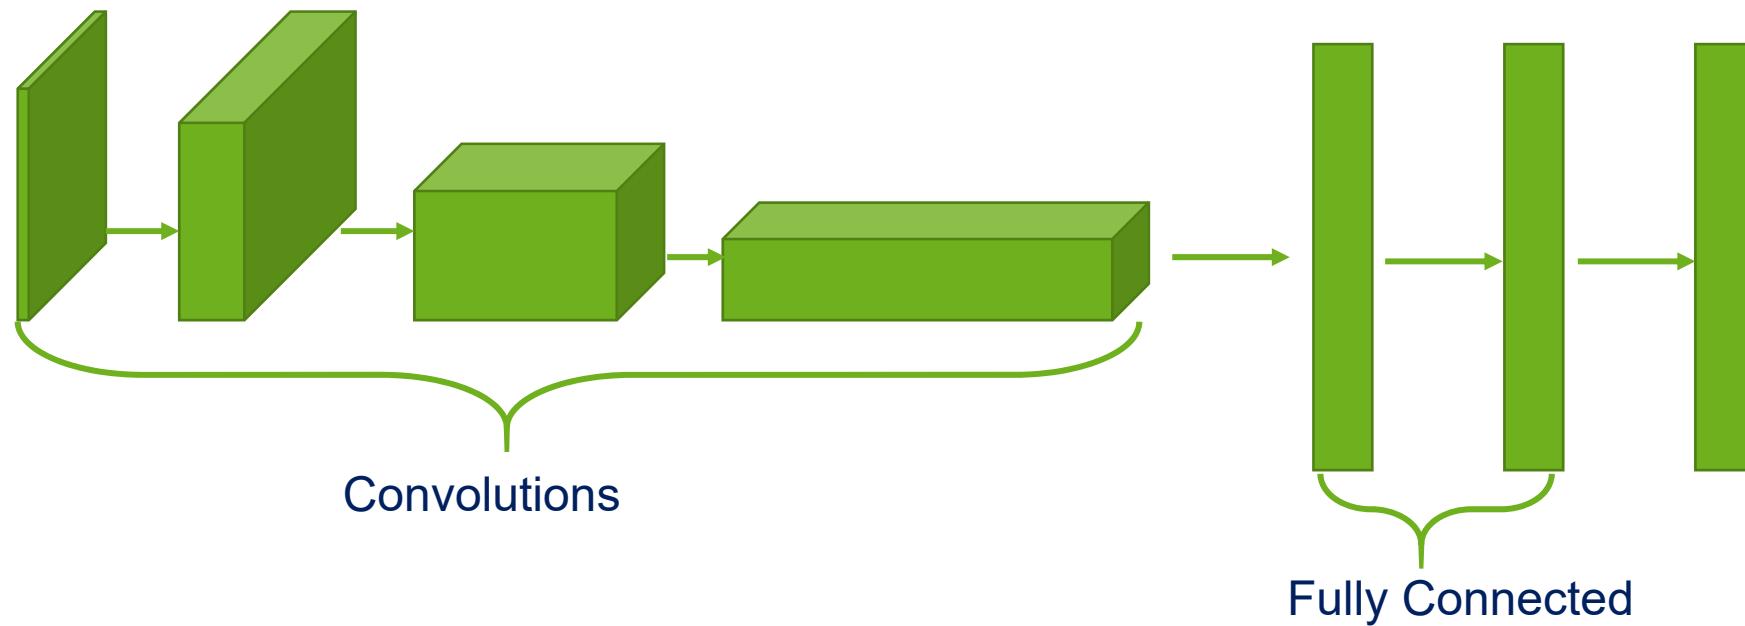


Convolutions





Convolutional Neural Network





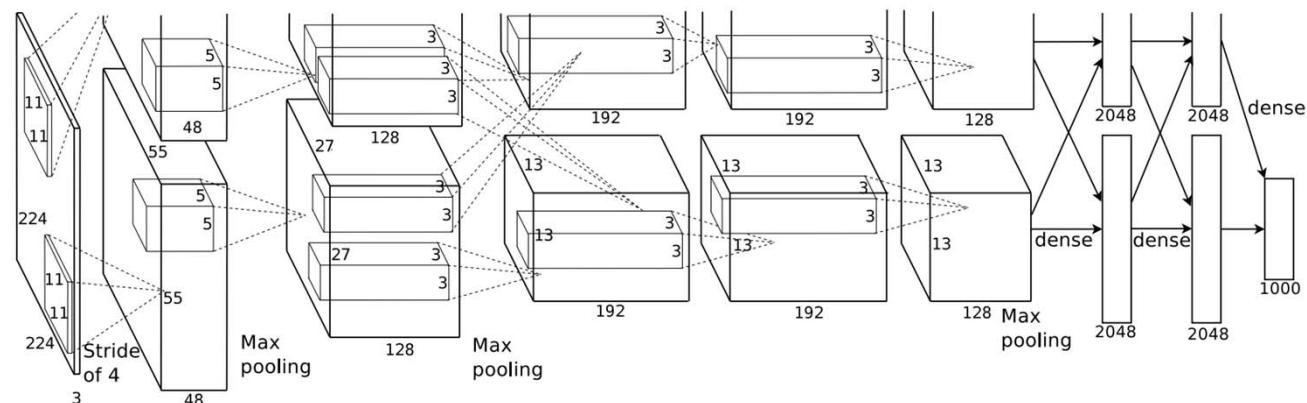
Common Architectures

- AlexNet
- VGG
- Inception
- MobileNets



AlexNet

- Created in 2012 for the ImageNet Large Scale Visual Recognition Challenge (ILSVRC)
- Task: predict the correct label from among 1000 classes
- Dataset: around 1.2 million images
- Considered the “flash point” for modern deep learning
- Demolished the competition.
- Top 5 error rate of 15.4%
- Next best: 26.2%





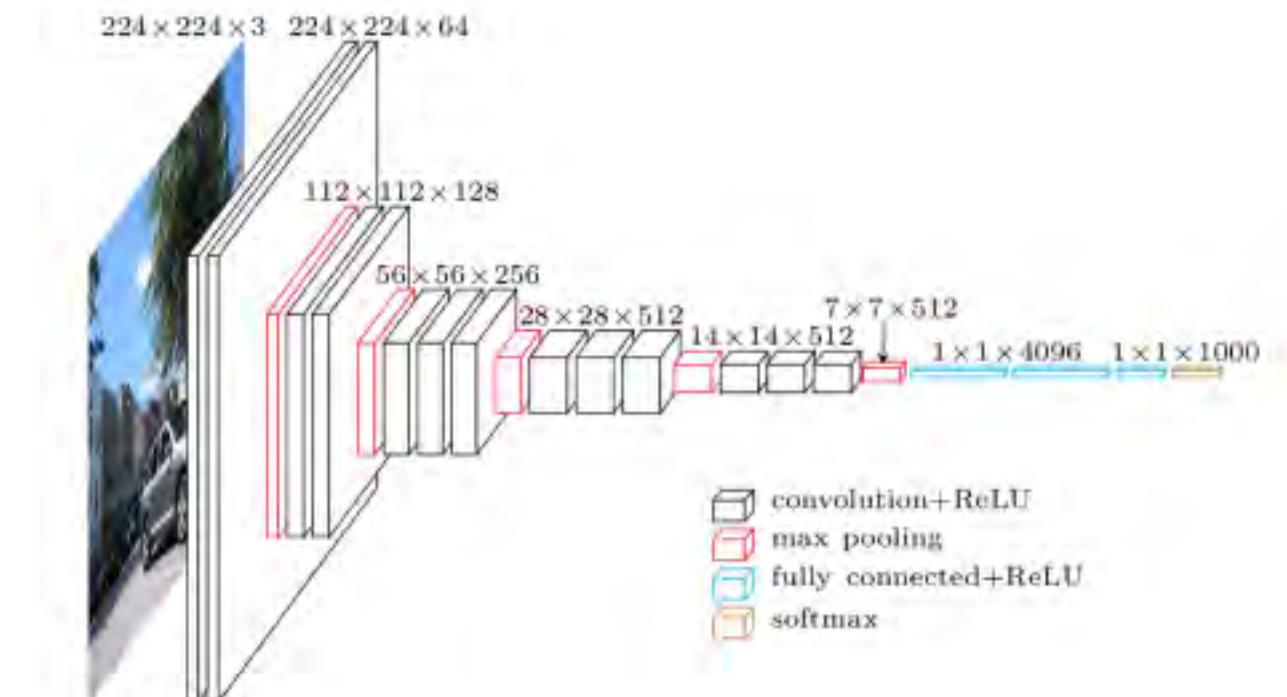
VGG

- Simplify Network Structure
- Avoid Manual Choices of Convolution Size
- Very Deep Network with 3x3 Convolutions
- These “effectively” give rise to larger convolutions

Reference:

Very Deep Convolutional Networks for Large-Scale Image Recognition

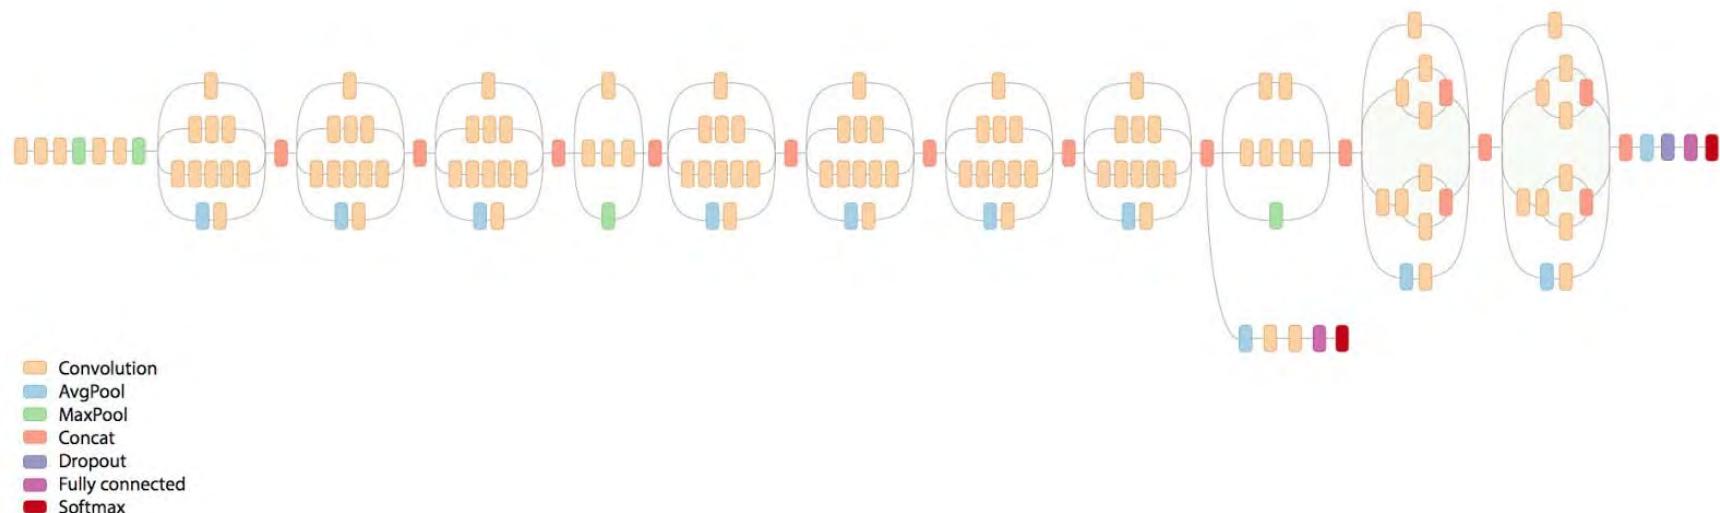
Karen Simonyan and Andrew Zisserman, 2014





Inception

- Szegedy et al 2014
- Idea: network would want to use different receptive fields
- Solution: Turn each layer into branches of convolutions
- Each branch handles smaller portion of workload
- Concatenate different branches at the end





MobileNets

- Efficient models for mobile and embedded vision applications.



MobileNet models can be applied to various recognition tasks for efficient device intelligence.





Recurrent Neural Network

Issue: Variable length sequences of words

- With images, we forced them into a specific input dimension
- Not obvious how to do this with text.
- For example, classify tweets as positive, negative, or neutral.
- Tweets can have a variable number of words.
- What to do?





Recurrent Neural Network

Issue: Ordering of words is important

- Want to do better than “bag of words” implementations
- Ideally, each word is processed or understood in the appropriate context.
- Need to have some notion of “context”.
- Words should be handled differently depending on “context”
- Also, each word should update the context.





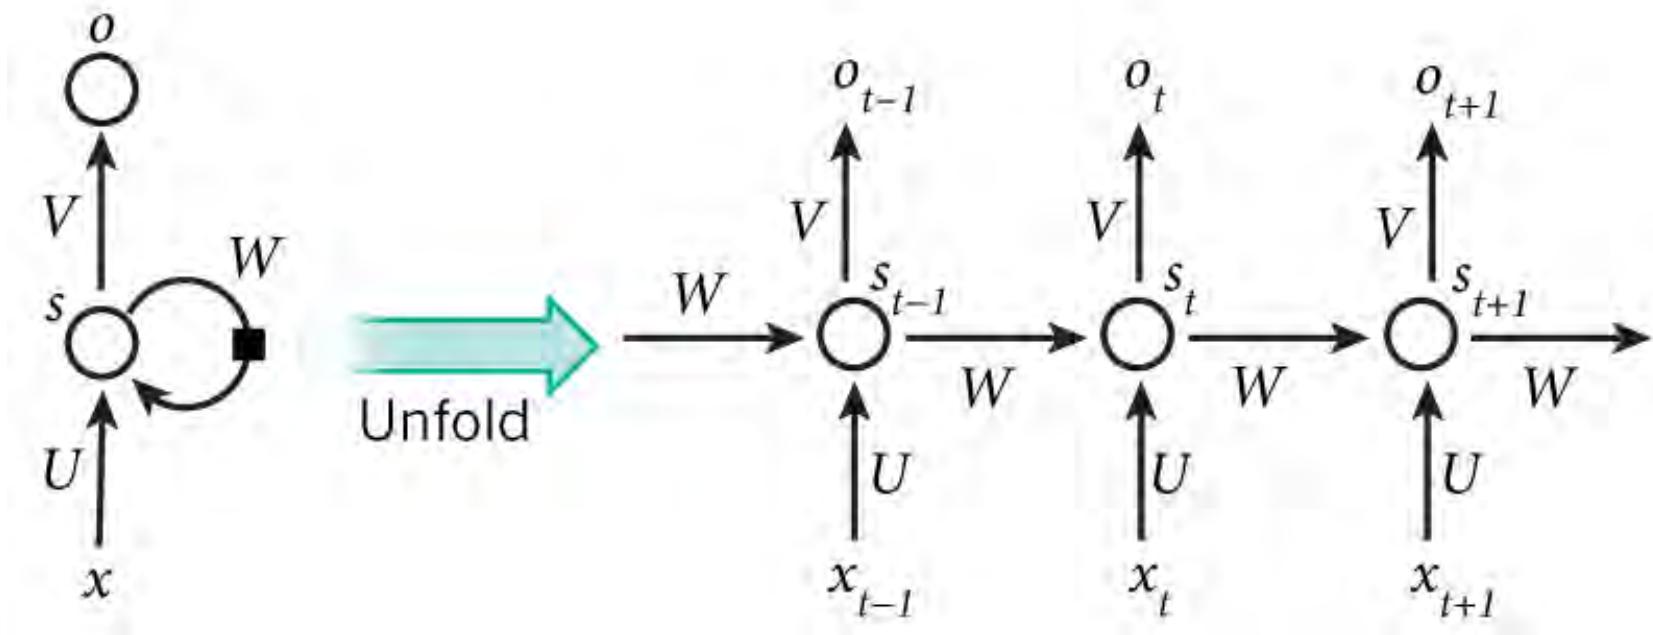
Recurrent Neural Network

Idea: Use the notion of “recurrence”

- Input words one by one
- Network outputs two things:
 - Prediction: What would be the prediction if the sequence ended with that word
 - State: Summary of everything that happened in the past
- This way, can handle variable lengths of text
- The response to a word depends on the words that preceded it.



Recurrent Neural Network



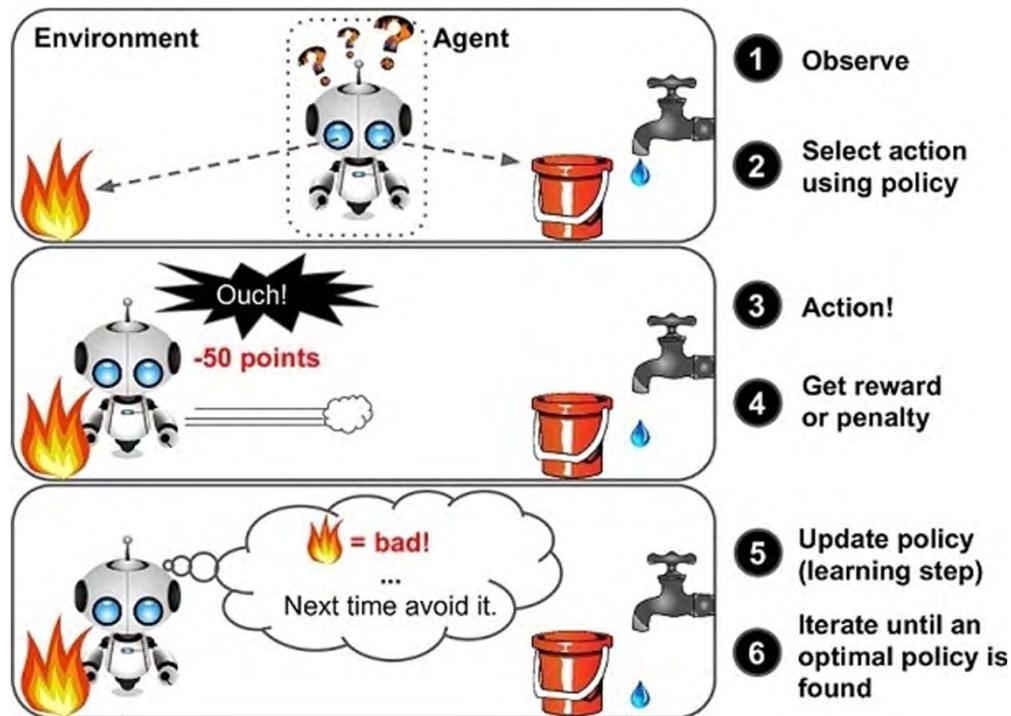


Recurrent Neural Network

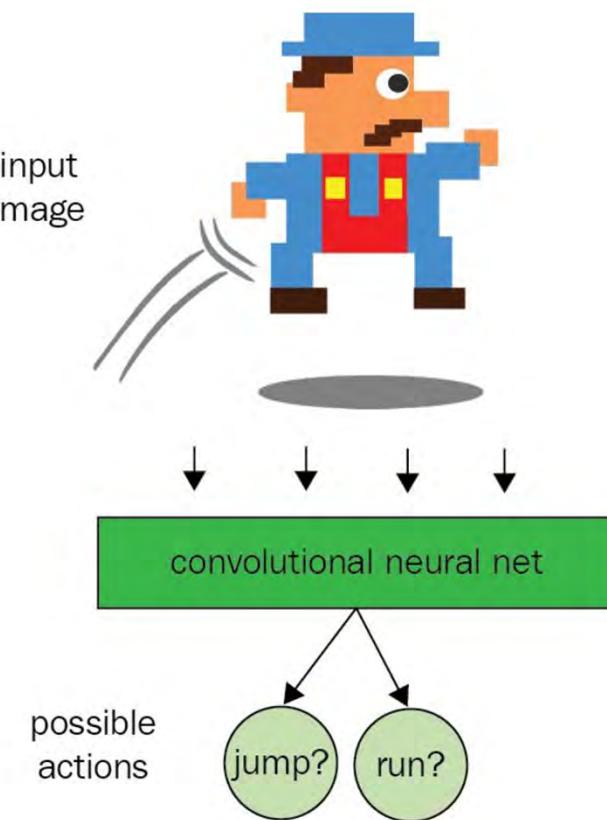
- For text/words as application.
- But, RNNs can be used for other sequential data
 - Time-Series Data
 - Speech Recognition
 - Sensor Data
 - Genome Sequences
- Nature of state transition means it is hard to keep information from distant past in current memory without reinforcement.
- **In the next lecture, we will introduce LSTMs, which have a more complex mechanism for updated the state.**



Reinforcement Learning



Convolutional Agent

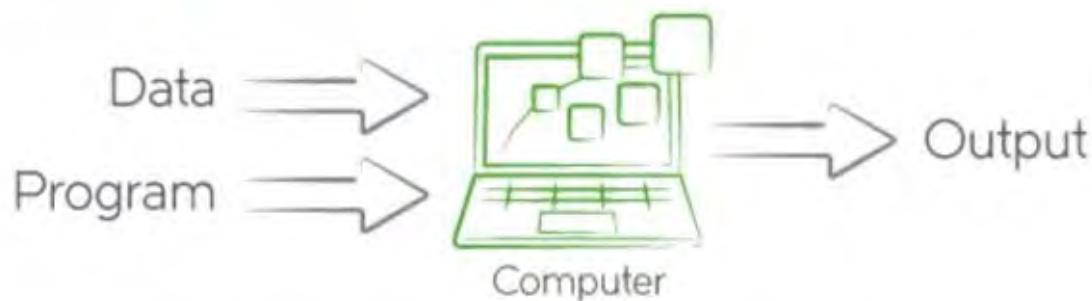


Source: Practical Reinforcement Learning

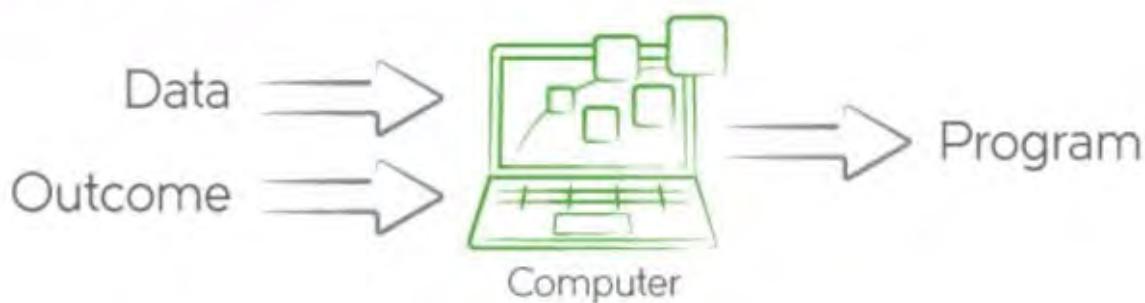


Traditional Programming vs Machine Learning

Traditional Programming

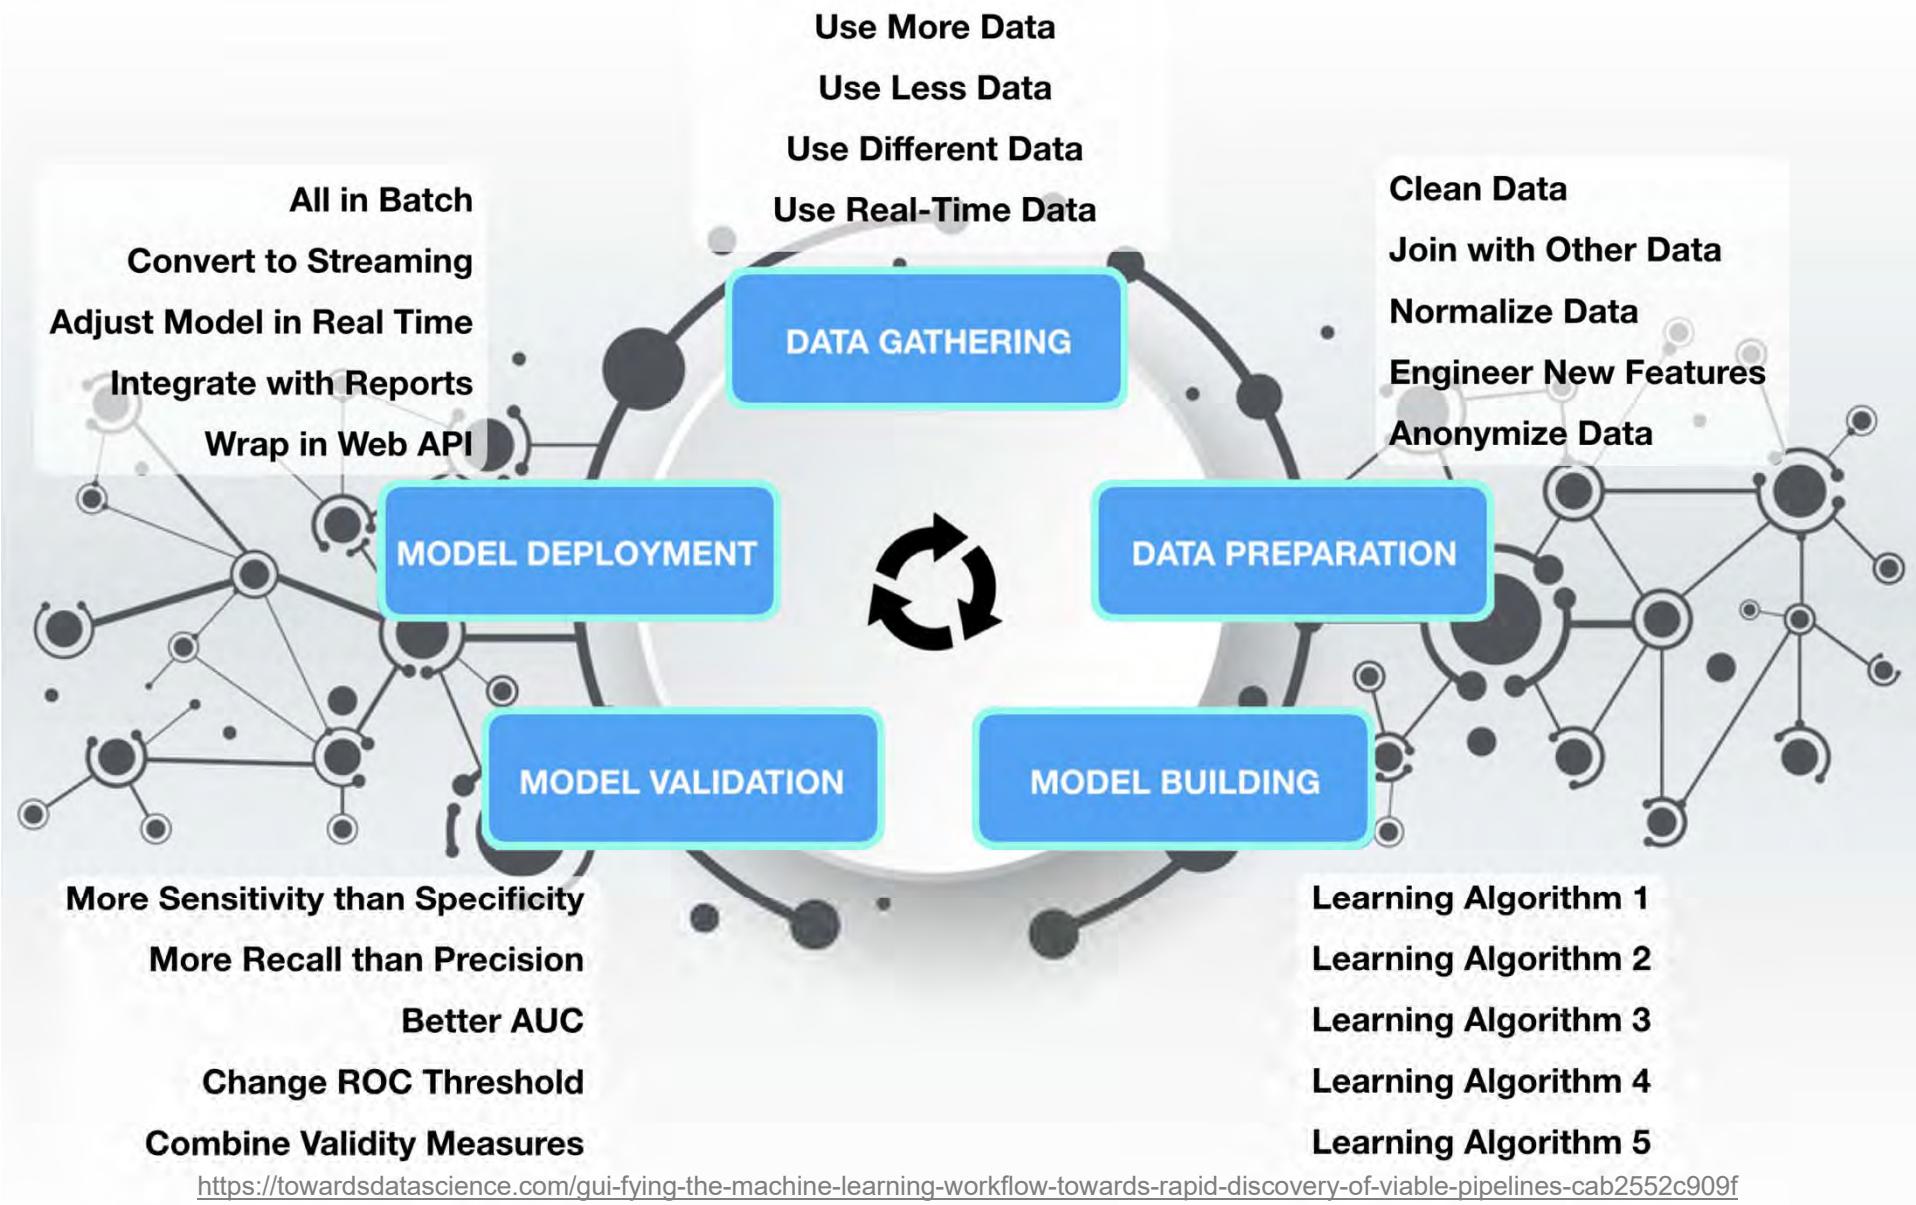


Machine Learning





AI/ML Workflow





AI Technologies Hands On

Meet LUIS

<https://aidemos.microsoft.com/luis>

Otherwise known as Language Understanding Intelligent Service.

LUIS is a service that interprets normal human language and understands what the desired outcomes and relevant pieces of information are. It's fast and easy to deploy, it learns and adapts as it interacts, and with very little training it can gain insights and trigger actions based on real user-inputs and contextual information.



Text Analytics

- Step #1: Enter Your Message

[Next Step >](#)

- Step #2: Sentiment & Key Phrases

This API analyzes your text to identify the keywords and discern the sentiment.

[Next Step >](#)

- Step #3: Entity Linking

This API ascertains which of the key words are entities and links them in Wikipedia.

[Next Step >](#)



- Step #4: Bing Entity Search

This API provides a summary of relevant information in the form of a card for each entity

[Next Step >](#)

[Learn to code >](#)

<https://aidemos.microsoft.com/text-analytics>

Computer Vision

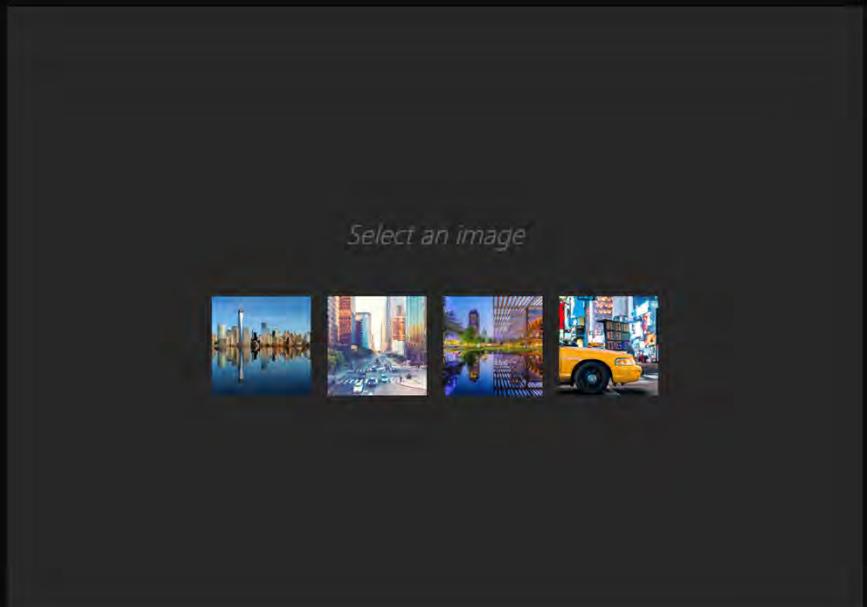
Microsoft AI - for a richer experience across a variety of visual mediums.

- **Analyze and describe images**

This feature will identify and tag the content of an image, give a written description, and give you confidence ratings on the results. It also identifies racy or adult content allowing easy moderation.

[Next Feature](#)

- Read text in imagery
- Read handwriting in imagery
- Recognize celebrities & landmarks



Face and Emotion Recognition

Facial recognition can be a great security tool. See how Microsoft AI is able to quickly analyze and compare photos. Upload or take two photos and we will use our Face API to tell you if the photos are of the same person. We'll also provide the approximate age, gender, and emotion. Really test the system by making different faces or getting a friend involved to take one of the photos.

[Add Photo](#)[Add Photo](#)

PHOTO 1

PHOTO 2

Speech Authentication

Voices are not always easy to decipher, but a critical piece of our identity. Test out our speech recognition capabilities, built using our Speech API, by recording a phrase three times or using our pre-recorded voices by selecting one of the silhouettes.

Select Phrase 

Select Recording



Video Indexer

Using AI to make video content better.

Microsoft AI tracks and identifies who appears in a video, it transcribes and translates what they're saying, it understands the topics discussed, and detects changes in both scene and sentiment. After processing all of this data it provides the controls to allow the viewer to consume content in a more intuitive & desirable way.



#MSBuild Day 1 Keynote: Part
1



#MSBuild Day 1 Keynote: Part
2



#MSBuild Day 1 Keynote: Part
3



#MSBuild Day 1 Keynote: Part
4

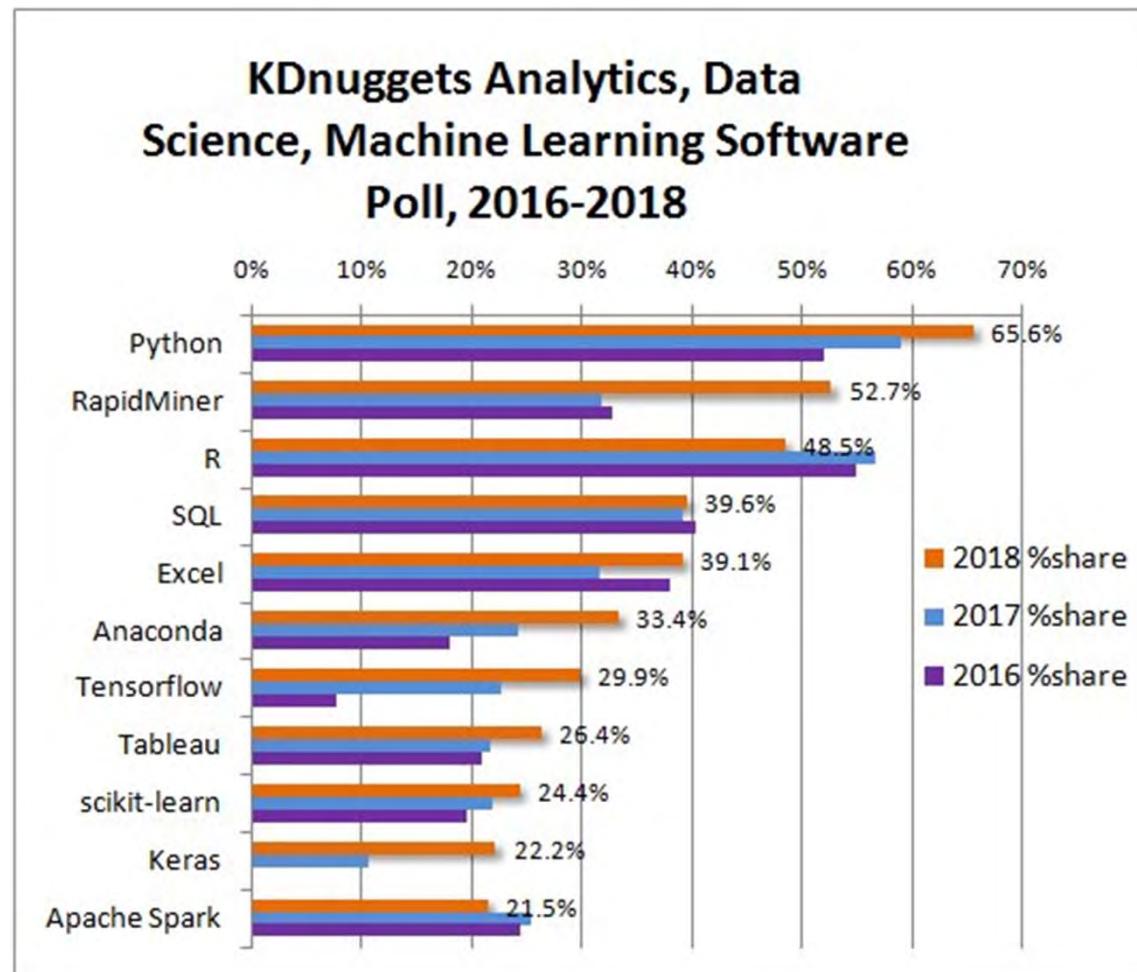


#MSBuild Day 1 Keynote: Part
5

Select one of the videos above to see it in action

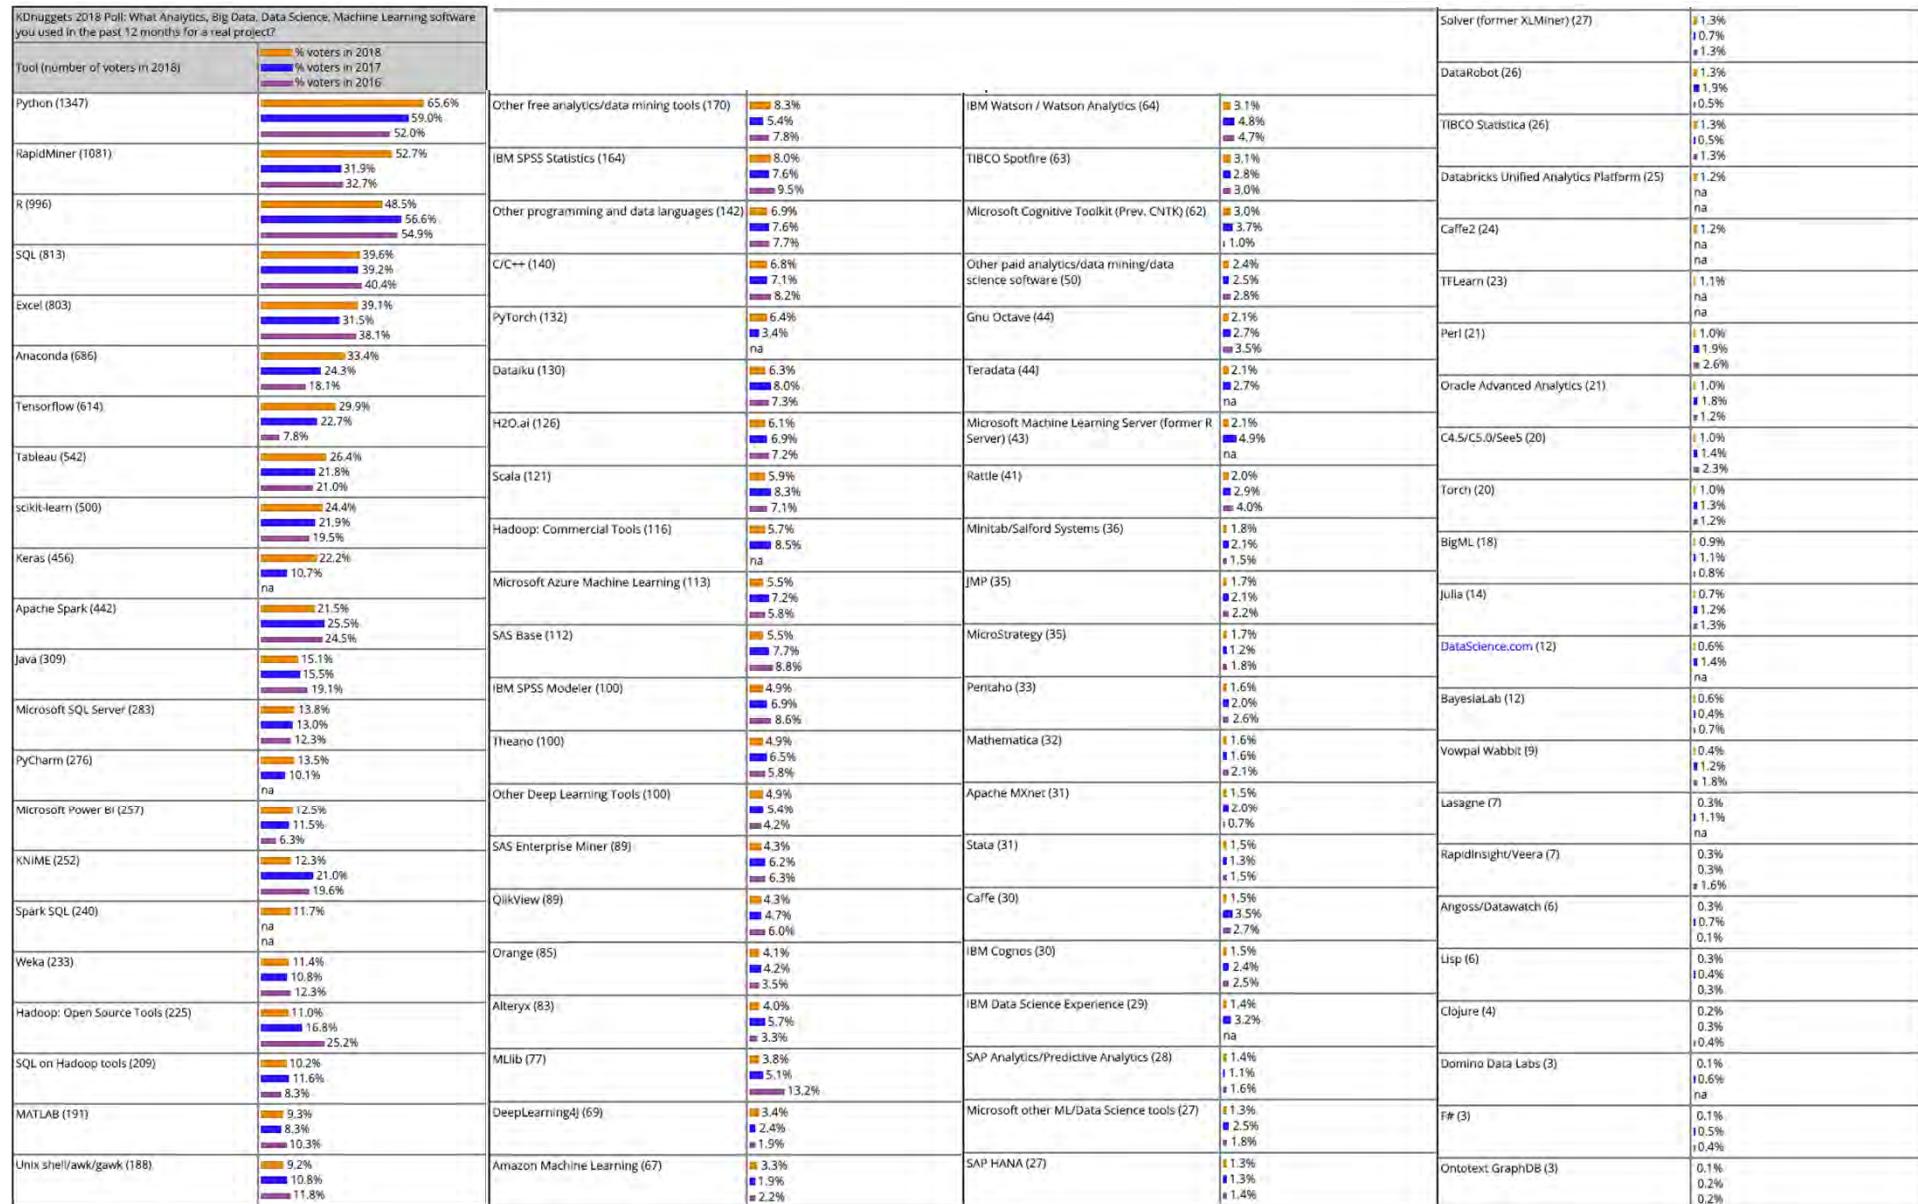


What are the popular tools?





What are the popular tools?





Graphical Tools (ML/DS)

RapidMiner

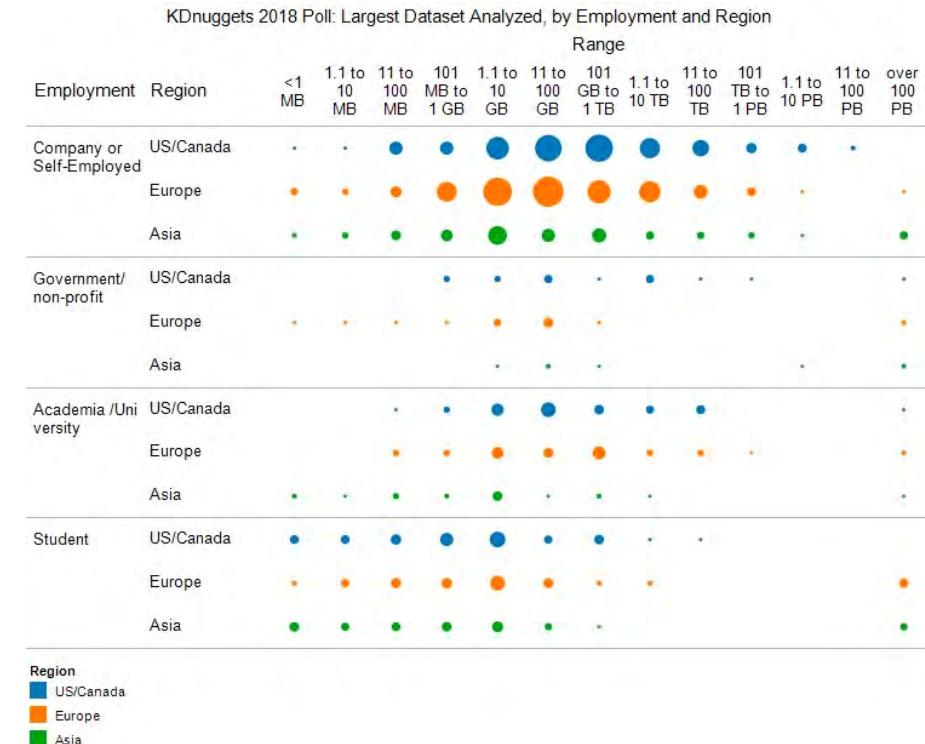
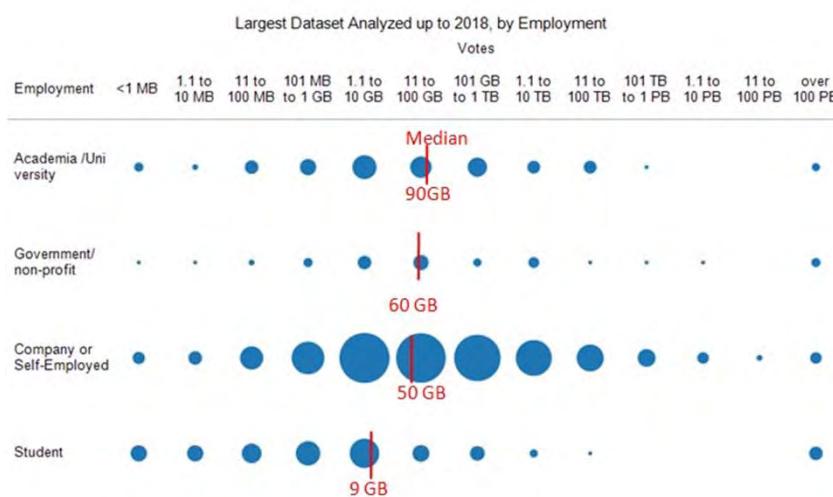
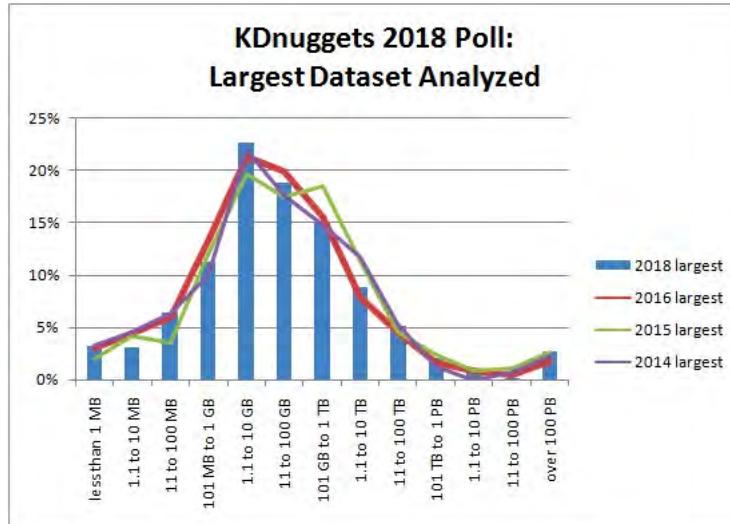
The screenshot shows the RapidMiner studio interface. On the left, there's a repository browser with various training resources and operators categorized by type (Data Access, Blending, Cleaning, Modeling, Scoring, Validation, Utility, Extensions). The main workspace displays a complex data mining process. It starts with two parallel retrieves: 'Loading NE Weather 2007' and 'Loading NE Flights 2007'. These feed into several sub-processes and data cleaning steps like 'Unify column types' and 'Origin equals BOS'. The processes involve various operators such as 'Generate Copy', 'Rename', 'Loop Attributes', 'Subprocess (2)', 'Filter Examples', and 'Subprocess (4)'. A 'Join' operator merges the weather and flight data. Finally, a 'Generate ID[DepID]' operator is used. A tooltip for the 'Subprocess' operator is open, providing its definition. The bottom part of the interface shows the KNIME Analytics Platform, featuring its own node-based workflow editor with nodes like CSV Reader, Scatter Plot, Logistic Regression Learner, and ROC Curve.

KNIME

The screenshot shows the KNIME Analytics Platform interface. On the left, there's a 'KNIME Explorer' showing projects and a 'Node Repository' listing various nodes categorized by type (IO, Manipulation, Views, Analytics, Database, Other Data Types, Structured Data, Scripting, Tool Integration, Community Nodes, KNIME Labs, Workflow Control, Social Media, Reporting, Chemistry, ChemAxon / Infocom). The main workspace displays a data mining workflow. It begins with a 'CSV Reader' node connected to a 'Column Filter' node, which then feeds into a 'Scatter Plot' node for exploratory analysis. Following this, the flow goes through 'Row Sampling', 'Transformations', 'Missing Value' handling, and 'Train Model' steps using a 'Logistic Regression Learner'. The workflow then branches into 'Validation Sample', 'Transformations', 'Missing Value' handling, and 'Treatment' steps. Finally, it uses a 'Score Model' node, a 'Regression Predictor', and a 'JavaScript ROC Curve' node to evaluate model performance. The bottom right shows the 'KNIME Console' displaying welcome messages and copyright information.



Datasets size





AI Software

As datasets get bigger the need for optimized software becomes increasingly important.

- The computation required increases with the size of the dataset.
- Software needs to be optimized for the underlying hardware to take full advantage of faster computers.
- The software and hardware need to be optimized for specific mathematical operations for optimal performance.



AI Software

It is usually inefficient for AI developers to write software to manage and analyze this data themselves.

- Developing software optimized for hardware and mathematical operations is time consuming.
- As methods become increasingly sophisticated the code for mathematics become more error prone.



AI Developer

An AI developer needs to leverage existing code so they can focus on the big picture and overall system framework.

- Use code that's optimized for the underlying hardware since AI tasks require a large amount of computation to complete.
- Simplify model development and training by providing high-level primitives for complex and error-prone mathematical transformations.
- Understand and use algorithms by stringing together API calls.
- Take advantage of software libraries and frameworks.





AI Frameworks

A **software framework** provides generic functionality that can act as a skeleton architecture to accomplish a particular task.

- A framework dictates the flow of control and defines the overall nature of the program where common design patterns can be reused.
- Frameworks can be extended by the users to provide specific functionality and customizations.
- Frameworks provide a standard way to build and deploy applications.
- Multiple frameworks can use the same underlying libraries.





Deep Learning Frameworks

Deep Learning frameworks vary in their level of functionality.

- Some allow you to define a neural network of arbitrary complexity from the most basic building blocks
- Others act as drivers or wrappers aimed at boosting developer productivity but are limited in their functionality due to the higher level of abstraction.



AI frameworks

Frameworks



TensorFlow



neon



Caffe



Chainer



torch



theano





Tensorflow

TensorFlow*: Python* based DL framework designed for ease of use and extensibility on modern deep neural networks.

- Open-sourced by Google in November 2015.
- As of May 2017, it now integrates optimizations for Intel® Xeon® processors.
- According to technology site KDNuggets*, as of 2018 TensorFlow is the most popular deep learning framework
- Based on defining static computation graphs and then having multidimensional arrays through the graph.





Tensorflow

TensorFlow* is based on the idea of defining static computational graphs and flowing data through them.

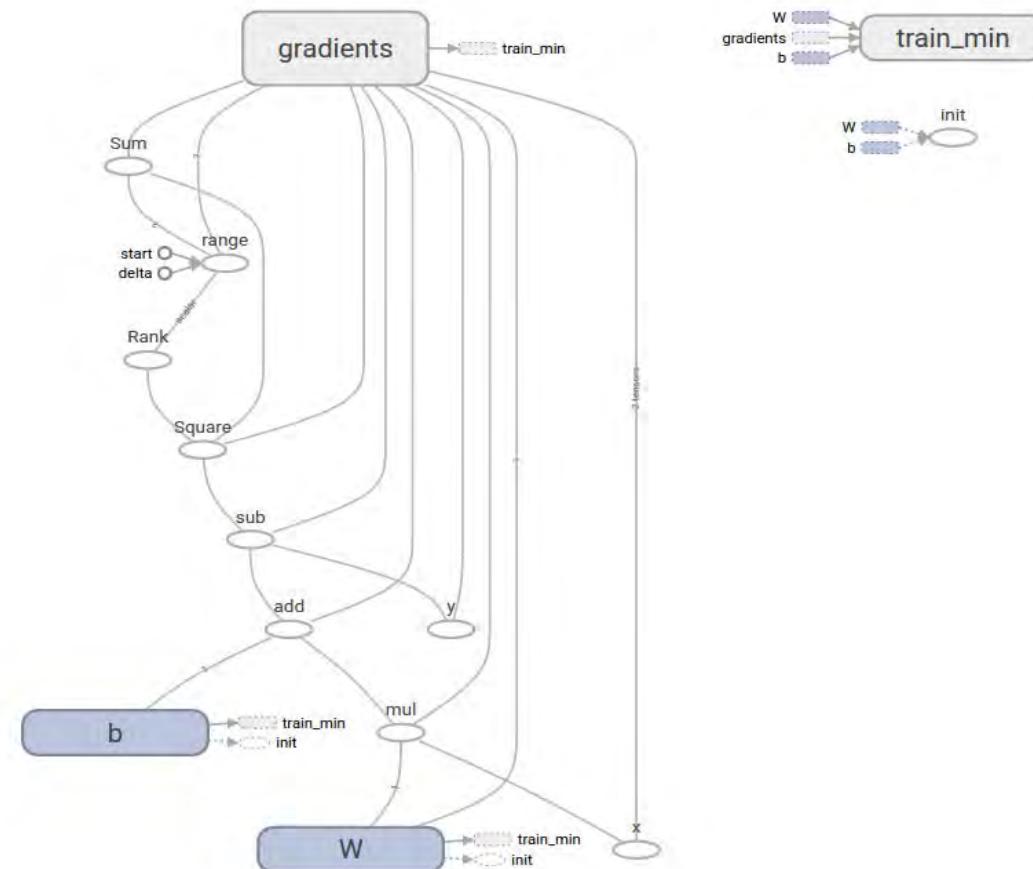
- Edges represent numerical data flowing through the graph.
- Nodes represent computations.
- Data is represented as tensors that are typed multidimensional arrays.
- These can be used to represent words, images, or any other data flowing through the network.
- TensorFlow uses “optimizers”, such as gradient descent, Adam, and RMSProp, to train neural networks.





Tensorflow

Computational graph
for linear regression.
Source: TensorFlow*





Caffe

Caffe* is a deep learning framework written in C++.

- Developed at the University of California Berkeley in 2014.
- Designed to run fast on GPUs, the most popular DL solution at that time.

Caffe





Caffe2

Caffe2* is an open source DL framework announced by Facebook in April 2017 with new features.

- Deployed at Facebook to help researchers train large machine learning models and deliver AI on mobile devices.
- Built with expression, speed, and modularity in mind.



Caffe2





Theano

Theano* is a Python* numerical computation library.

- Developed at the University of Montreal.
- Designed to run fast on GPUs.
- Improved performance on Intel® Xeon® and Intel® Xeon Phi™ processors with this fork of the popular Python library.
- Announced in September 2017 that major development will stop.



theano





MXNet

MXNet* is an open-source, deep learning framework.

- Apache* project originally developed within Carnegie Mellon University.
- Includes built-in support for the Intel® Math Kernel Library (Intel® MKL).
- Now includes Intel, Microsoft, and multiple universities as contributors.
- Characterized by nearly linear scaling with the number of GPUs used (doubling GPUs leads to nearly double the speed).
- Supports both imperative and symbolic programming.





PyTorch

- Python wrapper over Torch (C++ library) released in Jan 2017
- Also stores data and variables as Tensors
- Developed by Facebook and other companies
- Deep Learning operations are done outside of P
 - Graph based
 - Nodes are operations
 - Edges are multi-dimensional arrays called tensors

PYTORCH



neon™ Framework

neon™ is a deep learning solution optimized for Intel® architecture.

- Nervana™ Systems, founded in 2014, open-sourced neon™ in 2015.
- Nervana Systems joined Intel in August 2016.
- Posted top benchmarks on both GPU and CPU architectures.
- Used assembly level optimizations to produce faster speeds.





Keras

- Keras is high-level neural networks API, written in Python
- Built on top of either Theano, CNTK or TensorFlow
- Most powerful & easy to use for developing and evaluating deep learning models
- Allows for easy and fast Prototyping (through user friendliness, modularity, and extensibility)
- Support both convolutional networks and recurrent networks, as well as combinations of the two
- Runs seamlessly on CPU and GPU





Big Data

By 2020 the Internet of Things (IoT) will include:

- A projected 200 billion smart-and-connected devices.
- The data produced is expected to double every two years to total 40 zettabytes (40 trillion gigabytes).
- New frameworks are needed to process this data



Apache Spark

Apache Spark* is the most widely used framework for big data analytics.

- Open source big data cluster processing framework.
- Built around speed, ease of use, and advanced analytics.
- Big ecosystem that includes Spark Core*, SQL, Streaming, Mlib, and GraphX*.
- Integrates with cluster management and distributed storage.
- In-memory processing instead of saving data to hard disks.
- No major DL capabilities for Apache Spark.





BigDL

BigDL is a distributed deep learning library for Apache Spark*.

- Developed internally at Intel® and open-sourced December 2016.
- Efficiently scale-out to “Big Data Scale” using Apache Spark.
- High performance powered by Intel® MKL and multi-threaded programming, and implemented as a standalone library on Spark.
- Run DL applications as standard Spark™ programs, that can directly run on top of existing Spark or Hadoop* clusters.
- Feature parity with popular DL frameworks, such as Caffe* and TensorFlow*.
- Load pre-trained Caffe or Torch* models into Spark programs using BigDL.





BigDL: A Case Study

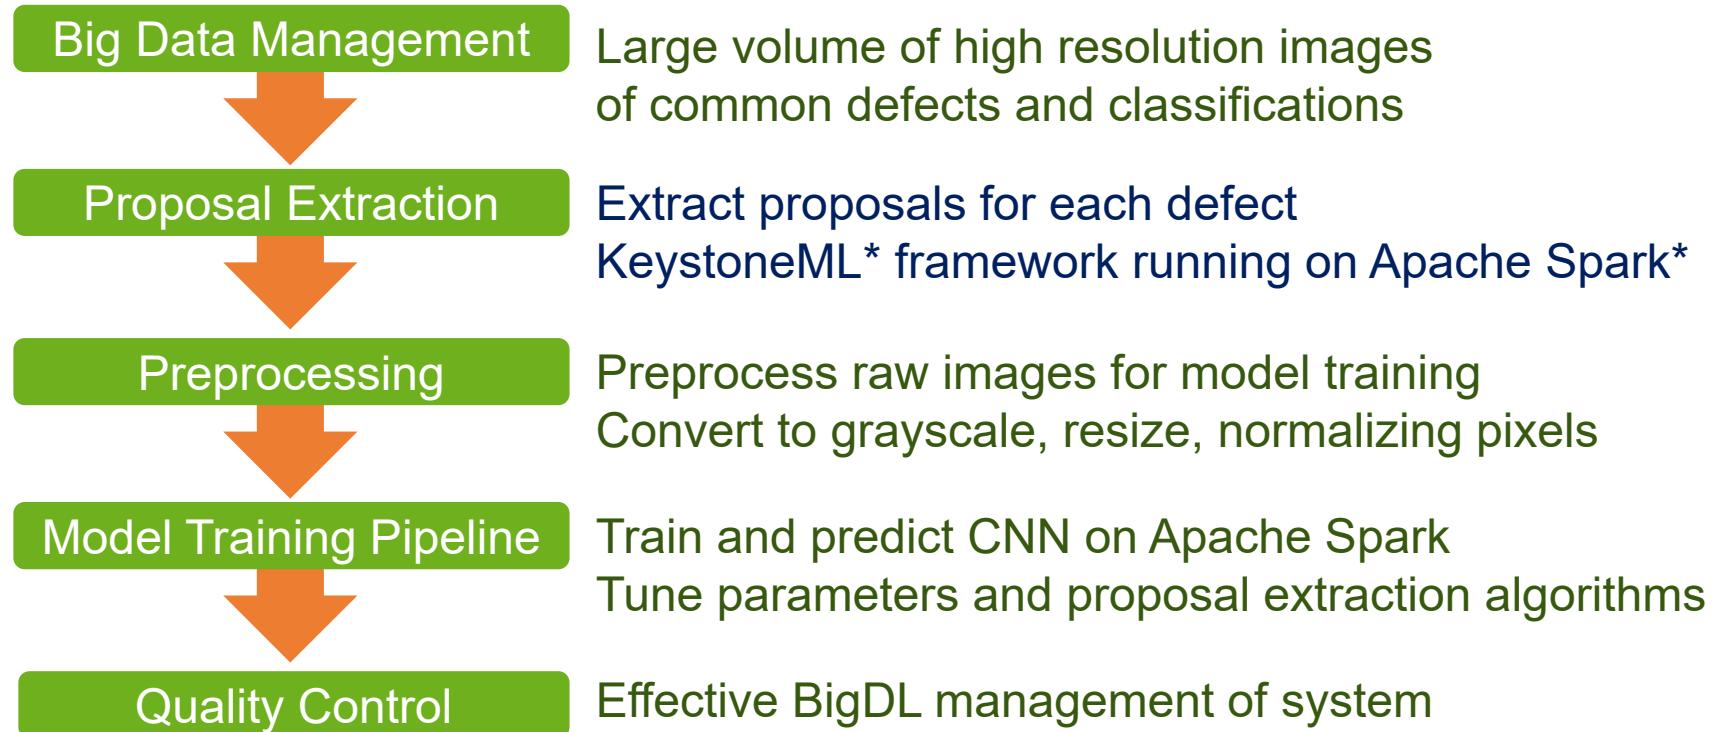
Intel worked with one of the largest steel producers to improve their end-to-end manufacturing pipeline.

- Goal: steel surface defect recognition for steel production and quality control from images.
- High resolution images were collected with cameras mounted on assembly line.
- Two part problem:
 1. Detect the defect in the image.
 2. Classify the defect type.





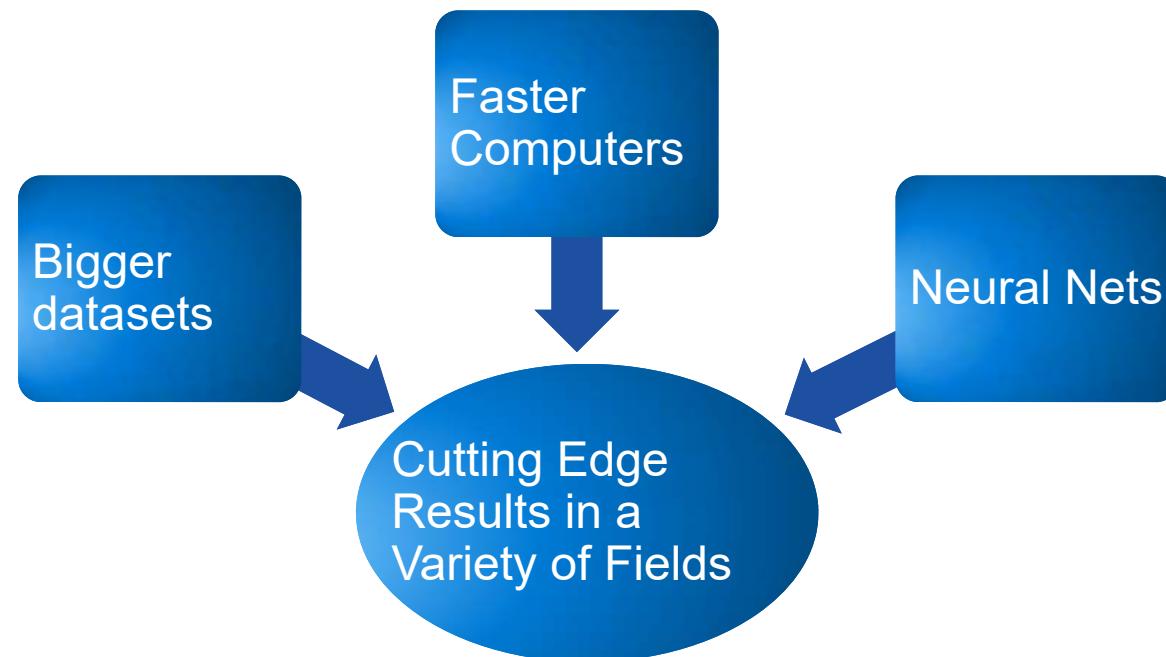
BigDL: Steel Defects Application





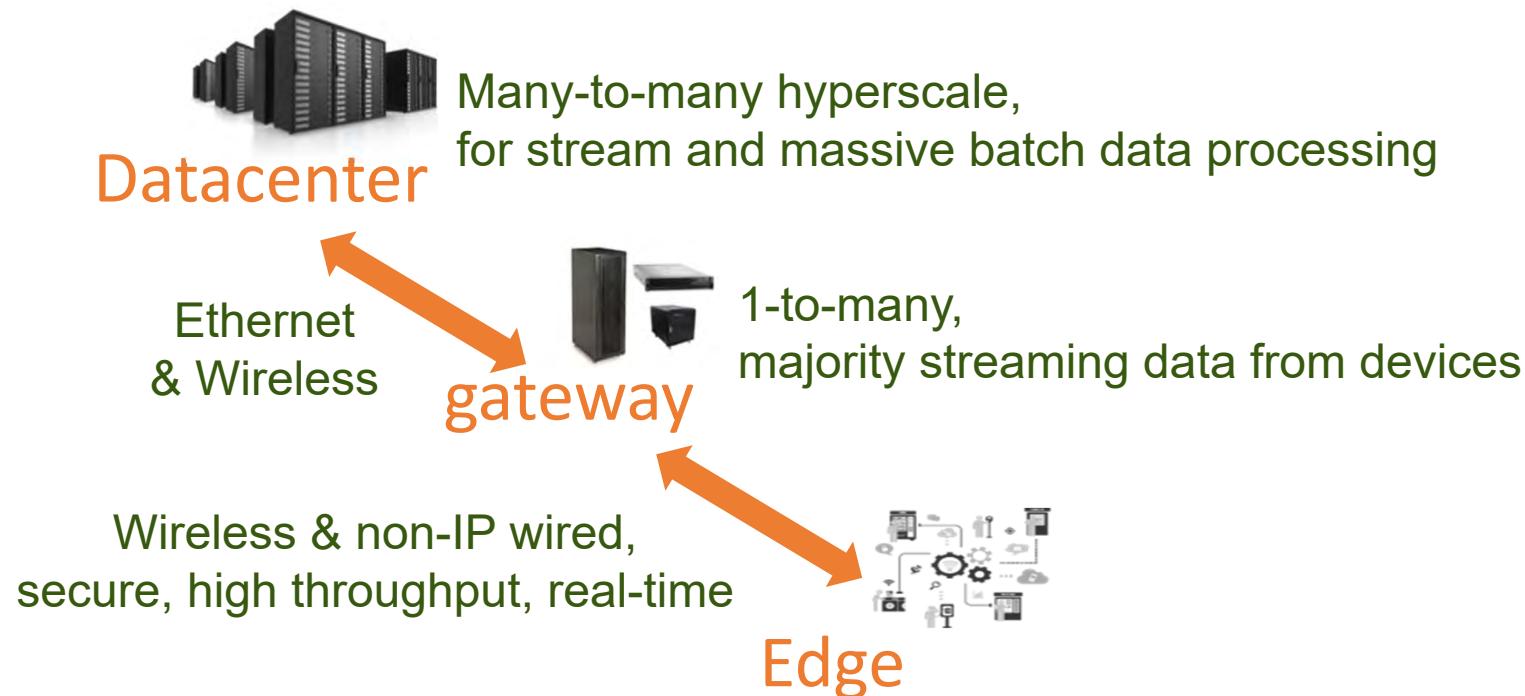
AI Hardware

Faster hardware is one of the key areas driving the modern era of AI.





End to End Computing for AI





End to End Computing for AI

There are various hardware requirements for different AI tasks.

- AI tasks can include data collection and fusion, training, and inference.
- The number of operations for training can be on the order of exaFLOPS, making this task more suited for datacenters.
- Inference takes fewer operations than training and can be done on both edge devices and at datacenters.
- Certain applications can have constraints on the processor size and power. For example, wearable devices and drones.



End to End Computing for AI

Datacenter

Datacenters are used to store and process data for applications, from websites to Internet of Things (IoT) systems.

- Low latency, high bandwidth – needed to access high compute resources.
- Reliability, lack of downtime, high performance.
- AI is the fastest-growing datacenter workload.
- A ~12X growth in demand by 2020 is expected.
- Training is typically done in a datacenter, allowing high processor power and physical size.



Datacenter

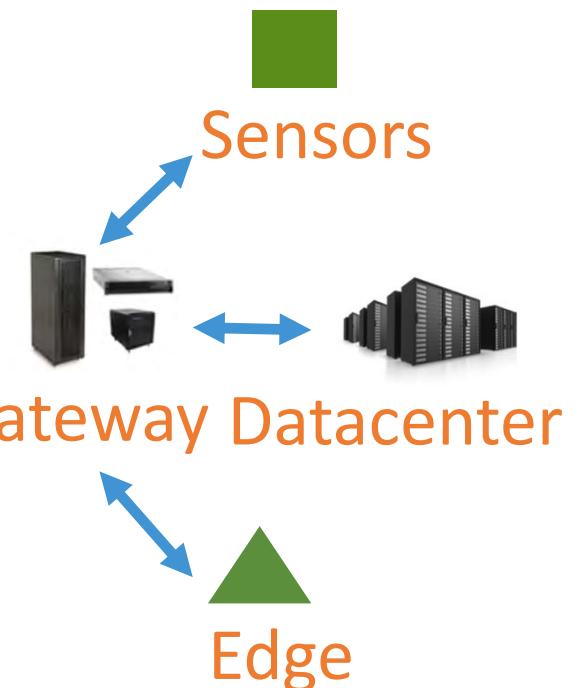


End to End Computing for AI

Gateway

Gateway computers route information from edge devices into and out of datacenters.

- High bandwidth – able to handle input from many devices and route correctly.
- Lightweight protocols – can't require extensive CPU usage, keeping gateway computer fast.
- Secure protocols – gateway computers are often a security point of failure.





End to End Computing for AI

Edge

Edge computing refers to computing happening as close as possible to where the computation is required.

- Sensors read information and compute a result directly on an Internet of Things (IoT) device, without sending information to data center.
- Reduces communication between sensors and datacenter.
- Uses resources that are not continuously connected to a network.
- Edge computing often used for inference rather than model training.





End to End Computing for AI

Example: Automated Driving

Autonomous vehicles produce ~4 terabytes of data per day.

One car, driving for one hour, requires ~5 exaFLOPS of computational power in order to safely keep it on the road.



DARPA Self-Driving
Car Challenge Won in 2005



End to End Computing for AI

Automated Driving

Vehicle requires:

- Human-Machine Interface (HMI) – to build trust between driver and vehicle and provide advanced virtualization and graphics capabilities
- Scalable, powerful in-vehicle computing
- Sensor processing
- Environment modeling
- Driving functions



Car

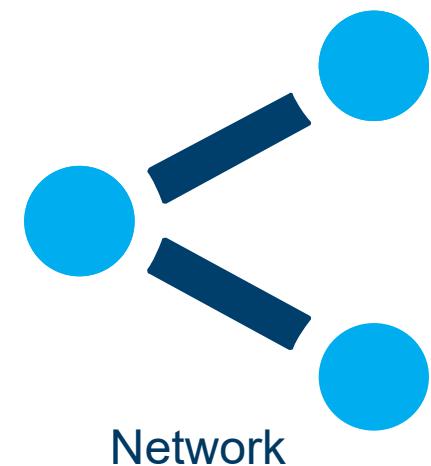


End to End Computing for AI

Automated Driving

Network requires:

- Vehicle-to-vehicle (V2V) and Vehicle-to-everything (V2X) communication
- Next-generation network connectivity
- Over-the-air-updates
- High-definition maps



Network

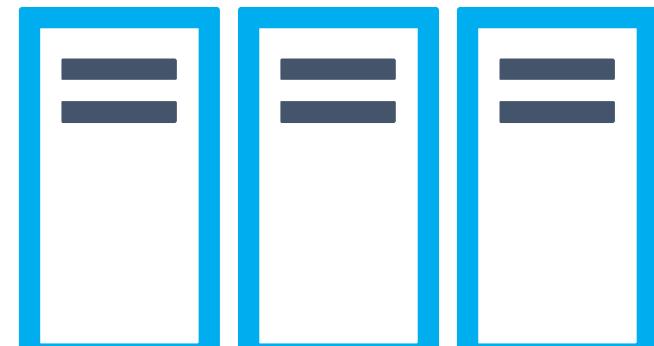


End to End Computing for AI

Automated Driving

Datacenter requires:

- High-performance data center and cloud
- AI to fleet management to data mining
- Model training
- Greater than real time inference



Datacenter



AI Hardware: FPGA

FPGA: Overview

Field Programmable Gate Arrays (FPGAs) are hardware that can be reconfigured – programmed in the field.

- FPGAs can become any digital circuit as long as the unit has enough logic blocks to implement that circuit.
- Enables the creation of custom hardware for individual solutions in an optimal way that other devices cannot efficiently support.
- High-throughput, low-latency processing of complex algorithms, such as neural networks.
- Flexible fabric enables direct connection of various inputs, such as cameras, without needing an intermediary.

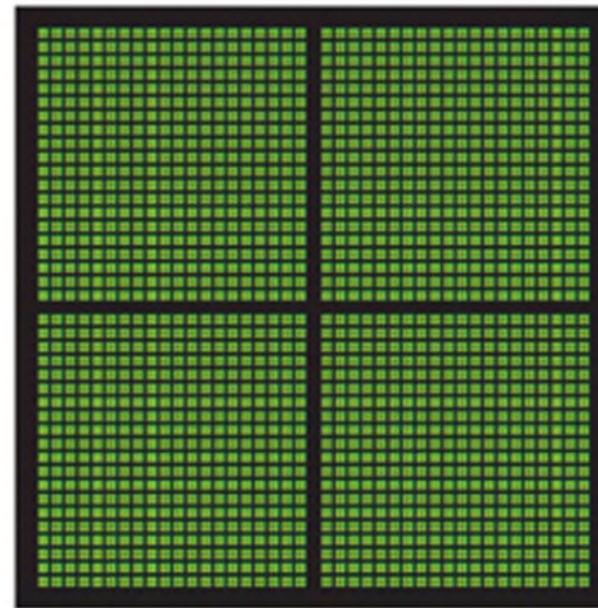


AI Hardware: GPU



CPU

DOZENS OF CORES



GPU

THOUSANDS OF CORES



AI Hardware: GPU

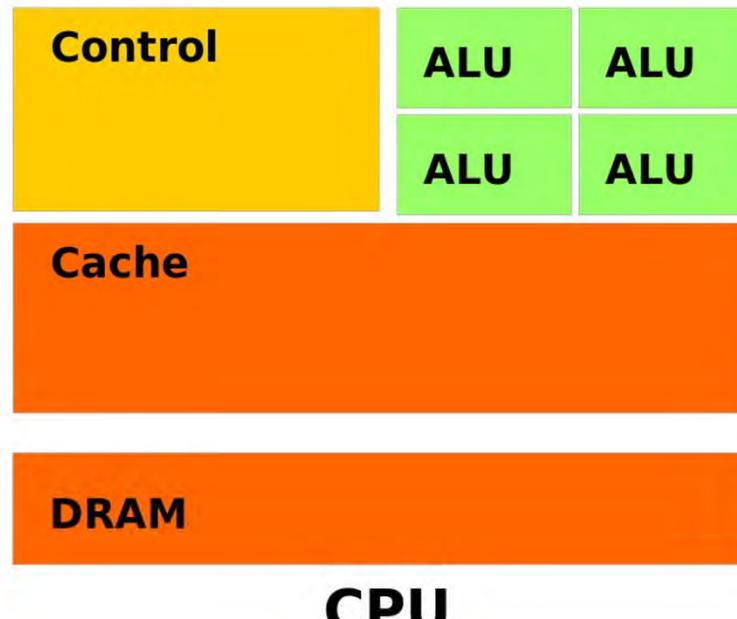
Originally designed to process graphics, GPUs have become a popular DL workhorse.

- Feature thousands of small, simple cores specialized for numeric, parallel computations.
- Many transistors dedicated to computation.
- GPUs excel at repeated similar instructions in parallel.
- Optimized for parallel data throughput computation.
- Major neural net breakthroughs since 2012 have been powered by GPU computations – performance has since increased >5x.
- Once the data is in the GPU memory the bottlenecks are small.
- NVIDIA is highly popular in discrete GPUs



AI Hardware: GPU

CPUs vs GPUs





AI Hardware: GPU

CPUs vs GPUs

While GPUs are specialized processors, CPUs are still the main computation engines on most computers.

- CPUs have dozens of cores compared to GPUs that have thousands of less powerful cores.
- CPUs have fewer arithmetic logic units (ALUs) and a lower compute density than GPUs
- CPUs are lower latency and have larger cache memory compared to GPUs





AI Hardware: GPU

CPUs vs GPUs

While GPUs are specialized processors, CPUs are still the main computation engines on most computers.

- GPUs are designed for parallel tasks and perform well when a single instruction is to be performed over a large amount of data.
- GPUs have additional overhead when copying data from main memory is required so CPUs can be better when a large number of memory swaps are needed.
- GPUs are not good for tasks that cannot be parallelized or when heavy processing on fewer data streams is needed. CPUs excel at serial tasks.
- CPUs are easy to program – popular programming languages compile to machine code, to be run on CPU by default.



AI Hardware: GPU

NVIDIA GPUs

NVIDIA's main business is in dedicated hardware for discrete graphics processors.

Several factors have made NVIDIA highly popular in this space.

- CUDA* framework which lets developers write code in C++, Python*, and other popular languages that run optimized operations on the GPU.
- Early incorporation of GPUs with cloud services, such as Amazon Web Services*.
- For several years, NVIDIA facilitated ecosystem adoption by optimizing their libraries and abstracting the complexities of the GPU through SDKs, tools, and libraries.

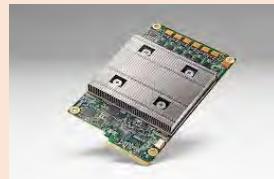




AI Custom Hardware

AI runs on both general purpose processors and on specialized hardware.

Specialized Hardware



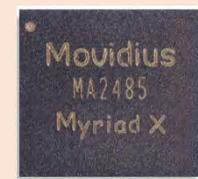
Google
TPU*



Apple Neural
Engine*



Intel®
Nervana™
NNP



Intel®
Movidius™
Myriad X



AI Custom Hardware

Google Tensor Processing Unit* (TPU)

Google's TPU* is the first major example of a chip designed for AI.

- Google's first-gen TPU (2015) was designed specifically for fast inference.
- Unlike a GPU, no capabilities for graphics tasks, such as rasterization or texture mapping.
- Second-gen (2017) increased the precision available for computations, making this version usable for training as well.





AI Custom Hardware

Apple Neural Engine*

Apple Neural Engine* system is on a chip designed for fast inference.

- Apple Neural Engine shipped on late 2017 iPhones.
- Separate from main CPU and GPU on iPhone.
- Specifically designed for fast neural net inference, capable of 600 billion operations per second.
- Highlighted importance of fast neural net inference to Apple's ability to provide its desired user experience.



AI Custom Hardware

Apple Neural Engine*

Apple Neural Engine* system is on a chip designed for fast inference.

- Apple Neural Engine shipped on late 2017 iPhones.
- Separate from main CPU and GPU on iPhone.
- Specifically designed for fast neural net inference, capable of 600 billion operations per second.
- Highlighted importance of fast neural net inference to Apple's ability to provide its desired user experience.

AI Services



- Google Search
 - Google Assistant ([hands on](#))
 - Google Photo (image recognition)
 - Speech Recognition
-
- Google's AI Services for Companies
 - <https://experiments.withgoogle.com/collection/ai>
 - Google's cloud-based AI Tools
 - <https://ai.google>
 - Google's AI Experiments:
 - <https://experiments.withgoogle.com/ai>
 - Do-it-Yourself AI:
 - <https://aiyprojects.withgoogle.com/voice/>



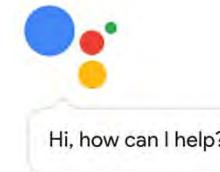
Voice Kit

Do-it-yourself intelligent speaker. Experiment with voice recognition and the Google Assistant.



Vision Kit

Do-it-yourself intelligent camera. Experiment with image recognition using neural networks.

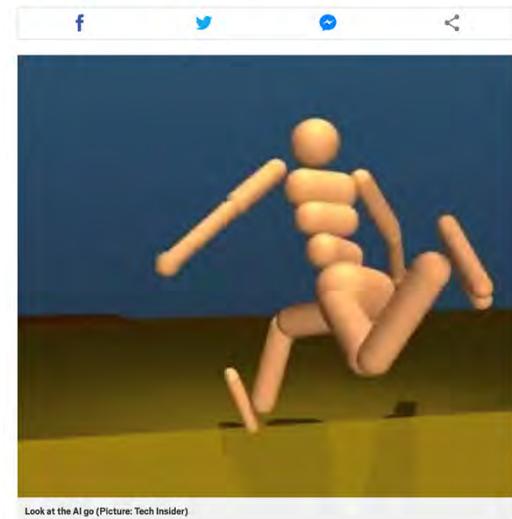


Meet your Google Assistant.

Ask it questions. Tell it to do things. It's your own personal Google, always ready to help.

Beware – Google's AI is so smart it just taught itself to walk without any human help

 **Jimmy Nsubuga** Monday 17 Jul 2017 6:31 pm



Look at the AI go (Picture: Tech Insider)

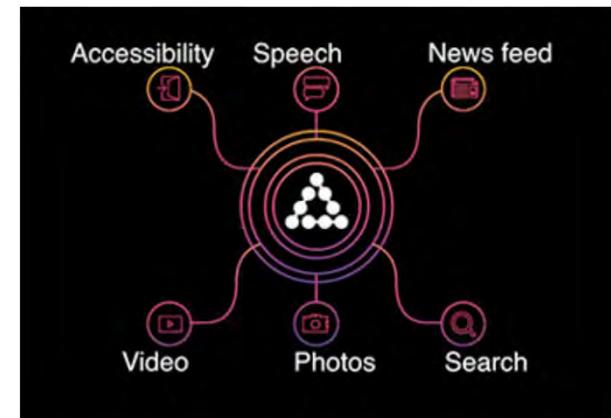


AI Services

- Facebook Photo search
- FB Learner Flow (<https://code.fb.com/ml-applications/introducing-fblearner-flow-facebook-s-ai-backbone/>)
- Text Analytics (Deep Text)
- Pattern Recognition to Prevent Suicides
- Improving 360 Degree Photos
- Computer Vision
- Facebook Personal Assistant M (experiment)
- Facebook Messenger Platform Chatbot
- Facebook's AI research Activities
 - <https://research.fb.com/category/facebook-ai-research/>

Facebook will use AI to help correct skewed 360-degree photos

The company has a technique for dealing with big file sizes, too.



TECHNOLOGY NEWS NOVEMBER 28, 2017 / 12:05 AM / A YEAR AGO

Facebook to expand artificial intelligence to help prevent suicide

David Ingram

3 MIN READ





AI Services

- Amazon Recommended Products
- Alexa Personal Assistant
- Cloud Storage
- Amazon's AI platform:
 - Amazon Lex
 - Amazon Polly
 - Amazon Recognition

LEARNING TOOLS

Get deep with machine learning

AWS DeepRacer

AWS DeepRacer is a fully autonomous 1/18th-scale race car designed to help you learn about reinforcement learning through autonomous driving.

- Experience the thrill of the race in the real world when you deploy your RL model onto AWS DeepRacer
- Build models in Amazon SageMaker and then train, test, and iterate on the track using the AWS DeepRacer 3D racing simulator
- Starting in 2019, compete in the world's first global autonomous racing league, to race for prizes and a chance to advance to win the coveted AWS DeepRacer Cup

[Learn more »](#)

AWS DeepLens

AWS DeepLens is the world's first deep learning-enabled video camera for developers. Integrated with Amazon SageMaker and many other AWS services, it allows you to get started with deep learning in less than 10 minutes through sample projects with practical, hands-on examples.

- Choose your deep learning model from the AWS DeepLens pre-trained model library, or your own model trained with Amazon SageMaker.
- Deploy your model to the device with a single click.
- Watch the results in real time in the AWS Management Console.

[Learn more »](#)



 Recommendations Personalize experiences for your customers with the same recommendation technology used at Amazon.com. AMAZON PERSONALIZE »	 Forecasting Build accurate forecasting models based on the same machine learning forecasting technology used by Amazon.com. AMAZON FORECAST »	 Image and Video Analysis Add image and video analysis to your applications to catalog assets, automate media workflows, and extract meaning. AMAZON REKOGNITION »
 Advanced Text Analytics Use natural language processing to extract insights and relationships from unstructured text. AMAZON COMPREHEND »	 Document Analysis Automatically extract text and data from millions of documents in just hours, reducing manual efforts. AMAZON Textract »	 Voice Turn text into lifelike speech to give voice to your applications. AMAZON POLLY »
 Conversational Agents Easily build conversational agents to improve customer service and increase contact center efficiency. AMAZON LEX »	 Translation Expand your reach through efficient and cost-effective translation to reach audiences in multiple languages. AMAZON TRANSLATE »	 Transcription Easily add high-quality speech-to-text capabilities to your applications and workflows. AMAZON TRANSCRIBE »

<https://aws.amazon.com/machine-learning/>



Amazon Rekognition

Amazon Rekognition Video をご紹介します。

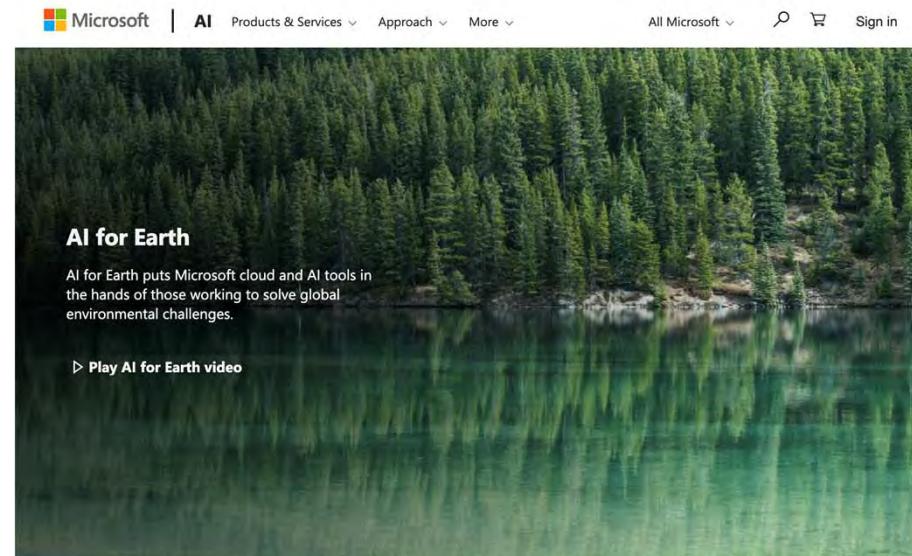


AI Services

- Cortana - <https://www.microsoft.com/en-us/windows/cortana>
- Presentation Translator -
<https://translator.microsoft.com/help/presentation-translator>
- HoloLens
- InnerEye - <https://www.microsoft.com/en-us/research/project/medical-image-analysis/>
- Azure Microsoft Cloud Service –
- AI for Earth -
<https://www.microsoft.com/en-us/ai/ai-for-earth>
- AI Language Translator -
<https://www.microsoft.com/en-us/translator/>
<https://www.bing.com/translator> (demo)



Microsoft



The screenshot shows the Microsoft AI for Earth landing page. At the top, there's a navigation bar with the Microsoft logo, a search icon, and links for 'Products & Services', 'Approach', 'More', 'All Microsoft', and 'Sign in'. Below the navigation is a large, scenic photograph of a forest reflected in a calm lake. Overlaid on this image is the text 'AI for Earth' and a brief description: 'AI for Earth puts Microsoft cloud and AI tools in the hands of those working to solve global environmental challenges.' A 'Play AI for Earth video' button is also visible. At the bottom of the page, under the heading 'Areas of focus', it says: 'AI for Earth awards grants to projects that use artificial intelligence to address four critical areas that are vital for building a sustainable future.' A link 'Learn about AI for Earth grants' is provided.

Microsoft | AI Products & Services Approach More All Microsoft Sign in

AI for Earth

AI for Earth puts Microsoft cloud and AI tools in the hands of those working to solve global environmental challenges.

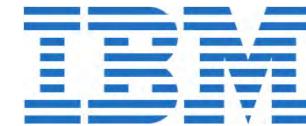
Play AI for Earth video

Areas of focus

AI for Earth awards grants to projects that use artificial intelligence to address four critical areas that are vital for building a sustainable future.

Learn about AI for Earth grants >

AI Services



- Cognos Analytice - <https://www.ibm.com/us-en/products/cognos-analytics>
- Tone Analyzer - <https://tone-analyzer-demo.ng.bluemix.net> (demo)
- Discovery - <https://discovery-news-demo.ng.bluemix.net>
- Visual Recognition -
<https://www.ibm.com/watson/services/visual-recognition/demo/#demo>
- Text to Speech - <https://text-to-speech-demo.ng.bluemix.net/> (audio streaming does not work on mobile browser)

IBM Watson Developer Cloud

Tone Analyzer

This service uses linguistic analysis to detect joy, fear, sadness, anger, analytical, confident and tentative tones found in text.

*This system is for demonstration purposes only and is not intended to process Personal Data. No Personal Data is to be entered into this system as it may not have the necessary controls in place to meet the requirements of the General Data Protection Regulation (EU) 2016/679.

Sample use cases

Choose an example to learn how you can adjust the tone of your content to change people's perceptions, or improve its effectiveness.

Learn more.

Tweets Online Review Email message Product Review in French Your own text

Analyzing Customer Engagement Data? Try out the [Tone Analyzer Customer Engagement Endpoint](#).

Analyze

Insurance (Custom Classifier)

Custom Classifier trained on insurance images

Category	Score
vandalism	0.64
flat_tire	0.53
broken_windshield	0.11
motorcycle_accident	0.06

International vehicle glass repair company
Belron uses Custom Models to automatically generate estimates of repair costs based on customer-submitted images of car damage.

Select an image on the left to evaluate how this Custom Model analyzes different images



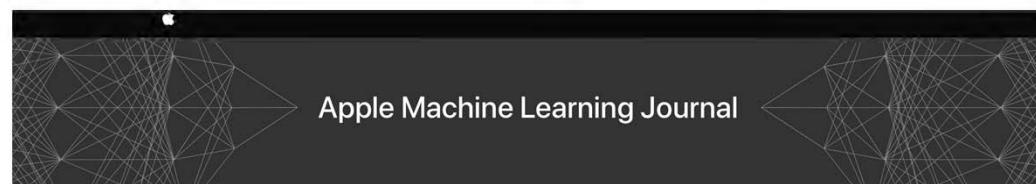
AI Services

- Speech Recognition on Siri
- QuickType
- A11 Bionic Chip – Core ML
- Apple Music
- Apple HomePod
- Apple Photos

Apple has new self-driving car hardware covered with iPod-style white plastic

Kif Leswing, Business Insider US

October 18, 2017



Apple acquires AI tech that seeks to understand your photos

Regaind can tell good pics from bad ones, and interpret what's going on.

71,188 views | Dec 26, 2016, 07:05am

Apple Publishes Its First Artificial Intelligence Paper



Aaron Tilley Forbes Staff

Optimizing Siri on HomePod in Far-Field Settings

Vol. 1, Issue 12 • December 2018
by Audio Software Engineering and Siri Speech Team



AI Services

Robots Run the Warehouses ([link](#))

Innovation

Alibaba lets AI, robots and drones do the heavy lifting on Singles' Day

This year's November 11 shopping ritual will engage a recommendation algorithm, robots, and chatbots capable of understanding human emotion

Topic | Singles' Day (11.11)

SMART CUSTOMER SERVICE

Ali Assistant is a chatbot that handles both spoken and written queries, acting as customer-service rep and personal shopping assistant. It is capable of handling up to **95%** of customer service enquiries.



Intelligent Machines

Alibaba's AI Fashion Consultant Helps Achieve Record-Setting Sales

AI will blur the line between online and offline retail.

BIG DATA

With nearly **500 million** active users across its websites and apps, Alibaba has a vast repository of consumer data that can be processed and analyzed by AI programs continuously in real time, leading to increasingly accurate predictions and a better shopping experience.



COMPUTING POWER

Alibaba has built up one of the world's largest networks of interconnected computer servers to run its e-commerce empire, backed by an operating system that can process more than **175,000 transactions per second**.

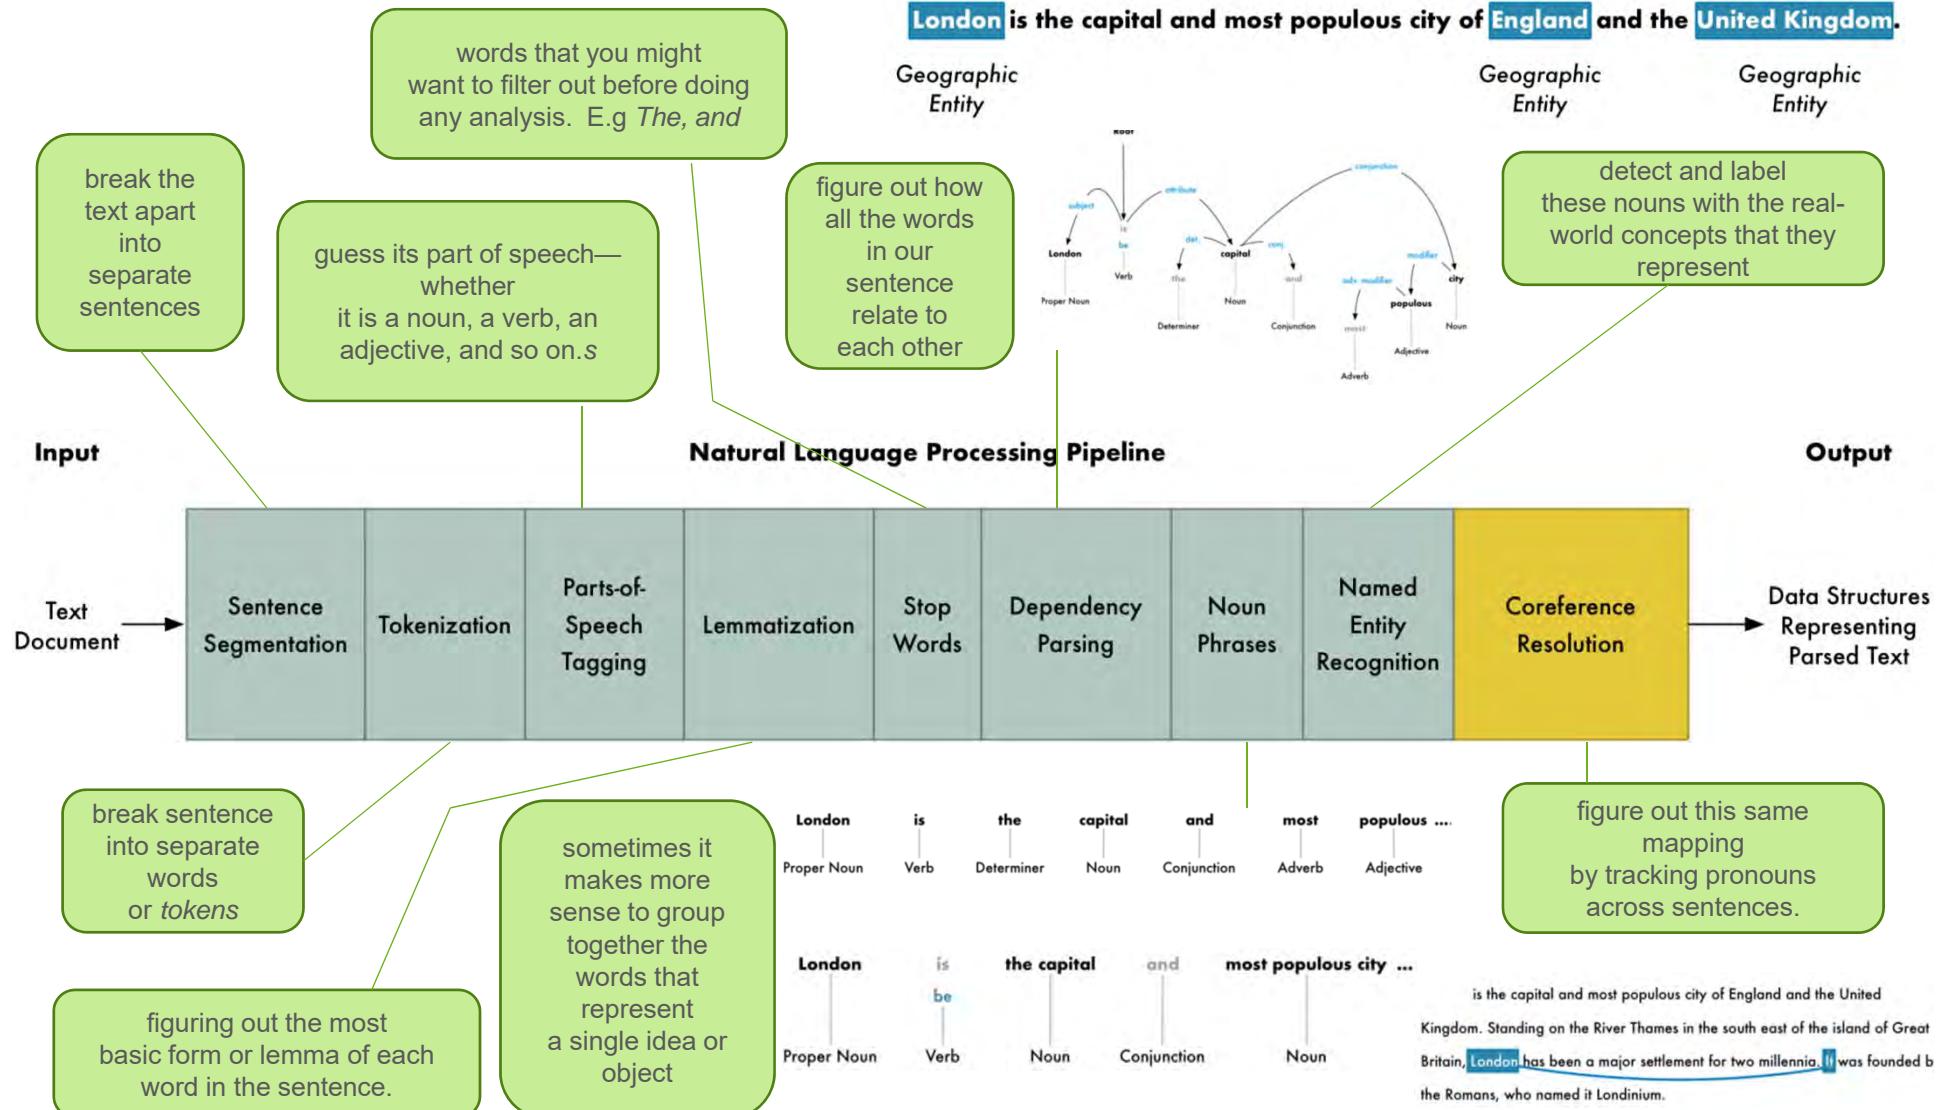


Function specific AI services

Sales	Outreach.io
Virtual Human	https://www.quantumcapture.com/ctrl-human
HR Services	https://leena.ai/HR-FAQ
Scheduling	https://x.ai/how-it-works/
Enterprise support functions	https://www.soapbox.ai/
Sales Bots	https://octaneai.com/
AI-Powered Transcription	http://capiro.ai/index.html
Hiring	https://hiringsolved.com/product
Programming	https://www.codota.com/enterprise https://kite.com/



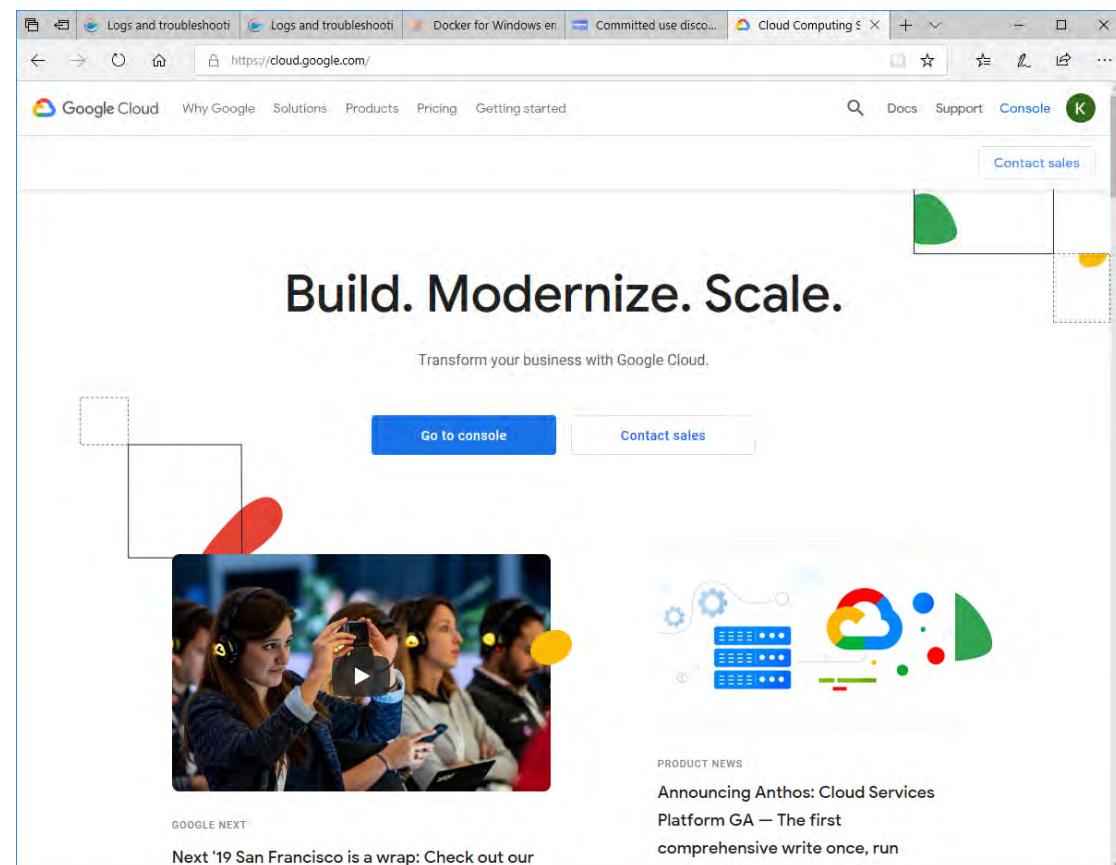
Hands-on NLP





Hands-on NLP

- Use Google Cloud Platform, we will:
 - Classify Text
 - Named Entity Recognition
 - Sentiment Analysis
 - Syntax Analysis

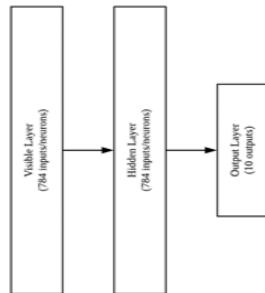


<https://cloud.google.com/natural-language/docs/apis>

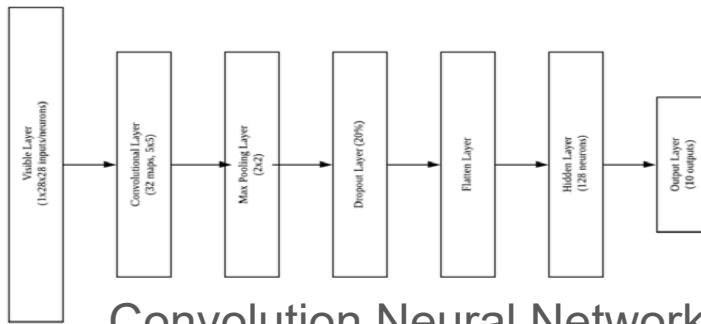
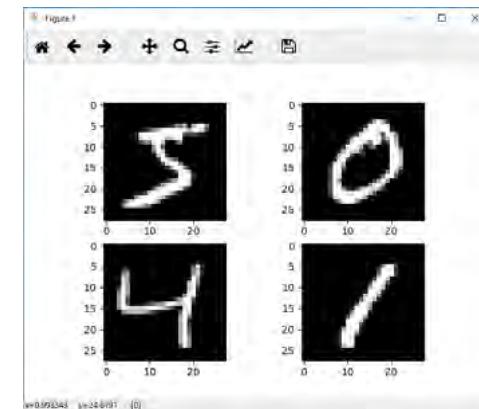


Hands on – Object Recognition

- Recognition handwritten digits (MNIST dataset)
- Using a Multilayer Perceptrons Neural Network
- Using a simple Convolution Neural Network



Multilayer Perceptrons
Neural Network



Convolution Neural Network

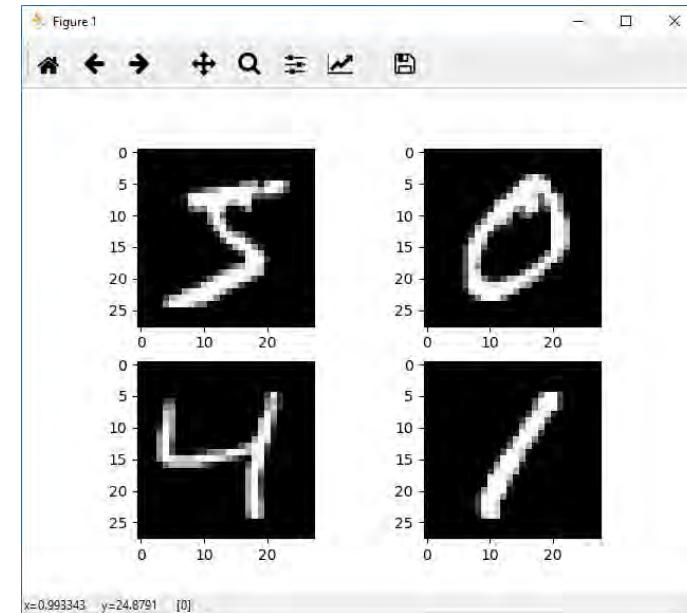




Hands on – Object Recognition

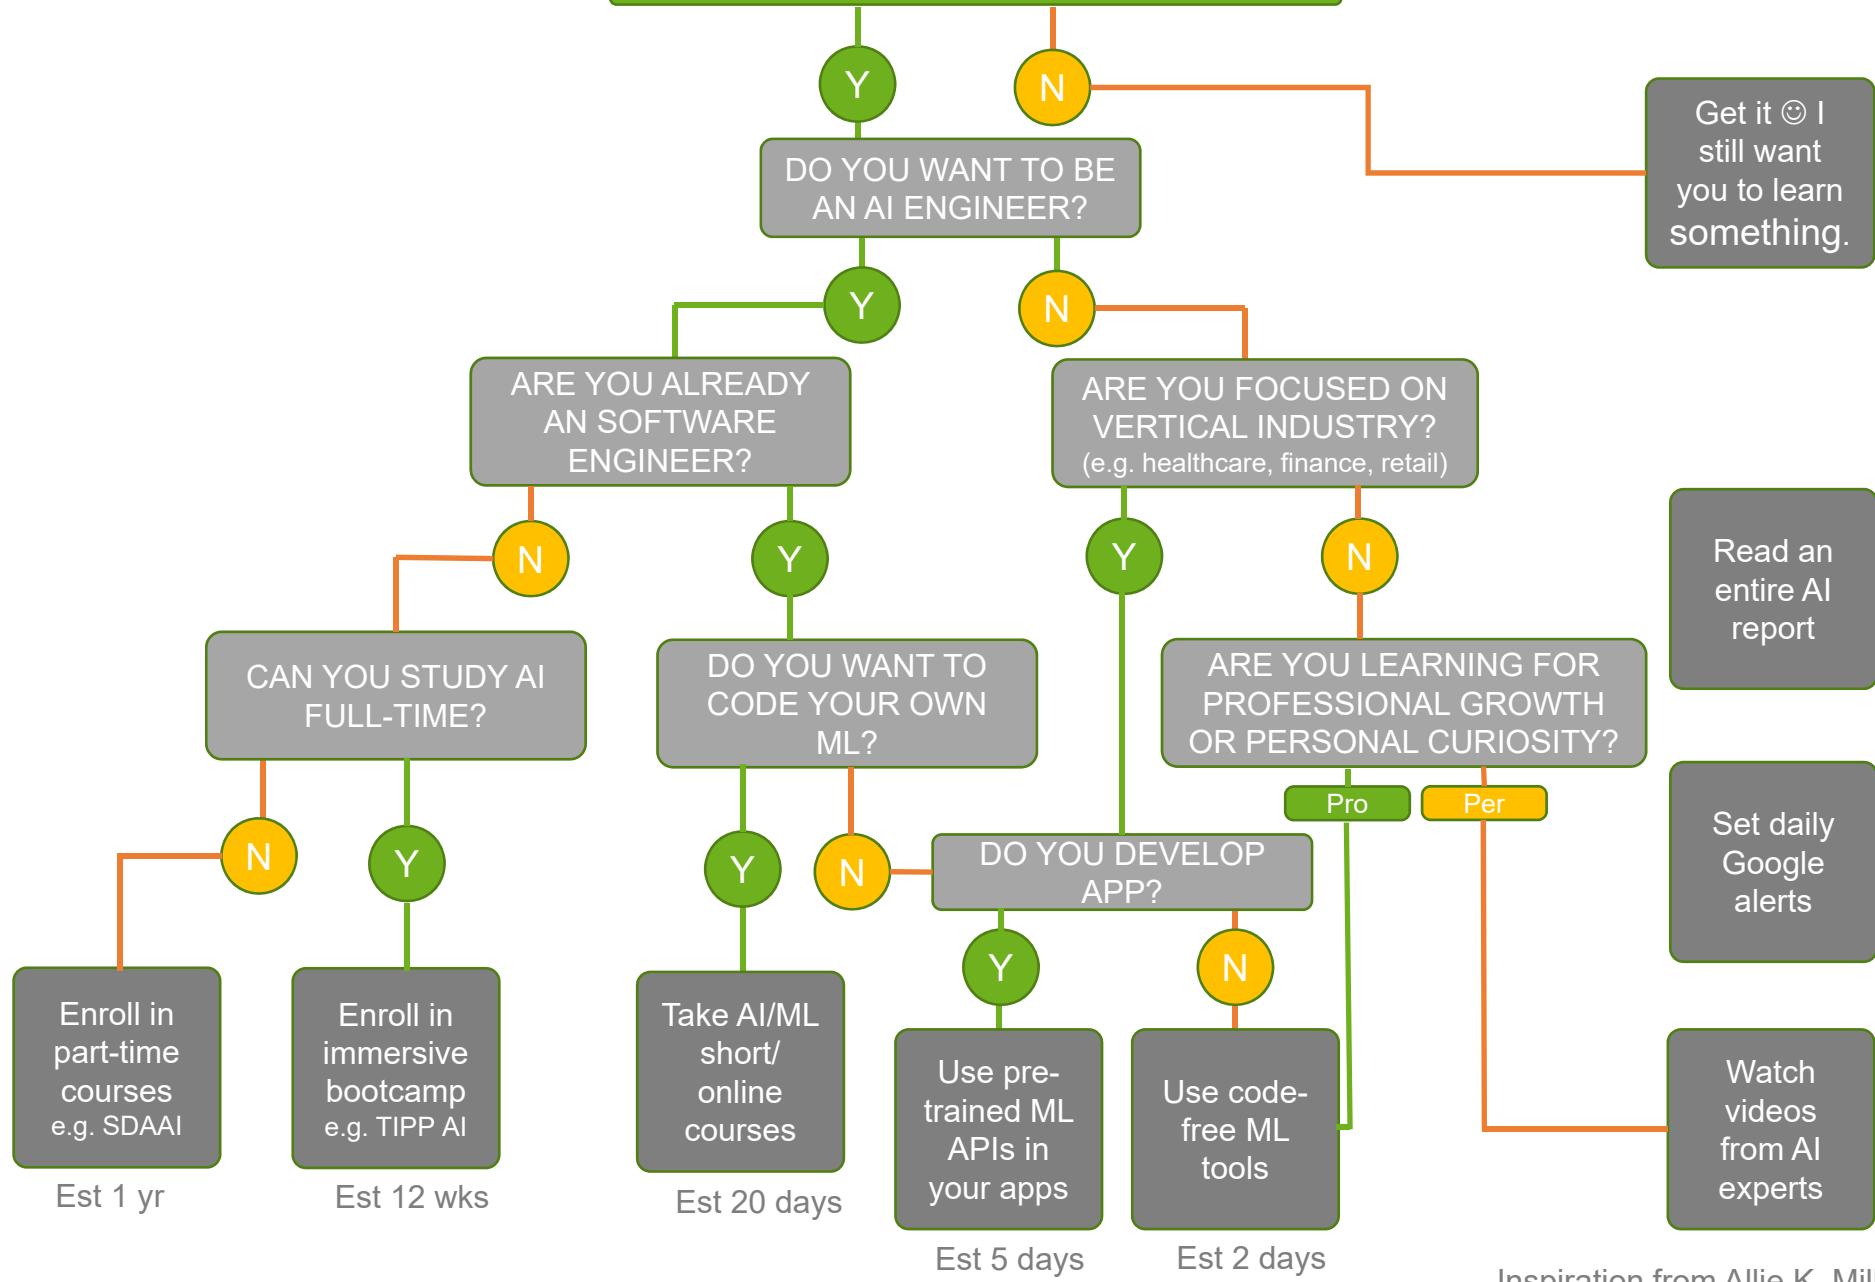
- Workflow

1. Load Data.
2. Define Model.
3. Compile Model.
4. Fit/Train Model.
5. Evaluate Model.
6. Predict with new data.





DO YOU WANT TO LEARN AI?





Where to go from here?

MOOC:
DataCamp
<https://www.datacamp.com/>

Subscribe now. Save 50% on DataCamp and skill up.
Offer ends in 04 days 04 hrs 56 mins 51 secs >

57 results for "python"

Intro to Python for Data Science

Intermediate Python for Data Science

Python Data Science Toolbox (Part 1)

Deep Learning in Python

Supervised Learning with scikit-learn

pandas Foundations

Edx
<https://www.edx.org/>

Udemy (freemium course)
<https://t.me/freecourse>

AI4T Telegram Channel
<http://bit.ly/2Y6j5GR>

edX Courses Programs & Degrees Schools & Partners edX for Business Search Sign In Register

Accelerate your future. Learn anytime, anywhere.

Find courses

What do you want to learn?

THE COUNTDOWN IS ON!
Get 15% off your purchase.

CYBER MONDAY

Prominent Subjects



\$5000 ELP SIGN-ON INCENTIVE

REPUBLIC POLYTECHNIC
Singapore's SkillsFuture

Specialist Diploma in APPLIED ARTIFICIAL INTELLIGENCE

➤ ARTIFICIAL Intelligence (AI) is a frontier technology that has been identified in Singapore's Industry Transformation Map (ITM) as an enabler of growth and jobs in Singapore. To achieve this objective, manpower skilled with AI knowledge and abilities is required to contribute and develop the AI ecosystem in Singapore.

The Specialist Diploma in Applied Artificial Intelligence (SDAAI) at Republic Polytechnic will equip participants with the requisite concepts and theories that underpin AI and to develop in participants, the fundamental practical skills and knowledge for developing AI solutions. Participants will undertake practical hands-on case studies and projects in clearly identified areas in which AI is applied today, namely the application of recommender systems, virtual assistants and pattern recognition and anomaly detection.

REGISTER NOW
<http://bit.ly/2IArFxw>

COMMENCEMENT DATE:
OCT 2018

SkillsFuture Earn and Learn Programme (ELP) information:
<http://bit.ly/2IArFxw>

SOI For more info:
Email: ACE-help@rp.edu.sg
Tel: 6510 3000

Specialist Diploma in Applied Artificial Intelligence

Total Hours = 270 hours

Time to Complete = 12 months

PDC in Fundamentals of Artificial Intelligence (120 hours)	PDC in Applications of Artificial Intelligence (150 hours)
Math for Machine Learning (30 hours)	Pattern Recognition and Anomaly Detection (30 hours)
Introduction to Programming (30 hours)	Recommender Systems (30 hours)
Introduction to Data Management for Machine Learning (30 hours)	Virtual Assistants (30 hours)
Machine Learning Fundamentals (30 hours)	Capstone Project* (60 hours)



Tech Immersion and
Placement Programme in

Applied Artificial Intelligence

In Collaboration With

Programme Overview



THIS programme is designed to equip participants with practical, industry-relevant skills and knowledge to undertake the role of an AI Application Developer or System Integrator or as an AI or Machine Learning Engineer.

It is structured as a 3-month / 12-week programme where participants complete 10 modules. The first 9 modules will be conducted during an 8-week facilitation and e-learning training phase before leading up to the final 4-week Industry Capstone Project.

Participants who successfully complete the programme will be awarded a Professional Certificate of Completion in Applied Artificial Intelligence from Republic Polytechnic (in collaboration with AI Singapore and Microsoft).

Programme Structure

8 WEEKS TRAINING:

- 1 Foundations to Artificial Intelligence
- 2 Python for Data Science
- 3 Mathematics for Machine Learning
- 4 Ethics and Law in Data and Analytics
- 5 Research Methods for Data Science
- 6 Machine Learning Fundamentals
- 7 Deep Learning Fundamentals
- 8 Reinforcement Learning Fundamentals
- 9 Applied AI Solutions Development

4 WEEKS PROJECT:

- 10 Applied AI Industry Capstone Project

Key Features

- Collaborative programme between Republic Polytechnic, AI Singapore and Microsoft Singapore
- Industry-relevant programme outcomes based on the Skills Framework for Infocomm Technology (ICT) professionals for AI
- Practical and innovative pedagogy that infuses classroom learning with real-world applications
- Real-world and industry-sourced Capstone Projects
- Guest talks and seminars from industry experts and practitioners
- Bundled Microsoft's Professional Program in Artificial Intelligence certification vouchers
- Jobs placement opportunities and facilitation upon completion of programme
- Subsidy support under IMDA's Tech Immersion and Placement Programme (TIPP)

Lifelong Learning

- <https://www.rp.edu.sg/soi/lifelong-learning>

Short Courses



SOI offers an extensive variety of short, industry-relevant courses for ICT skills upgrading and skills acquisition. Our courses are categorized under different areas, ranging from Artificial Intelligence (AI), Business Intelligence/Business Analytics (BI/BA), Business Processes (BP), Unmanned Aerial Vehicle (UAV), IT Security, New/Digital Media, Software Development to the Internet of Things (IoT). To view our short course offerings, click on the relevant tab below.

AI

Data Analytics

IT Security

DevOps

Software Development

New/Digital Media

UAV

RPA

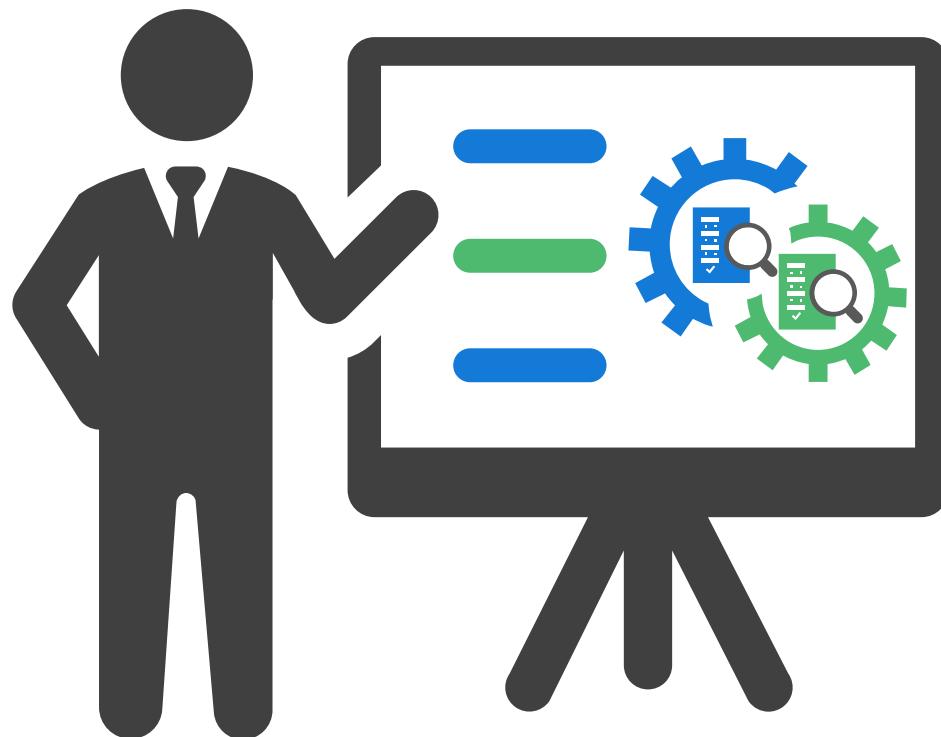
+ Artificial Intelligence for Everyone - A Practical Experience (1 day Beginner)

+ Artificial Intelligence for Techies - A Hands-On Approach (1 day Beginner)

+ An Introduction to Code-Free Machine Learning (1 day Beginner)



Summary



Email
seow_khee_wei@rp.edu.sg

Telegram
[@kwseow](https://t.me/kwseow)

Source code:



Thank you